

Executive AI Transformation Proposal



सिद्धिमूलं प्रबन्धनम्
भा. प्र. सं. इन्दौर
IIM INDORE

भारतीय प्रबंध संस्थान इंदौर

INDIAN INSTITUTE OF MANAGEMENT INDORE

Executive AI Transformation Proposal

Submission Details

Field	Detail
Name	Himanshu Sirohi
Batch / Cohort	Batch 01
Date of Submission	26 September 2026
Programme	Executive Programme in Generative and Agentic AI for Business Applications
Institution	Indian Institute of Management Indore
Milestone	Milestone 14 - Executive AI Transformation Proposal (consolidating Milestones 1-13)

Use Case and Confidentiality Note

This capstone is built on a single use case throughout: the Oracle Cloud Financials implementation lifecycle delivered by the Oracle Cloud ERP practice of a global professional services firm, covering design sessions, fit-gap, legacy code analysis, functional design, configuration and BR100 documentation, integration and report build, SIT, UAT and go-live.

Employer and client names are genericised. The delivery organisation is referred to throughout as Helix Consulting Group (HCG), a placeholder. Roles, artifacts, workflow structure, constraints and delivery volumes reflect real Oracle Cloud Financials implementation practice; all financial figures are illustrative estimates constructed for this submission and are labelled as such.

Milestone 1 - AI Opportunity Brief

1. Business Context

Helix Consulting Group (HCG) is a global professional services firm. This proposal concerns its Oracle Cloud ERP practice - the Finance delivery group implementing Oracle Cloud Financials (GL, AP, AR, Fixed Assets, Cash Management, Source-to-Pay) for large enterprises migrating off legacy Oracle E-Business Suite R12. Names are genericised; the workflow and constraints reflect live delivery practice.

The practice runs about ten implementations a year. A typical engagement lasts 12-18 months with 18-25 consultants and produces roughly 60 RICE objects, 25 BR100 configuration workbooks, 800 test scripts and 600 logged defects. Work is priced fixed-fee or capped, so margin is a direct function of how many hours it takes to produce a correct artifact first time. This workflow is where cost, quality and client trust are simultaneously determined.

2. Current Workflow Summary

- Requirements and fit-gap - Functional Lead runs design sessions with client process owners; produces the Requirements Traceability Matrix (RTM) and fit-gap dispositions.
- Solution design - Functional Consultant reverse-engineers the legacy EBS PL/SQL the requirement replaces, then authors a Functional Specification Document (FSD).
- Configuration - Functional Consultant applies Oracle Cloud setups, recorded in a BR100 workbook.
- Build - Technical team develops integrations (OIC, FBDI, REST) and BI Publisher / OTBI reports.
- Test preparation - Test Lead authors SIT and UAT scripts from the FSDs and BR100s.
- SIT, UAT, cutover - execute, triage defects, retest, obtain business sign-off, go live.

3. Workflow Challenge

Legacy comprehension is the bottleneck. A 2,000-line PL/SQL package takes two to three days to work through. Miss one validation branch and the defect surfaces in UAT, where it is far costlier.

Meaning drifts across handoffs. A requirement passes through a design session, an RTM row, an FSD, a BR100 and a test script - each a re-typing. About 22% of SIT defects trace to design ambiguity rather than coding error.

Documentation drifts from the system. BR100s are transcribed by hand while setups keep changing. By UAT a material share no longer matches the instance, weakening the client's audit position.

Scarce judgement is spent on mechanics. Functional Leads, Architects and Test Leads spend much of their time transcribing and consolidating rather than deciding.

Milestone 1 - AI Opportunity Brief

4. Why Generative AI Is Relevant

Every artifact here is structured text produced from other structured text. An FSD derives from legacy code, a design session and a template; a test script from an FSD; a BR100 from an instance configuration; a triage note from a defect and the artifacts governing it. That is the shape of work generative AI does well - bounded transformation between document formats, with a reference corpus and expert review.

HCG already owns the grounding asset: hundreds of approved FSDs, BR100s and test scripts in house style, plus the legacy codebases. Retrieval over that corpus produces drafts matching the practice's conventions without fine-tuning. The governing principle is narrow: AI drafts, humans decide.

5. Is Agentic AI Required?

Partially - and the distinction matters for cost and risk. Single-turn prompting suffices for isolated tasks such as summarising one procedure. It does not suffice end to end, because producing an FSD requires dependent actions: locate the legacy object, read it, retrieve the template and two comparable approved FSDs, retrieve the mapped RTM rows, draft, validate that every rule cites a source, and write back with traceability. That is planning, tool use and memory - an agentic system.

What is deliberately not agentic: no agent configures an Oracle environment, closes a defect, signs off a test cycle or triggers cutover. Autonomy covers drafting and analysis, never state change in a client system. That boundary is what makes the initiative governable.

6. Type of Business Impact

Efficiency (primary). Hours per artifact fall. Under fixed fees this converts directly to margin and releases senior capacity - the binding constraint on concurrent engagements.

Risk and quality (secondary). Exhaustive legacy coverage and enforced traceability attack the defect class that does most damage: logic missed at design and found in UAT.

Decision support (tertiary). A traceability store built as a by-product gives leadership live visibility of coverage gaps that today surface only when SIT fails.

Revenue growth is not claimed as a primary impact and is not relied on in the business case.

7. Initial Value Hypothesis

Five repetitive tasks - legacy analysis, FSD authoring, BR100 documentation, test script authoring and defect triage - consume about 2,300 hours per engagement, or 23,000 a year. The hypothesis is that AI-assisted drafting with mandatory human review removes 40-50% of that effort without reducing quality: roughly 10,500 hours a year, about Rs 1.9 crore at a blended Rs 1,800 per hour. A secondary hypothesis is a 20-25% reduction in ambiguity-driven defects. Both are validated in Milestone 4 and gated on measured pilot results in Milestone 13.

Milestone 2 - Workflow Baseline & Risk Canvas

1. Current Workflow Steps

#	Step	Owner	Tool	Duration
1	Design sessions, requirement capture	Functional Lead + client	Teams, Confluence, Excel RTM	2-4 wks per area
2	Fit-gap disposition	Functional Lead + Architect	Excel RTM, Oracle docs	3-5 days per area
3	Legacy EBS code reverse-engineering	Functional Consultant + Tech Consultant	SQL Developer, source repo	8 hrs per object
4	FSD authoring and review	Consultant, Functional Lead	Word, Confluence	12 hrs, 5-7 days
5	Oracle Cloud configuration + BR100	Functional Consultant	Oracle Cloud, Excel BR100	10 hrs per workbook
6	Technical build	Technical Consultant	OIC, BI Publisher, FBDI	Varies
7	SIT / UAT test script authoring	Test Lead + Consultant	ALM	30 min per script
8	SIT execution, defect triage, routing	Test Lead + Defect Manager	ALM	45 min per defect
9	UAT and business sign-off	Client business users	ALM	4-6 weeks
10	Cutover, go-live, hypercare	Programme + Release Mgr	Checklist, Oracle Cloud	2-4 weeks

Steps 3-8 are the scope of this proposal. Steps 1-2 and 9-10 remain fully human-led.

2. Manual vs Automated Tasks

Manual today	Partly automated today	Correctly manual - must stay human
Design session capture; legacy code analysis; FSD authoring; BR100 transcription; test script authoring; defect triage; status reporting	FBDI data loads (templates exist, mapping manual); integration deployment (CI/CD exists, specification manual); regression execution	Applying configuration in the instance; defect closure; UAT sign-off; cutover execution - all require change control and segregation of duties

Milestone 2 - Workflow Baseline & Risk Canvas

3. Workflow Bottlenecks

- Legacy comprehension gates everything downstream - design is rate-limited by reading speed, not design skill.
- Sequential handoffs - RTM to FSD to BR100 to test script, each waiting on the last, so elapsed time far exceeds the effort contained.
- Review queues at the Functional Lead - every FSD and BR100 converges on two or three reviewers.
- Test authoring lags design, so SIT entry is delayed by test preparation rather than by build.
- Defect triage congestion - peak SIT produces 40-60 defects a day against a triage capacity of ~25.
- Configuration-to-document drift - the BR100 is written once while the instance changes.

Milestone 2 - Workflow Baseline & Risk Canvas

4. Baseline Workflow Metrics

Metric	Baseline	Basis
Legacy analysis effort	8 hrs per RICE object	60 objects = 480 hrs
FSD authoring effort / elapsed	12 hrs; 5-7 days	60 FSDs = 720 hrs per engagement
BR100 authoring effort	10 hrs per workbook	25 workbooks = 250 hrs
Test script authoring	30 min per script	800 scripts = 400 hrs
Defect triage effort	45 min per defect	600 defects = 450 hrs
Total, five in-scope tasks	2,300 hrs per engagement	23,000 hrs a year across 10 engagements
Defects traced to design ambiguity	~22% of SIT defects	~130 defects per engagement
Requirements with no test script at SIT entry	~12%	Gap found during SIT
BR100 entries drifted from the instance by UAT	~15%	Sampled at cutover rehearsal
Defects misrouted at least once	~18%	Rework in the triage step
Engagements slipping UAT entry by a week+	1 in 3	Programme status history

Milestone 2 - Workflow Baseline & Risk Canvas

5. Expected Generative AI Outputs

- Structured legacy logic extract** - inputs, validations, derivations, outputs and exception paths, each rule carrying a package and line reference and a VERIFIED or INFERRED label.
- FSD first draft** - scope, assumptions, business rules, field-level mapping and error handling, grounded in the extract, the RTM and two approved FSDs.
- BR100 configuration draft** - generated from a read-only extract of the live instance, so the workbook reflects the environment rather than recollection.
- SIT and UAT test scripts** - positive, negative and boundary cases tagged to a requirement ID.
- Defect triage note** - classification, a cited root-cause hypothesis and a proposed owner queue.
- Daily SIT status summary** - defect movement, ageing, coverage and blockers, for the Test Lead to issue.

6. AI Limitations and Risks

Limitation / risk	Why it matters here	Initial mitigation
Hallucinated business rules	A fabricated validation becomes a built integration and a signed deliverable.	Source citation mandatory; VERIFIED / INFERRED labels; uncited rules fail validation.
Incomplete legacy coverage	A missed branch is invisible - the draft looks complete.	Deterministic coverage check against a parsed object list.
Client data confidentiality	Legacy code and financial configuration are client IP.	Private endpoint, contractual zero retention, per-engagement partitioning.
Cross-client contamination	Grounding Client B's FSD on Client A's corpus would breach contract.	Hard separation of retrieval indexes by engagement.
Oracle quarterly update drift	Oracle Cloud changes quarterly; a stale corpus gives confidently wrong answers.	Knowledge base versioned to the release; drift tracked via override rate.
Prompt injection	Legacy comments and client documents are untrusted input.	Source content passed as data, never as instruction.
Over-reliance and skill erosion	Juniors may lose the ability to read legacy code unaided.	Review treated as a trained competency; juniors verify extracts against source.
Non-determinism	The same input can yield different drafts, complicating audit.	Low temperature; versioned prompts; every output version logged.

Milestone 3 - Prompt & Output Evaluation Toolkit

Problem Understanding and the Need for Advanced Prompting

These five prompts support the design-to-test segment: reading legacy EBS logic, disposing of fit-gaps, documenting configuration, authoring test scripts and triaging defects. Basic prompting fails here for four reasons.

- Volume and nesting - a legacy AP package runs to 2,000+ lines of nested conditionals. An open-ended summary drops branches, and the omission is invisible: the draft reads as complete.
- Asymmetric cost of error - a missed validation becomes a built integration, a signed FSD and a UAT defect. Prompts must make incompleteness visible rather than read well.
- House-style conformance - output must match FSD and BR100 conventions closely enough to review quickly, which needs worked examples, not instructions about tone.
- Machine consumption - outputs feed other agents and a traceability store, so they must be schema-valid data, not prose.

Hence decomposition for coverage, few-shot chain-of-thought for reasoning and style, tree-of-thoughts where a real design choice exists, ReAct where a claim must be verified before assertion, and self-consistency where a wrong answer is expensive.

Milestone 3 - Prompt & Output Evaluation Toolkit

Prompt 1 - Summarisation: Legacy PL/SQL Logic Extract

1. Purpose: Convert a legacy EBS package into an exhaustive, verifiable statement of its business logic for reuse as the business-rules section of an FSD. Workflow step 3.

2. Prompt Text:

"You are a senior Oracle Financials consultant analysing legacy EBS R12 code for migration to Oracle Cloud. Two approved extracts are provided as worked examples; follow their structure exactly. Do not summarise the package as a whole. Work through it in this order, stating your reasoning before each conclusion: (a) every procedure and function with line ranges; (b) every input parameter and its source column; (c) every validation, with the exact condition and the action on failure; (d) every derivation, stating source, rule and target; (e) every exception handler; (f) every INSERT, UPDATE or DELETE and the table affected. Label each item VERIFIED with package and line number where the code states it directly, INFERRED where you are reading intent, or UNRESOLVED where a branch is unclear. Do not guess, and do not omit an item because it looks trivial. Return JSON conforming to the logic_extract schema."

3. Input: Package source; schema definitions for every table touched; two approved extracts as few-shot examples; the logic_extract schema.

4. Output: JSON array - item_type, description, condition, action, source_reference, status (VERIFIED / INFERRED / UNRESOLVED).

5. Validation: Schema-valid; every procedure in the source appears in the extract (deterministic check against a parsed object list); every VERIFIED item has a resolvable line reference.

6. Human Review: Consultant checks all UNRESOLVED and INFERRED items against source plus a 20% VERIFIED sample before the extract may feed the FSD agent. Techniques: decomposition, few-shot CoT.

Milestone 3 - Prompt & Output Evaluation Toolkit

Prompt 2 - Classification: SIT / UAT Defect Triage

1. Purpose: Classify each logged defect by cause and route it to the correct owner queue, replacing a 45-minute manual triage. Workflow step 8.

2. Prompt Text:

"You are triaging a defect raised during SIT of an Oracle Cloud Financials implementation. Eight previously triaged defects are provided as worked examples, each showing the reasoning behind its classification. Classify into exactly one of CONFIGURATION, DATA, CODE, TEST_SCRIPT, EXPECTED_BEHAVIOUR or INSUFFICIENT_INFORMATION. Before classifying, state what the tester expected, what happened, which FSD or BR100 section governs this behaviour, and what it says. If no governing artifact can be retrieved you must return INSUFFICIENT_INFORMATION - do not infer intended behaviour from the description alone. Then propose an owner queue and severity, citing the artifact section that supports your classification. Return JSON conforming to the defect_triage schema."

3. Input: Defect record; the governing FSD or BR100 section reached through the traceability store; eight labelled examples; the schema.

4. Output: JSON - classification, confidence, reasoning, governing_artifact_reference, proposed_owner_queue, proposed_severity, escalate_flag.

5. Validation: One of six permitted values; an artifact cited for every class except INSUFFICIENT_INFORMATION; accuracy measured weekly against the Defect Manager's decision.

6. Human Review: Three independent classifications, majority taken. On disagreement or low confidence the defect goes to the Defect Manager untriaged. Techniques: few-shot CoT, self-consistency.

Milestone 3 - Prompt & Output Evaluation Toolkit

Prompt 3 - Decision Support: Fit-Gap Disposition

- 1. Purpose:** Recommend how a requirement should be met in Oracle Cloud - standard, configuration, extension or custom build - with trade-offs explicit for the Solution Architect. Workflow step 2.
- 2. Prompt Text:**

"You are advising an Oracle Cloud Financials Solution Architect on a fit-gap disposition. Consider four approaches in parallel before recommending any: (1) Oracle standard functionality; (2) configuration only - setups, flexfields, approval rules, value sets; (3) a supported extension using OIC, REST, BI Publisher or AI Agent Studio; (4) a custom build outside the Oracle stack. For each, reason step by step and state how the requirement would be met, the effort, the impact on the quarterly update path, the impact on the client's audit and SOX position, and what would make it fail. Cite the Oracle documentation or precedent you rely on, and mark any uncitable point as an assumption requiring confirmation. Then recommend one, stating explicitly what is given up by not choosing the next-best. If the requirement is ambiguous, do not recommend - return the question the Functional Lead must resolve with the client."
- 3. Input:** RTM requirement row; legacy extract for the process replaced; retrieved Oracle documentation; comparable prior dispositions; the client's customisation policy.
- 4. Output:** Four-option comparison table (approach, effort, upgrade impact, audit impact, failure mode, evidence), then a recommendation with a named trade-off, or a clarification question.
- 5. Validation:** All four options evaluated; every claim about standard behaviour cites documentation; the trade-off given up is stated; ambiguity yields a question, not a guess.
- 6. Human Review:** Advisory only - the Architect decides and records it in the RTM; AI output is supporting analysis, never the decision of record. Technique: tree-of-thoughts with CoT per branch.

Milestone 3 - Prompt & Output Evaluation Toolkit

Prompt 4 - Structured Reporting: BR100 Configuration Documentation

- 1. Purpose:** Generate a BR100 section from a read-only extract of the live instance so documentation reflects the environment rather than recollection, attacking the 15% drift. Step 5.
- 2. Prompt Text:**

"You are documenting Oracle Cloud Financials configuration for a BR100 workbook. You are given a read-only extract of the setup as it currently exists in the instance, the approved BR100 template, and two approved sections as style examples. For every field output: setup task name, navigation path, field name, configured value as read from the extract, and design rationale. The configured value must be reproduced exactly - you must not normalise, correct or complete a value even where it appears to be an error; if a value looks wrong, record it as it is and add a DISCREPANCY flag with your reason. The rationale may be drafted from the linked FSD section; where no FSD rationale exists, write RATIONALE_NOT_DOCUMENTED rather than composing one. Return a table conforming to the br100_row schema."
- 3. Input:** Read-only configuration extract (REST or BI Publisher); approved BR100 template; the linked FSD section; two approved examples.
- 4. Output:** Tabular - setup_task, navigation_path, field_name, configured_value, design_rationale, discrepancy_flag.
- 5. Validation:** Every configured_value matches the source character for character by deterministic string comparison, not model judgement; every row maps to a template setup task.
- 6. Human Review:** Consultant reviews all DISCREPANCY and RATIONALE_NOT_DOCUMENTED rows. Because values are verified deterministically, review targets rationale quality, not transcription. Techniques: structured output control, template grounding, few-shot.

Milestone 3 - Prompt & Output Evaluation Toolkit

Prompt 5 - Risk Mitigation: Pre-Sign-Off FSD Assurance Check

1. Purpose: Act as an adversarial reviewer of a drafted FSD before it reaches the Functional Lead. This prompt exists to catch the failure mode of the other four. Workflow step 4.

2. Prompt Text:

"You are an independent reviewer whose sole objective is to find reasons this FSD should not be signed off. You have the legacy extract, the RTM, and a retrieval tool over the approved corpus. Do not rewrite the document and do not comment on style. Work claim by claim: for each business rule stated in the FSD, (1) state the rule; (2) search the extract and RTM for its source; (3) report what you found; (4) mark it SUPPORTED with the reference, UNSUPPORTED if no source exists, or CONTRADICTED if the source differs, quoting both. Then run three completeness checks: every VERIFIED extract item is addressed somewhere in the FSD; every mapped RTM requirement is covered by a stated rule; every UNRESOLVED item is resolved or carried as an open question. Report only findings. If you find nothing, say so rather than manufacturing an issue."

3. Input: The drafted FSD; the reviewed logic extract; RTM rows; retrieval access to the approved corpus and Oracle documentation.

4. Output: Findings table - claim, status, evidence, severity - plus three completeness results and a READY_FOR_REVIEW or RETURN_TO_DRAFT verdict.

5. Validation: Every rule assessed, count reconciled against the document; every finding cites evidence; the verdict follows deterministically - any UNSUPPORTED or CONTRADICTED item forces RETURN_TO_DRAFT.

6. Human Review: RETURN_TO_DRAFT returns the FSD to the drafting agent and never reaches the Functional Lead's queue; READY_FOR_REVIEW is a precondition for human review, not a substitute. Techniques: ReAct, self-consistency on the verdict.

Milestone 3 - Prompt & Output Evaluation Toolkit

Advanced Techniques Applied, and Why

Technique	Where	Why this and not something simpler
Decomposition	Prompt 1	An open-ended summary drops branches invisibly. A forced six-part enumeration makes the output auditable against a parsed object list, so incompleteness becomes detectable.
Few-shot chain-of-thought	Prompts 1, 2, 4	Worked examples fix reasoning pattern and house style, so reviewers check a document shaped like ones they already trust - which is what makes review fast enough to be worth doing.
Tree-of-thoughts	Prompt 3	A fit-gap has four genuinely competing answers. Evaluating all four before recommending surfaces the trade-off; a linear chain commits to the first plausible option and rationalises it.
ReAct	Prompt 5	Verification needs reasoning interleaved with retrieval - the model must look each claim up before asserting support. A single pass would assert it from memory.
Self-consistency	Prompts 2, 5	Used where a wrong answer is expensive and the output is a discrete verdict that can be voted on: three samples, majority vote, disagreement escalates.
Structured output control	All five	Every output is schema-constrained so a downstream agent or deterministic check can verify it without a human reading it first.

Milestone 3 - Prompt & Output Evaluation Toolkit

Workflow Design, Guardrails and Output Control

The pipeline - five stages of one sequence, not independent tools:

Legacy source + RTM -> [P1] -> logic extract -> GATE 1 (consultant verifies UNRESOLVED and INFERRED) -> [P3] -> GATE 2 (Architect decides) -> FSD draft -> [P5] -> pass, or auto-return to draft -> GATE 3 (Functional Lead signs off) -> [P4] BR100 and test scripts -> GATE 4 -> SIT -> [P2] triage -> GATE 5 (Defect Manager confirms).

P5 sits deliberately between generation and human review - the only stage whose purpose is to reject the output of the stage before it. It fails closed: RETURN_TO_DRAFT cannot be overridden by the drafting agent, only by a human.

Guardrails.

- Citation obligation - uncited rules fail validation before a human sees them.
- Three-state labelling - VERIFIED, INFERRED, UNRESOLVED give the model a permitted way to say it does not know, which is what makes the other two labels informative.
- No silent correction - P4 forbids normalising a value that looks wrong; discrepancies are flagged.
- Deterministic checks outside the model - coverage against a parsed object list, string comparison of configured values, requirement-to-test reconciliation. Where a check can be code, it is code.
- Refusal paths - INSUFFICIENT_INFORMATION, UNRESOLVED, RATIONALE_NOT_DOCUMENTED and P3's clarification question each let the model decline. Most hallucination here is a model with no acceptable way to say the input was inadequate.
- Input sanitisation - source content is passed as data, never concatenated into instructions.
- Scope limits - no prompt can write to an Oracle environment, close a defect or approve an artifact.

Milestone 3 - Prompt & Output Evaluation Toolkit

Evaluation, Accuracy and Trade-offs

Compared on a held-out set of 40 legacy objects and 200 previously triaged defects against consultant-verified ground truth. Pilot-stage estimates, re-measured in Phase 1.

Approach	Accuracy	Consistency	Cost	Latency	Verdict
Zero-shot single prompt	Low	Low	1x	Fast	Rejected - failure is invisible.
Few-shot (examples only)	Medium	Medium	1.2x	Fast	Better style, still drops branches.
Decomposition + few-shot CoT	High	Medium-high	2.5x	Moderate	Selected for extraction and drafting.
Add self-consistency (n=3)	Highest	High	7x	Slow	Selected for triage and verdicts only.
ReAct with retrieval	High on verification	High	4x	Slow	Selected for the assurance check.
Tree-of-thoughts	High on option quality	Medium	5x	Slow	Selected for fit-gap only.

Best performer. Decomposition plus few-shot chain-of-thought with schema-constrained output - the largest accuracy gain per unit cost. Self-consistency and ReAct are reserved for the two points where a wrong answer is most expensive.

Trade-offs accepted. Self-consistency roughly triples cost and latency on the stages it covers, accepted because a misrouted defect costs a consultant hour and a wrongly signed FSD far more. Applying it everywhere was rejected.

Declined. Fine-tuning on the FSD corpus - retrieval gives comparable style conformance, keeps client data out of model weights, and makes deletion on contract exit complete.

Complexity cost. A five-stage pipeline with deterministic validators is harder to operate than a chat interface. That cost is real and is carried in the Milestone 11 run cost, not hidden.

Milestone 4 - Efficiency & Financial Impact Model

Five Repetitive Tasks, Time Saved and Annual Cost Reduction

Basis: 10 Oracle Cloud Financials implementations a year. Per engagement: 60 RICE objects, 25 BR100 workbooks, 800 test scripts, 600 SIT/UAT defects. Blended loaded cost of Rs 1,800 per hour across the Consultant, Senior Consultant and Manager grades that perform this work.

#	Repetitive task	Baseline hours/yr	% time saved	Annual hours saved	Cost/hour	Annual cost reduction
1	Legacy EBS PL/SQL reverse-engineering (60 objects x 8 hrs x 10)	4,800	50%	2,400	Rs 1,800	Rs 43.20 lakh
2	FSD authoring (60 FSDs x 12 hrs x 10)	7,200	40%	2,880	Rs 1,800	Rs 51.84 lakh
3	BR100 configuration documentation (25 x 10 hrs x 10)	2,500	60%	1,500	Rs 1,800	Rs 27.00 lakh
4	SIT/UAT test script authoring (800 x 0.5 hrs x 10)	4,000	55%	2,200	Rs 1,800	Rs 39.60 lakh
5	SIT/UAT defect triage and routing (600 x 0.75 hrs x 10)	4,500	35%	1,575	Rs 1,800	Rs 28.35 lakh
	TOTAL	23,000	46% avg	10,555		Rs 189.99 lakh (Rs 1.90 crore)

Why the percentages differ. BR100 documentation saves most (60%) because the configured values are read directly from the instance and verified by string comparison - the manual step being replaced is transcription, not judgement. Defect triage saves least (35%) because a human still confirms every routing, and roughly a fifth of defects return INSUFFICIENT_INFORMATION and are triaged manually anyway. FSD authoring is held at 40% rather than a higher figure because review effort rises even as drafting effort falls.

Milestone 4 - Efficiency & Financial Impact Model

Estimated Implementation Cost (Year 1)

Cost category	Amount	Basis
Foundation model consumption (private endpoint)	Rs 18 lakh	Managed model service behind a private network endpoint; includes the governance premium over a public API.
Orchestration and agent platform build	Rs 45 lakh	2 FTE x 6 months - agent framework, prompt library, traceability store, evaluation harness.
Enterprise integration	Rs 15 lakh	Jira, Confluence, ALM test management, Git, read-only Oracle Cloud REST and BI Publisher access.
Knowledge base and retrieval build	Rs 12 lakh	Indexing the approved FSD, BR100 and test corpus; Oracle documentation; per-engagement partitioning.
Security, governance and legal review	Rs 10 lakh	Architecture review, client contractual amendments for AI use, DPIA, penetration test.
Training and change management	Rs 8 lakh	Review-skills training for 120 consultants; playbooks; adoption support.
TOTAL YEAR 1	Rs 108 lakh	Rs 1.08 crore, of which Rs 90 lakh is one-time build
Recurring run cost (Year 2 and Year 3)	Rs 42 lakh/yr	Model consumption Rs 22 lakh, platform run 0.5 FTE Rs 15 lakh, tooling Rs 5 lakh.

Milestone 4 - Efficiency & Financial Impact Model

Estimated 3-Year ROI

	Year 1	Year 2	Year 3	3-year total
Realisation assumption	50% (pilot H1, partial scale H2)	100%	110% (broader adoption)	
Gross benefit	Rs 95 lakh	Rs 190 lakh	Rs 209 lakh	Rs 494 lakh
Cost	Rs 108 lakh	Rs 42 lakh	Rs 42 lakh	Rs 192 lakh
Net cash flow	(Rs 13 lakh)	Rs 148 lakh	Rs 167 lakh	Rs 302 lakh
Cumulative	(Rs 13 lakh)	Rs 135 lakh	Rs 302 lakh	

3-Year ROI: $(494 - 192) / 192 = 157\%$.

Payback period: Approximately 13 months - Year 1 closes marginally negative and the cumulative position turns positive early in Year 2.

What is deliberately excluded. No benefit is claimed for reduced defect leakage, avoided go-live slippage or improved win rates, although all three are expected. Quality benefits are real but harder to attribute cleanly, and a business case that rests only on hours saved is one the practice can be held to. Including them would roughly double the headline ROI and halve its credibility.

Sensitivity. At 30% average time saved rather than 46%, the three-year ROI falls to roughly 68% and payback extends to around 22 months - still positive. The initiative does not require the optimistic case to be worth doing, which is the test that matters for approval.


Governance premium. Running the model behind a private endpoint with a contractual no-retention clause costs more than calling a public API. That premium is named as a line item rather than buried, because it is the direct price of the client-confidentiality position in Milestone 12 - it is a requirement, not an inefficiency.

Milestone 5 - Future-State Workflow Redesign

1. Current-State Workflow Diagram

Ten steps across five roles, five of them bottlenecks. The shape that matters is the single-file chain from step 3 to step 8: nothing downstream can start early, so every delay is additive rather than absorbed.

Current-State Workflow - Oracle Cloud Financials Implementation (manual design-to-test chain)




Milestone 5 - Future-State Workflow Redesign

2. AI-Enabled Future-State Workflow

The gates are unchanged - what changes is what arrives at each one. The assurance check sits between the drafting agent and the human and fails closed: an uncited or untraceable draft returns to the agent and never reaches the reviewer's queue.

AI-Enabled Future-State Workflow - agents draft, named humans decide at every gate



Milestone 5 - Future-State Workflow Redesign

3. New Capabilities Enabled by AI

Exhaustive legacy comprehension at speed. Every procedure, branch and exception handler in a legacy package is enumerated and labelled, rather than as much as a consultant could read in the time available. This is a capability the practice did not previously have at any price - the constraint was never skill, it was reading hours.

Traceability as a by-product rather than a project. Requirement to design to configuration to test case to defect is captured automatically because each agent writes its lineage as it works. Traceability today is a manual reconciliation exercise done late, badly, and only when audit requires it.

Documentation that reflects the live system. BR100 content is generated from a read-only extract of the actual instance and verified by string comparison, so configuration drift becomes detectable rather than discovered at cutover rehearsal.

Parallel rather than sequential test preparation. Because scripts are generated from the FSD the moment it is signed off, test authoring stops being a phase that follows design and becomes an immediate consequence of it.

Adversarial self-review before human review. An automated reviewer looks for unsupported claims before the Functional Lead does. Reviewer time shifts from finding errors to judging design.

Coverage and design-debt visibility in real time. Delivery leadership can see uncovered requirements, unresolved legacy branches and open discrepancies during design, instead of discovering them when SIT fails.

Milestone 5 - Future-State Workflow Redesign

4. Business Impact

Margin and capacity (the primary effect). Roughly 1,055 consultant hours are released per engagement - about Rs 19 lakh at blended cost. In a fixed-fee model this is margin. Just as importantly, the hours released are disproportionately senior hours, and senior capacity is the practice's binding constraint on how many engagements it can run concurrently.

Client experience. The client-facing improvement is not speed for its own sake - it is that fewer things are discovered late. A design ambiguity caught by an assurance check during FSD drafting is a five-minute clarification; the same ambiguity found in UAT is a defect, a re-test, a change request and a difficult conversation. Reducing ambiguity-driven defects by 20-25% removes a recurring source of client frustration at the worst possible point in a programme.

Revenue and competitive positioning. Oracle now embeds agentic capability directly in Fusion Cloud and bundles it into the subscription, so an implementation partner can no longer differentiate on configuring AI features the client already owns. Differentiation moves to how fast and how accurately the partner can deliver. A practice that can evidence a traceable, AI-assisted delivery method with measured quality outcomes has a defensible position in a competitive bid; one that cannot is competing on rate card alone.

Risk position. Enforced traceability and instance-derived configuration documentation strengthen the client's own audit and SOX evidence base - a benefit to the client's control environment, not just to delivery efficiency.

What does not change. Deal size, pricing model and team composition are unaffected in the three-year horizon. No headcount reduction is proposed; released capacity is redeployed to additional engagements, which is both the commercially better outcome and the one that makes adoption possible.

Milestone 6 - Agent Architecture Blueprint

1. Agent Objective and Business Context

A multi-agent Delivery Artifact Intelligence platform inside the Oracle Cloud Financials implementation lifecycle. Its objective is stated narrowly on purpose:

"Produce complete, source-cited first drafts of the design, configuration and test artifacts, with full traceability from requirement to test case, such that a named human expert can verify and approve each artifact in materially less time than authoring it - and such that no artifact reaches an approval gate with an uncited claim in it."

Four specialised agents share one memory rather than one general agent: the tasks need different tools, corpora and autonomy levels, and a single agent would need the union of all four permission sets - exactly what the governance model forbids.

Agent	Scope	Primary output	Human approver
Legacy Analysis Agent	Legacy EBS PL/SQL, forms, reports	Logic extract with VERIFIED / INFERRED / UNRESOLVED labels	Functional Consultant
Design Authoring Agent	Extract + RTM + templates -> FSD; instance extract -> BR100	FSD and BR100 section drafts	Functional Lead
Test Authoring Agent	FSD and BR100 -> SIT and UAT scripts	Test scripts tagged to requirement IDs	Test Lead
Delivery Support Agent	Defects, status, cutover checklist	Triage notes, status, readiness checklist	Defect Manager / Programme Manager

Milestone 6 - Agent Architecture Blueprint

2. Agent Planning Workflow

Worked through for the Design Authoring Agent producing an FSD, the most complex path:

- Step 1 - Receive trigger. A reviewed logic extract is marked complete at Gate 1, raising an event.
- Step 2 - Assemble context. Retrieve the mapped RTM rows, the approved FSD template, two comparable approved FSDs, and the Oracle Cloud documentation for the target module.
- Step 3 - Plan. Decompose the FSD into template sections and map source material to each. Sections with no source are listed as open questions rather than drafted.
- Step 4 - Draft section by section, carrying the citation for every asserted rule to the output.
- Step 5 - Self-check: every VERIFIED item addressed, every mapped RTM row covered, every rule cited.
- Step 6 - Assurance check. The independent reviewer prompt runs; RETURN_TO_DRAFT loops back to step 4 with findings, up to two automatic iterations.
- Step 7 - Escalate or submit. After two failed iterations, escalate to the Solution Architect rather than looping further. On pass, the draft enters the Functional Lead's queue.
- Step 8 - Write lineage. Record requirement ID, extract version, prompt and model version, retrieval set and output version - before human review, so a rejected draft is as auditable as an approved one.
- Step 9 - Handle the decision. On approval, publish and trigger the Test Authoring Agent. On rejection, capture the reviewer's reason as a labelled signal for prompt improvement.

Milestone 6 - Agent Architecture Blueprint

3. Tool Ecosystem

Tool / system	Access	Purpose
Legacy source repository (EBS PL/SQL, forms)	Read-only	Source for the Legacy Analysis Agent.
Confluence / document store	Read; write only after approval	Templates, approved corpus, published FSDs and BR100s.
Jira	Read / write with human confirmation	Requirement and task tracking, RTM linkage.
ALM test management tool	Read / write with human confirmation	Test script creation, defect records, triage notes.
Oracle Cloud instance (REST, BI Publisher)	Read-only, non-production	Configuration extracts for BR100 generation and defect context.
Traceability store (purpose-built)	Read / write	Shared lineage graph - the system's working memory.
Vector store over the approved corpus	Read, partitioned by engagement	Retrieval grounding for house style and precedent.
Oracle documentation and release notes	Read	Grounds standard-functionality claims; versioned to the release.
Git and CI/CD	Read-only status	Build status for the Delivery Support Agent.
Identity provider (SSO, RBAC)	Authentication	Scoped per-agent service accounts.

No tool here grants write access to an Oracle Cloud environment, defect closure, or artifact approval.

Milestone 6 - Agent Architecture Blueprint

4. Memory Design

Short-term (task memory). One artifact's working context - retrieved source, draft in progress, assurance findings, iteration count. Discarded on completion. Bounded so a long package cannot silently push the template out of context.

Medium-term (engagement memory). Everything known about one engagement: RTM, logic extracts, approved artifacts, open questions, design decisions, defect history. Partitioned per engagement at the storage and retrieval layer - a hard boundary, not a filter.

Long-term (practice memory). Cross-engagement knowledge HCG owns outright: template library, defect taxonomy, Oracle documentation, anonymised patterns. Client-identifying content is excluded by construction, not by policy.

The traceability store. The structural memory that makes the system governable. A directed graph: requirement -> fit-gap decision -> extract item -> FSD section -> configuration -> test script -> defect. Every node records its creator, source, prompt and model version, and approver.

The graph turns three otherwise unanswerable questions into queries: which requirements lack test coverage; which legacy logic no design document addressed; and, for any line in a signed FSD, what the AI drafted and who changed it.

Milestone 6 - Agent Architecture Blueprint

5. Autonomy Level and Human Oversight

Activity	Autonomy level	Human involvement
Legacy code analysis	L3 - supervised autonomy	Runs unprompted; consultant verifies all UNRESOLVED and INFERRED items plus a 20% sample.
FSD and BR100 drafting	L2 - assisted	Drafts only. Functional Lead signs off. No publication without approval.
Assurance check	L3 - supervised autonomy	Auto-returns failing drafts; humans review only what passes.
Test script generation	L3 - supervised autonomy	Test Lead approves the set; sampling-based review permitted from Phase 3.
Defect triage	L2 in Phase 1, L3 from Phase 2	Defect Manager confirms every routing, then a sample once accuracy is evidenced.
Status reporting and cutover checklist	L2 - assisted	Drafted for a human to issue under their own name.
Configuration changes in any environment	L0 - not permitted	Human-applied. No agent holds write access to an Oracle Cloud instance.
Defect closure, test sign-off, cutover execution	L0 - not permitted	Human decision, always. These are the boundaries the design does not move.

The boundary is state change. An agent may read and draft anything it is entitled to, but may not alter a client system, close a record or approve an artifact. That single rule makes oversight tractable: a reviewer checks a document, never audits an action already taken.

Milestone 6 - Agent Architecture Blueprint

6. Estimated Execution Cost

Agent task	Input tokens	Output tokens	Cost per task	Volume/engagement	Cost/engagement
Legacy logic extract (per object)	~60,000	~10,000	Rs 40	60	Rs 2,400
FSD draft (per document)	~45,000	~8,000	Rs 30	60	Rs 1,800
Assurance check (per FSD)	~50,000	~4,000	Rs 28	60	Rs 1,680
BR100 section set (per workbook)	~35,000	~7,000	Rs 25	25	Rs 625
Test script (per script)	~6,000	~1,200	Rs 4	800	Rs 3,200
Defect triage (3x self-consistency)	~8,000	~900	Rs 3	600	Rs 1,800
Sub-total					Rs 11,505
With 3x allowance for retries and rejected drafts					~Rs 35,000

Assumptions. Blended Rs 0.30 per 1,000 input and Rs 1.50 per 1,000 output tokens on a managed private-endpoint deployment; orchestration and logging overhead loaded at ~20%.

The figure that matters. About Rs 35,000 of inference per engagement against roughly Rs 19 lakh of consultant effort released - 0.2% of the value. The real cost is the platform build and run team in Milestone 11, not tokens.

Implication. Because tokens are negligible against consultant cost, it is rational to spend them on quality - self-consistency, an independent assurance pass, automatic redrafting. The constraint is reviewer patience and latency, not budget.

Milestone 7 - Human-in-the-Loop & Responsibility Matrix

1. Automated vs Human-Controlled Workflow Steps

Workflow step	AI-handled	Human-controlled	Why the split falls here
Design sessions and requirement capture	-	Fully human	Client relationship and elicitation judgement.
Fit-gap disposition	Option analysis and evidence	The decision	AI compares four options; the Solution Architect chooses and owns the consequence.
Legacy code analysis	Full extraction and labelling	Verification of flagged items	Extraction is mechanical; judging an ambiguous branch is not.
FSD authoring	First draft with citations	Review, edit, sign-off	The document is a client deliverable; a named human must stand behind it.
Oracle Cloud configuration	Documentation of what was set	Applying every setup	Change control and segregation of duties in a financial system.
BR100 documentation	Generation from instance extract	Discrepancy review and confirmation	Values verified deterministically; rationale needs human judgement.
Technical build	-	Fully human	Out of scope.
Test script authoring	Generation and requirement tagging	Approval of the set	Coverage adequacy is a risk judgement, not a completeness check.
Test execution	-	Fully human	Execution against a real instance stays with QA.
Defect triage	Classification and routing proposal	Confirmation of routing	Misrouting is cheap to correct and expensive to automate wrongly.
Defect closure and UAT sign-off	-	Fully human	Contractual acceptance; never delegable to a system.
Cutover and go-live	Checklist and status preparation	Every decision and trigger	Irreversible action in a production financial system.

Milestone 7 - Human-in-the-Loop & Responsibility Matrix

2. Human Approval Checkpoints

Gate	What is approved	Approver	Consequence of approval
Gate 1	Legacy logic extract	Functional Consultant	Extract becomes the trusted source for FSD drafting. Nothing downstream may cite unverified content.
Gate 2	FSD	Functional Lead (client process owner counter-signs)	FSD is published, becomes the contractual design of record, and triggers test authoring.
Gate 3	Configuration and BR100 section	Functional Consultant, reviewed by Functional Lead	Configuration is applied by the human; the BR100 becomes audit evidence.
Gate 4	Test script set	Test Lead	Scripts enter the execution cycle and define what 'tested' means for that requirement.
Gate 5	Defect triage and routing	Defect Manager	Defect is assigned and enters a team's queue with an agreed classification.
Gate 6	Cutover readiness checklist	Programme Manager with Solution Architect	Go / no-go for a production event. AI prepares evidence only.

No agent output crosses a gate without a named individual recording a decision - against a person, never a role mailbox, so accountability always attaches to someone who can explain it.

Milestone 7 - Human-in-the-Loop & Responsibility Matrix

3. Escalation Rules

Uncertain outputs.

- Self-consistency disagreement - where three samples do not agree on a classification or verdict, the item is routed to a human untriaged rather than resolved by majority alone.
- Confidence below threshold - the artifact is flagged LOW_CONFIDENCE and placed at the front of the human queue with the reason stated.
- UNRESOLVED legacy items - any branch the agent could not interpret is carried forward as an explicit open question on the FSD and blocks Gate 2 until the Functional Lead disposes it.
- Missing source - INSUFFICIENT_INFORMATION and RATIONALE_NOT_DOCUMENTED outputs are routed to a human rather than completed by inference. The system is designed so that saying nothing is always available as a correct answer.

Complex cases.

- Two consecutive assurance failures on the same artifact - stop redrafting, escalate to the Solution Architect with all findings. Looping is capped deliberately; a third attempt rarely differs materially from the second.
- Two consecutive human rejections of the same artifact - escalate to the Delivery Lead and flag the prompt for review. Repeated rejection is treated as a system defect, not a drafting accident.
- Cross-module or architectural impact detected - route to the Solution Architect before drafting continues.
- Contradiction between the legacy extract and the client's stated requirement - escalate to the Functional Lead as a business question. The system must never silently pick one over the other.

Operational failures.

- Tool or API unavailable - three retries with exponential backoff, then queue the task and notify the Environment Manager. The agent never proceeds on partial data or fabricates a substitute.
- Schema validation failure - one regeneration attempt, then escalate to the AI Platform Owner.
- Traceability store write failure - the task is halted, not completed. An artifact without lineage is treated as an invalid artifact.
- Model or endpoint outage - the workflow falls back to fully manual authoring, which remains a supported path throughout all three phases. The practice must never be unable to deliver because the platform is down.

Milestone 7 - Human-in-the-Loop & Responsibility Matrix

4. Override Mechanisms

A reviewer reaching a gate can take any of four routes, and none of them is treated as the exception:

- Accept unchanged - the draft is published as written and the lineage records that nothing was altered.
- Amend before accepting - the reviewer edits directly. Both the original and the amended version are retained, and the distance between them is the platform's primary quality signal.
- Return for redrafting - the reviewer states the defect and the agent redrafts with that feedback in context, capped at two cycles before the item escalates.
- Discard and author manually - no justification required and no friction imposed. If discarding were harder than accepting, the gate would stop functioning as a control at all.

Two structural facts make those routes real. Nothing an agent produces becomes live without a deliberate human action, because write permission to the approved document store, the Oracle instance and a defect's final status sits outside every agent identity. And the wholly manual route stays open permanently - adoption is measured and encouraged, never mandated, since a reviewer who cannot say no is not reviewing.

Override rate is monitored as a health metric rather than a compliance one. A rising override rate for a particular artifact type is read as a signal that the prompt, the knowledge base or an Oracle release change needs attention - not as a signal that reviewers need persuading.

Milestone 7 - Human-in-the-Loop & Responsibility Matrix

5. Responsibility Matrix

Area	Responsible (does the work)	Accountable (owns the outcome)	Consulted / Informed
AI platform operation, uptime, model endpoint	AI Platform Engineer	AI Platform Owner	Firm IT Security; Delivery Leads
Prompt library, evaluation harness, agent quality	AI Platform Engineer with practice SMEs	AI Platform Owner	Functional and Test Leads
Knowledge base currency (Oracle releases, templates)	Practice Knowledge Manager	Oracle Practice Lead	AI Platform Owner
Legacy extract verification	Functional Consultant	Functional Lead	Solution Architect
FSD content and design correctness	Functional Consultant (AI-assisted)	Functional Lead	Client process owner; Solution Architect
Configuration applied in the instance	Functional Consultant	Functional Lead	Client IT change control
Test coverage adequacy	Test Analyst	Test Lead	Functional Lead
Defect classification and routing	Delivery Support Agent, confirmed by Defect Manager	Defect Manager	Test Lead; Technical Lead
Client data protection and AI contractual compliance	Engagement Manager	Engagement Partner	Firm Legal; Client DPO
AI risk, governance and incident response	AI Governance Committee	Oracle Practice Lead	Firm Risk; Engagement Partners
Final deliverable accepted by the client	Engagement Manager	Engagement Partner	Client sponsor

Final accountability never rests with the AI system or with the platform team. The Engagement Partner is accountable for every artifact delivered to a client, exactly as before the initiative. The platform team is accountable for the tool, not for the work product - and that distinction is stated explicitly so that it cannot be blurred under delivery pressure.

Milestone 8 - Autonomous Workflow & Failure Design

1. Event-Triggered Autonomous Workflow

Triggering event. A tester logs a defect in the ALM tool during SIT. No human requests triage; the record's creation is the trigger. The clearest autonomous case - high volume, time-critical, and cheaply reversible if wrong.

Agent sequence.

- 1. Detect and parse - webhook fires; the agent reads summary, steps, expected versus actual, module, environment, severity and attachments.
- 2. Resolve traceability - look up the test script, then follow the graph to the governing FSD section, requirement and BR100 entry. If none resolves, that is itself a finding.
- 3. Gather context - read-only calls to the SIT instance for the configuration, plus recent deployment and data-load history.
- 4. Classify - CONFIGURATION, DATA, CODE, TEST_SCRIPT, EXPECTED_BEHAVIOUR or INSUFFICIENT_INFORMATION, generated three times independently with a majority vote.
- 5. Form a cited hypothesis - state the probable cause with a direct reference to the governing artifact section and the observed configuration value.
- 6. Check for duplicates - search open defects for the same root cause and link rather than add.
- 7. Propose routing and severity - owner queue and a severity recommendation with reasoning.
- 8. Write the triage note - post the structured note to the defect record in a draft state, visibly attributed to the agent.
- 9. Notify and hand over - the Defect Manager receives the proposed routing. Gate 5 applies: the agent cannot assign, re-prioritise or close.
- 10. Record the outcome - the Defect Manager's decision is written to the traceability store as that week's accuracy measurement.

Outcome. Triage arrives in minutes with evidence attached. The Defect Manager's job shifts from investigation to confirmation - the source of the deliberately modest 35% saving in Milestone 4.

Milestone 8 - Autonomous Workflow & Failure Design

2. Tool Orchestration

Interaction	System	Mechanism	Permission
Defect event and record read	ALM test management	Webhook + REST	Read; write comments only, never status
Traceability resolution	Traceability store	Internal graph query	Read / write lineage
Governing artifact retrieval	Confluence / vector store	Retrieval API, engagement-partitioned	Read within this engagement only
Configuration lookup	Oracle Cloud Financials SIT instance	REST and BI Publisher	Read-only service account, non-production only
Deployment and data-load history	CI/CD pipeline, FBDI logs	Read-only API	Read
Duplicate detection	ALM defect index	Search API	Read
Notification	Teams / email	Outbound webhook	Send only
Model invocation	Managed model behind private endpoint	Private network call	Scoped per-agent service identity
Audit logging	Central log store	Append-only write	Write only, no delete

Orchestration runs on a stateful workflow engine, so every step is retryable, idempotent and independently logged. Each agent authenticates as its own scoped identity - the Delivery Support Agent cannot reach the document write path. Permissions are enforced at the identity layer, not by prompt instruction, because an instruction is not a control.

Milestone 8 - Autonomous Workflow & Failure Design

3. Failure Simulation Scenarios

#	Failure scenario	System response	Residual risk
1	Defect description too vague to classify	Returns INSUFFICIENT_INFORMATION; auto-comments requesting reproduction steps; flags to Defect Manager after four hours unanswered.	Tester ignores the request; ageing report catches it.
2	Defect maps to no requirement in the traceability store	Escalates to Functional Lead as a coverage gap, not just a triage failure - an untraceable defect is itself a finding about the design.	Genuine gaps may be numerous early in SIT.
3	Confidently wrong root-cause hypothesis	Citation forces a reference to a specific artifact section, so the Defect Manager can check it in seconds. Wrong hypotheses are caught at Gate 5 and logged as accuracy misses.	A plausible, correctly cited but wrong hypothesis can still mislead a busy reviewer.
4	Oracle Cloud instance unreachable	Three retries with backoff, then triage proceeds using artifact evidence only, explicitly marked CONTEXT_INCOMPLETE. Never fabricates a configuration value.	Lower accuracy during environment outages; flagged, not hidden.
5	Prompt injection via a crafted defect description or code comment	Source content is passed as data, never concatenated into instructions; schema-constrained outputs mean an injected instruction cannot produce an out-of-schema action, and the agent lacks permission to take the action an injection would seek.	Novel techniques; covered by quarterly security review.
6	Oracle quarterly update changes standard behaviour	Knowledge base is versioned to the client's release; rising override rate on EXPECTED_BEHAVIOUR classifications triggers a refresh.	A detection lag of days after each quarterly update.
7	Traceability store write failure	Task halts and is queued rather than completing. An artifact without lineage is invalid by definition.	Throughput loss during store outages.
8	Systematic accuracy degradation after a model version change	Regression suite of 200 previously triaged defects runs against every model or prompt change before promotion; a drop beyond tolerance blocks the change.	Slow drift within tolerance; caught by monthly trend review.

Milestone 8 - Autonomous Workflow & Failure Design

4. Escalation Path for Operational Issues

Severity	Trigger condition	Notified	Action and timeframe
S1 - Critical	Suspected client data exposure or cross-engagement contamination	AI Platform Owner, Engagement Partner, Firm Information Security, Client DPO	Immediate suspension of affected agents; containment within 1 hour; client notification per contract.
S2 - High	AI-originated error found in an approved client deliverable	Engagement Manager, Functional Lead, AI Platform Owner	Artifact flagged in the traceability store; impact assessment on all downstream artifacts within 24 hours; agent stage suspended for that artifact type pending root cause.
S3 - Medium	Accuracy below tolerance, or override rate breaching threshold for two weeks	AI Platform Owner, Practice Knowledge Manager	Prompt and knowledge base review within 5 working days; affected stage reverts to L2 assisted until restored.
S4 - Low	Platform unavailability or degraded latency	AI Platform Engineer, Delivery Leads	Fall back to the manual workflow; restore within the platform SLA; no delivery impact beyond the efficiency loss.

Two principles. Any S1 or S2 suspends the agent stage before root cause is known - the default on a suspected fault is to stop. And delivery is never blocked by the platform, because the manual workflow remains fully supported, which is what makes stopping affordable.

Milestone 8 - Autonomous Workflow & Failure Design

5. Monitoring Strategy for System Performance

Metric category	Metric	Target	Frequency / alert
Quality	Draft acceptance rate (no or minor edit)	>= 70% by end of Phase 2	Weekly; alert below 50%
Quality	Edit distance between AI draft and approved version	Declining trend	Weekly trend review
Quality	Triage accuracy vs Defect Manager decision	>= 85%	Weekly; alert below 75%
Quality	Override rate by artifact type	Stable or falling	Weekly; alert on 2-week rise
Integrity	Traceability validation pass rate	100% (blocking gate)	Every artifact; alert on bypass
Integrity	Uncited claims reaching a human reviewer	Zero	Immediate alert
Integrity	Legacy coverage check pass rate	100%	Every extract
Safety	Cross-engagement retrieval attempts	Zero	Immediate S1 alert
Safety	Schema validation failure rate	< 2%	Daily
Operational	Latency per artifact type	FSD draft < 10 min; triage < 3 min	Daily percentiles
Operational	Cost per artifact and per engagement	Within Milestone 6 estimates	Monthly
Operational	Platform availability	>= 99% in business hours	Continuous
Business	Hours saved per engagement vs baseline	Tracks Milestone 4 model	Per engagement close
Business	Defects traced to design ambiguity	20-25% reduction	Per engagement close

Oversight cadence. Weekly automated quality dashboard; monthly AI Governance Committee review of accuracy, incidents and override patterns; per-engagement close-out reconciling realised savings against the business case. The close-out matters most - it is what keeps the Milestone 4 numbers honest rather than an estimate nobody revisits.

Milestone 9 - Strategic Validation & Investment Brief

1. Description of Three AI Initiatives

Initiative 1 - Delivery Artifact Intelligence - Four agents producing source-cited drafts of legacy extracts, FSDs, BR100s, test scripts and triage notes on a shared traceability store, with a named human approving every artifact. Objective: cut consultant hours per engagement and defects caused by design ambiguity. Strategic relevance: it improves margin and quality on the core delivery process, and applies to every engagement rather than a subset.

Initiative 2 - Client-Facing Oracle Agentic Applications Enablement - A service offering helping clients activate, govern and measure the AI agents Oracle now ships inside Fusion Cloud, plus custom agents built in Oracle AI Agent Studio. Objective: new billable revenue and stronger renewal positioning. Strategic relevance: Oracle launched Fusion Agentic Applications in March 2026 bundled into existing subscriptions, so clients own capability they have not deployed - activation, guardrail design and change management are an immediate services gap.

Initiative 3 - Bid and Estimation Intelligence - AI-assisted RFP drafting and effort estimation grounded in the practice's history of completed engagements, actual versus estimated effort and reusable components. Objective: faster, more accurate bids and fewer under-scoped fixed-price deals. Strategic relevance: estimation error is a recurring margin leak, and the data already exists internally.

These are not unrelated. Initiative 1 produces the corpus Initiative 3 needs to estimate credibly, and builds the agent-governance competence Initiative 2 must sell. Sequencing matters as much as scoring.

Milestone 9 - Strategic Validation & Investment Brief

2. Strategic Comparison Framework

Criterion	1. Delivery Artifact Intelligence	2. Oracle Agentic Enablement	3. Bid & Estimation Intelligence
Feasibility	High. Data is firm-owned, sits in internal systems, needs no client permission to pilot. Domain experts in-house.	Medium. Depends on client instances, appetite, data agreements and Oracle release timing - none firm-controlled.	Medium-high. Data exists but is fragmented and inconsistently recorded.
Return on investment	High and measurable. Rs 1.90 crore a year at steady state against Rs 1.08 crore Year 1 cost; 157% three-year ROI.	Potentially high but unproven. Depends on untested demand and work not yet won.	Moderate. Real value in avoided under-scoping, but hard to attribute - visible only as a loss that did not happen.
Implementation risk	Medium-low. Client IP is the material risk, addressed by architecture. Failure is contained - revert to manual authoring.	High. Agents in a client production ERP; failure is visible to client and Oracle, with immediate reputational cost.	Low-medium. Internal only. Main risk is anchoring bids on a poor model.
Strategic impact	High. Improves the core delivery process on every engagement and creates the traceability asset the others depend on.	Medium-high. New revenue line, but on capability Oracle bundles - partners compete on governance, not technology.	Medium. Improves commercial discipline; does not change what the practice delivers.
Time to first value	3-6 months (one-engagement pilot)	9-12 months (needs a client)	6-9 months (data remediation first)
Dependency on others	None material	Client and Oracle release cadence	Internal finance data quality

Milestone 9 - Strategic Validation & Investment Brief

3. Scoring Matrix

Scored 1 (poor) to 5 (strong). Implementation risk is scored inversely, so 5 means low risk. Weights reflect that the practice's constraint today is delivery margin and senior capacity, not revenue opportunity.

Criterion	Weight	1. Delivery Artifact Intelligence	2. Oracle Agentic Enablement	3. Bid & Estimation Intelligence
Feasibility	25%	5	3	4
Return on investment	30%	5	4	3
Implementation risk (inverse)	20%	4	2	4
Strategic impact	25%	5	4	3
Weighted score	100%	4.80	3.35	3.45
Rank		1	3	2

Sensitivity: the ranking is stable. Even halving the ROI weight and doubling strategic impact - which would favour the revenue option - Initiative 1 still leads, because it scores highest on feasibility and strategic impact independently of the financial case.

Milestone 9 - Strategic Validation & Investment Brief

4. Final Prioritisation Decision

Initiative 1, Delivery Artifact Intelligence, is recommended for implementation first.

Reasoning.

- It is the only one the firm can start immediately - no client permission, no new revenue, no data it does not own. Initiative 2 needs a willing client; Initiative 3 needs data remediation first.
- Its benefit compounds across every engagement rather than only the deals that include it.
- Its failure mode is cheap. If the agents underperform, the practice reverts to manual authoring and loses the platform investment. If a client-facing agent misbehaves in a production ERP, the practice loses a client.
- It builds the prerequisites for the other two: the traceability store makes Initiative 3 estimable, and the governance framework is precisely the competence Initiative 2 would sell. Doing it first makes the others cheaper and more credible; the reverse is not true.

What is being given up. Initiative 2 is the revenue play, and deferring it by twelve months cedes early-mover positioning on Oracle agentic enablement while the market is forming. This is a real cost and it is accepted knowingly. The mitigation is that Initiative 2 is deferred rather than declined - it enters the roadmap at Phase 3 (Milestone 13), by which point the practice can sell agent governance from demonstrated experience rather than from a slide.

Decision rule for revisiting. If a strategic client formally requests agentic enablement before Phase 3, Initiative 2 is brought forward as a scoped single-client engagement rather than a practice-wide offering.

Milestone 9 - Strategic Validation & Investment Brief

5. Investment and Change Management Considerations

Investment required. Rs 1.08 crore in Year 1, Rs 42 lakh a year thereafter (Milestones 4 and 11). About Rs 90 lakh of Year 1 is one-time build.

Resource requirements. An AI Platform Owner (0.5 FTE), two Platform Engineers for build reducing to 0.5 FTE for run, a Practice Knowledge Manager (0.3 FTE), and ~15% of a Solution Architect during the pilot. Critically, the SMEs writing prompts and worked examples are senior consultants - the scarcest input, and their time must be protected rather than fitted around billable work.

Stakeholder engagement. The Oracle Practice Lead sponsors. Engagement Partners must agree before any client content enters the platform - a contractual matter, not an internal one. Pilot Delivery Leads must accept a temporary productivity dip. Firm Legal and Information Security are gating approvers before the pilot, not consulted afterwards.

The change management problem, stated honestly. The hardest part of this initiative is not technical. It is that reviewing a draft is a different skill from authoring one, and it is less satisfying. A consultant who has always written FSDs from scratch may find reviewing an AI draft feels like proofreading, and the failure mode is not rejection - it is approving too readily because the document reads well. Rubber-stamping is the single largest risk to the entire quality argument in this proposal.

Mitigations.

- Review is trained and assessed as a distinct competency, with a checklist per artifact type.
- Edit distance and override rate are monitored per reviewer. A reviewer who never edits anything has a conversation with their Delivery Lead, not congratulations on efficiency.
- Junior consultants continue to verify extracts against source code as a deliberate training exercise, protecting the legacy-reading skill the practice still needs.
- Utilisation targets are not raised during the pilot. Asking people to adopt a new method and deliver more simultaneously guarantees neither.
- Released hours are redeployed to additional engagements, not headcount reduction, stated at the outset. An efficiency programme believed to be a redundancy programme is not adopted honestly.

Milestone 10 - Industry Benchmark & Competitive Positioning Note

1. Company 1 Benchmark - Accenture

Why this comparator. Accenture is the closest large-scale analogue to HCG's situation: a global professional services firm whose product is consultant hours, applying generative AI both to what it sells and to how it delivers.

Implementation and scale. Accenture has pursued generative and agentic AI as a core business line rather than an internal experiment, supported by a large acquisition and alliance programme and a substantial reskilling commitment across a workforce of several hundred thousand people. In 2025 it consolidated five previously separate business units into a single reinvention services organisation, explicitly to make it easier to embed AI and data across everything it sells.

Reported outcomes. In the twelve months ending 31 August 2025 Accenture reported that generative and agentic AI revenue tripled year on year to approximately USD 2.7 billion, with related bookings close to USD 5.9 billion, against total annual revenue of roughly USD 70 billion. Separately, Accenture's own research across 2,000 executives found that around three-quarters of organisations said their generative AI and automation investments had met or exceeded expectations, and that organisations with mature AI-led processes achieved materially higher productivity than their peers.

Relevance to HCG. Two things are instructive. First, the scale of the revenue relative to the total - a few percent - shows that even for the most aggressive adopter, AI is a growing line within a services business rather than a replacement for it. Second, the reorganisation is the more telling signal: Accenture concluded that capturing AI value required changing how delivery is organised, not just adding tools to an unchanged process. That is the same conclusion this proposal reaches at a much smaller scale.

Milestone 10 - Industry Benchmark & Competitive Positioning Note

2. Company 2 Benchmark - Infosys

Why this comparator. Infosys is the closer operational analogue: an India-centric delivery model with the same cost structure, the same blended-rate economics and the same regulatory environment HCG operates in, applying AI directly to the software delivery lifecycle rather than only to client-facing products.

Implementation and scale. Infosys delivers AI work through its Topaz platform and has taken an explicitly AI-first position across consulting and delivery. It has deployed GitHub Copilot to more than 30,000 developers, and through its AWS collaboration has applied Topaz with Amazon Q Developer across the SDLC specifically for automated documentation, code generation, debugging, testing and legacy code modernisation - the same task categories this proposal targets. It has also established collaborations with multiple foundation model providers to support software development, legacy modernisation and agent building.

Reported outcomes. Infosys reported AI revenues at 8.2% of revenue in Q1 FY27 on quarterly revenue of approximately USD 5.08 billion, and has cited specific operational results from AI deployment including roughly 50% faster contract validation and an 18% improvement in IT operations efficiency. Management has publicly framed AI productivity impact as a defining factor for FY27.

Relevance to HCG. The most useful detail is the task list. Automated documentation, testing and legacy modernisation are exactly the categories in this proposal, which confirms the opportunity is real and already being industrialised by competitors. The second useful detail is the scale of developer deployment - 30,000 seats indicates that the hard problem at scale is adoption and workflow integration across a large consultant base, not model capability.

Milestone 10 - Industry Benchmark & Competitive Positioning Note

3. Key Insights for the HCG AI Solution

Insight 1 - Tool deployment is not the same as value capture. Accenture restructured its delivery organisation and Infosys deployed AI at seat scale across tens of thousands of developers. Neither treated AI as a licence purchase. This proposal responds by designing the workflow and the approval gates first and the agents second, and by budgeting explicitly for change management rather than assuming adoption follows availability.

Insight 2 - The SDLC documentation and testing layer is where competitors are actually working. Infosys's SDLC focus on documentation, testing and legacy modernisation validates the five tasks selected in Milestone 4. It also means this is now table stakes rather than differentiation - a reason to move now, and a reason not to overstate the competitive advantage in the business case.

Insight 3 - Published outcome figures are heterogeneous and rarely comparable. Both firms report in different units: revenue share, seat counts, and point improvements on specific processes. Almost none of it is directly transferable to a per-engagement hours-saved model. This proposal therefore derives its numbers from HCG's own baseline rather than borrowing a benchmark percentage, and gates scaled investment on measured pilot results in Phase 1.

Insight 4 - AI revenue remains a small share even at the leaders. At single-digit percentages of total revenue for both firms, the realistic near-term prize is improving the economics of existing delivery, not replacing it. That supports prioritising Initiative 1 over the revenue-generating Initiative 2 in Milestone 9.

Insight 5 - Governance is becoming the differentiator. Both comparators emphasise responsible AI and sovereignty and regulatory capability alongside capability itself. For a firm handling client financial system IP, a demonstrable governance model is not overhead - it is increasingly the thing clients are actually buying.

Milestone 10 - Industry Benchmark & Competitive Positioning Note

4. Competitive Positioning Reflection

HCG will not match Accenture's or Infosys's scale of AI deployment, and should not try to. The positioning question is narrower and more answerable: what can a mid-size Oracle Cloud Financials specialist do that a generalist at scale cannot?

Where HCG can be genuinely differentiated.

- Depth over breadth. A general-purpose SDLC assistant does not know what a BR100 is, how an Oracle Cloud approval rule is configured, or why a legacy EBS AP interface derives a distribution the way it does. HCG's corpus of Oracle Financials artifacts is small in absolute terms but dense and specialised, and that is exactly what retrieval grounding rewards.
- Traceability as a client-facing asset. Requirement-to-test lineage with an audit trail of who approved what, and what the AI originally drafted, is directly valuable to a finance client's own SOX and audit position. Most competitors present AI as a speed story; presenting it as a control story is a more credible pitch to a CFO and a harder one to replicate.
- Honest positioning on what stays human. Oracle now ships agentic capability inside Fusion Cloud, bundled into the subscription. Clients are being told a great deal about autonomy. A partner that can articulate precisely where autonomy is appropriate and where it is not - and can show its own governance model as evidence - is offering something clients are short of.

Where HCG will not lead, stated plainly. The firm will not have a proprietary model, a named AI platform brand, or research output. It will be a fast, disciplined adopter of commodity capability applied to a narrow domain it knows better than a generalist does. That is a defensible position, but only if the delivery quality claim is true - which is why the measurement regime in Milestone 8 and the pilot gate in Milestone 13 are load-bearing parts of the strategy, not administrative additions.

Milestone 11 - Enterprise Architecture & Vendor Strategy

1. AI Architecture Overview

Five layers, deployed inside HCG's own cloud tenant. The design principle is that client content never traverses the public internet and never leaves a boundary HCG controls.

Layer 1 - Model. A managed foundation model reached through a private endpoint inside HCG's tenant (Azure OpenAI, AWS Bedrock or Vertex AI - the pattern is deliberately vendor-neutral). A smaller model handles classification and extraction; a larger one drafting and assurance.

Layer 2 - Orchestration. A stateful workflow engine running the four agents: prompt templates, tool invocation, retry and escalation logic, the evaluation harness and the deterministic validators. Every agent authenticates as a distinct least-privilege principal.

Layer 3 - Knowledge and retrieval. A vector store and metadata index over the approved corpus, templates and Oracle documentation, physically partitioned by engagement. Retrieval is scoped by the caller's engagement identity, so cross-engagement retrieval is structurally unavailable.

Layer 4 - Traceability and audit. A graph store holding requirement-to-test lineage, plus an append-only log of every prompt, retrieval set, model version, output version and human decision.

Layer 5 - Integration and presentation. Connectors to Jira, Confluence, ALM, Git and read-only Oracle Cloud endpoints, surfaced inside the tools consultants already use. Adoption depends on this.

Deployment. Containerised on managed Kubernetes in the firm's tenant, infrastructure as code, three environments, with model and prompt versions promoted only after passing the Milestone 8 regression suite.

Milestone 11 - Enterprise Architecture & Vendor Strategy

2. Vendor Ecosystem

Layer	Vendor category	Role	Selection consideration
Cloud platform	Hyperscaler (Azure / AWS / GCP)	Hosts the solution in the firm's own tenant	Whichever the firm standardises on; avoid a second cloud for one workload.
Foundation model	Managed service from the same hyperscaler	Drafting, extraction, classification, assurance	Private endpoint support and a zero-retention clause are non-negotiable.
Vector store	Managed vector database	Retrieval grounding over the artifact corpus	Must support hard tenant isolation per engagement, not just metadata filtering.
Orchestration framework	Open-source agent framework on managed compute	Agent execution and workflow state	Avoid proprietary orchestration locking the differentiated layer to a vendor.
Document and work management	Atlassian (Confluence, Jira) - incumbent	Artifact storage, requirement tracking	Already licensed; integrate, do not replace.
Test management	ALM tool - incumbent	Test scripts, defect records	Already licensed.
Source control and CI/CD	Git platform - incumbent	Legacy source access, build status	Already licensed.
Identity and access	Enterprise IdP	SSO, RBAC, per-agent service identities	Agent identities must be first-class in the existing IAM model.
Observability	Managed logging and monitoring	Metrics, audit log, alerting	Append-only retention satisfying the Milestone 12 audit requirement.
Oracle	Strategic alliance partner	Fusion Cloud, AI Agent Studio, documentation	Relevant to Initiative 2; Initiative 1 depends only on documentation and read-only APIs.

Milestone 11 - Enterprise Architecture & Vendor Strategy

3. Build vs Buy vs Partner Decision

Answered separately per layer. A single blanket answer would be the wrong answer for at least three of the five.

Component	Decision	Justification
Foundation model	PARTNER	Running model infrastructure is context work for a consulting firm, not core work; every engineer-hour there is one not spent on differentiating delivery capability. The partner relationship is what buys the private endpoint and no-retention clause - self-hosting costs more and protects less.
Orchestration, prompts, validators, evaluation harness	BUILD	The differentiated asset. The prompt library, VERIFIED / INFERRED discipline, traceability validators and assurance check encode HCG's Oracle Financials delivery method. A bought product would make the firm's method identical to its competitors' - the opposite of the Milestone 10 argument.
Traceability store and lineage model	BUILD	No off-the-shelf product models requirement-to-BR100-to-test-script lineage for an Oracle Cloud implementation, and it persists after any model or vendor is replaced.
Vector store and retrieval infrastructure	BUY	A commodity managed service. Building one would add operational burden with no differentiation. The differentiation is in what is indexed and how it is partitioned, not in the database.
Document, work and test management, source control, CI/CD, identity	BUY	Already licensed and embedded. Replacing any would be a separate programme and would destroy adoption.
Client-facing agents in Fusion Cloud (Initiative 2)	PARTNER (Oracle)	Oracle AI Agent Studio inside Fusion inherits Fusion's own access controls, approval hierarchies and audit framework. A parallel external agent layer against a client ERP would be strictly worse on governance.

Accepted trade-off. Partnering on the model layer creates vendor dependency. Mitigated by keeping orchestration model-agnostic: prompts and validators are written against an abstraction, and the regression suite makes switching measurable. Dependency is not eliminated, but made survivable.

Milestone 11 - Enterprise Architecture & Vendor Strategy

4. Data and Integration Architecture

Data sources.

- Legacy EBS source code and object definitions - read-only from the client-provided repository or a secured copy inside the engagement partition.
- Requirements and design content - Jira RTM rows, Confluence design session notes and approved artifacts.
- Oracle Cloud configuration - read-only REST and BI Publisher extracts from non-production instances only. No agent reads from, or writes to, a production environment.
- Test and defect data - ALM scripts, execution results and defect records.
- Practice knowledge - templates, anonymised patterns, Oracle documentation and release notes, held in the shared long-term layer with no client-identifying content.

Data flow for a single artifact. Event raised in Jira or ALM -> orchestrator authenticates as the agent's scoped identity and resolves the engagement partition -> retrieval from that partition only -> model invoked over the private endpoint -> output schema-validated -> deterministic checks run -> lineage written to the traceability store -> draft surfaced in the consultant's existing tool -> human decision captured -> on approval only, write-back to the document store.

Integration patterns. Event-driven webhooks for triggers; REST for reads and gated writes; a message queue between agents so that a slow or failed stage does not block the pipeline; idempotency keys on every write so that retries cannot duplicate artifacts or triage notes.

Data lifecycle and exit. Each engagement's data has a defined retention period tied to the client contract. On engagement closure the partition is exported for the client if required and then deleted, including its retrieval indexes. Because no client content is used for fine-tuning, deletion is genuinely complete - there is no residue in model weights. This is a direct consequence of the retrieval-over-fine-tuning decision in Milestone 3, and it is the reason that decision was made.

Milestone 11 - Enterprise Architecture & Vendor Strategy

5. Security Layer

Control area	Implementation
Network isolation	Model accessed only over a private network endpoint within HCG's cloud tenant. No public internet egress for any path that carries client content. Private DNS; deny-by-default egress rules.
Contractual data protection	Zero-data-retention and no-training clauses with the model provider. Network isolation and contractual non-retention are two distinct controls and both are required - neither substitutes for the other.
Tenant and engagement isolation	Separate logical partitions per engagement across storage, retrieval indexes and logs. Retrieval is scoped by engagement identity at query construction, so a cross-engagement query cannot be formed. Any attempted cross-partition access raises an immediate S1 alert.
Identity and access management	SSO with MFA for users; every agent authenticates as a distinct least-privilege principal. No shared credentials. No agent identity holds write access to an Oracle Cloud instance, the approved document store or a defect's final status.
Encryption	TLS 1.2+ in transit; AES-256 at rest with customer-managed keys. Traceability store and audit logs encrypted with separate key material.
Input handling	Legacy code, client documents and defect text are treated as untrusted data and are never concatenated into the instruction portion of a prompt. Outputs are schema-constrained, so an injected instruction cannot produce an out-of-schema action.
Secrets management	Managed key vault; no credentials in code, configuration or prompts; automated rotation.
Audit logging	Append-only, immutable, tamper-evident. Every prompt, retrieval set, model and prompt version, output version and human action recorded with identity and timestamp.
Data residency	Model endpoint and all storage pinned to the region required by the client contract, with India as the default for domestic engagements.
Security assurance	Architecture review and penetration test before the pilot; quarterly agent-specific security review covering prompt injection and permission drift; annual third-party assessment.

Milestone 11 - Enterprise Architecture & Vendor Strategy

6. Estimated Implementation Budget

Category	Year 1	Year 2	Year 3	Notes
Foundation model consumption	Rs 18 L	Rs 22 L	Rs 22 L	Includes the governance premium for private-endpoint deployment.
Platform build (2 FTE x 6 months)	Rs 45 L	-	-	One-time: orchestration, validators.
Platform run (0.5 FTE)	-	Rs 15 L	Rs 15 L	Maintenance, prompt tuning, incidents.
Enterprise integration	Rs 15 L	-	-	Jira, Confluence, ALM, Git, Oracle APIs.
Knowledge base and retrieval build	Rs 12 L	-	-	Corpus indexing and partitioning.
Security, governance, legal	Rs 10 L	-	-	Architecture review, DPIA, penetration test.
Training and change management	Rs 8 L	-	-	Review-skills training, 120 consultants.
Cloud infrastructure and tooling	In above	Rs 5 L	Rs 5 L	Compute, vector store, logging.
TOTAL	Rs 108 L	Rs 42 L	Rs 42 L	Three-year total Rs 192 L

About 83% of Year 1 spend is one-time build; ongoing platform cost is ~Rs 42 lakh a year against Rs 1.90 crore of annual benefit. A 20% build contingency is recommended and deliberately excluded above, so it stays a visible governance decision rather than padding hidden in the estimates.

Milestone 12 - AI Governance & Control Framework

1. AI Risk Identification

#	Risk	Description and impact	Likelihood / Impact
R1	Client data leakage	Client source code, financial configuration or commercial terms exposed outside the contractual boundary. Contractual breach, reputational damage, possible termination.	Low / Critical
R2	Cross-engagement contamination	Client B's artifact grounded on Client A's content. A confidentiality breach that could occur silently and be discovered by a client rather than by HCG.	Low / Critical
R3	Hallucinated design content	A fabricated or wrong business rule reaches a signed FSD, is built, and surfaces as a production defect in a financial system.	Medium / High
R4	Incomplete legacy coverage	A conditional branch is missed. Unlike a hallucination this leaves no trace in the output - the draft simply looks finished.	Medium / High
R5	Rubber-stamp review	Reviewers approve fluent drafts without genuine scrutiny, removing the control the entire governance model depends on.	Medium / High
R6	Prompt injection	Malicious or malformed content in legacy comments or client documents attempts to alter agent behaviour.	Low / Medium
R7	Knowledge drift after Oracle updates	Oracle Cloud changes quarterly; a stale knowledge base produces confidently incorrect standard-behaviour claims.	High / Medium
R8	Bias in triage and estimation	Systematic misrouting by module or under-estimation for certain work types, distorting workload and commercial decisions.	Medium / Medium
R9	Skill erosion	Junior consultants lose the ability to read legacy code unaided, weakening the firm's capability and its ability to review AI output.	Medium / High
R10	Vendor dependency	Concentration on one model provider; pricing, availability or policy changes affect delivery.	Medium / Medium
R11	Regulatory and contractual non-compliance	Client contracts that pre-date AI adoption may not permit processing client data with a third-party model.	Medium / High
R12	Audit and SOX control impact	AI-generated configuration documentation used as control evidence without adequate provenance.	Low / High

Milestone 12 - AI Governance & Control Framework

2. Governance Controls

Data privacy (R1, R2, R11).

- Private network endpoint inside HCG's tenant; no public internet path for client content.
- Contractual zero-data-retention and no-training clauses with the model provider - a separate control from network isolation, and both are required.
- Hard partitioning by engagement across storage, retrieval indexes and logs; retrieval scoped by engagement identity at query construction so a cross-engagement query cannot be formed.
- Retrieval rather than fine-tuning, so client content never enters model weights and deletion on contract exit is genuinely complete.
- Engagement Partner sign-off and a client-facing AI usage clause before any client's content is onboarded. Consent is obtained, not assumed.

Hallucination mitigation (R3, R4, R5).

- Mandatory source citation - every asserted business rule carries a package and line reference or a document section. Uncited claims fail validation before reaching a human.
- VERIFIED / INFERRED / UNRESOLVED labelling on every extracted item, giving the model a permitted way to express uncertainty. Most hallucination is a model with no acceptable way to decline.
- Deterministic checks outside the model - coverage against a parsed object list, string comparison of configured values, requirement-to-test reconciliation. Where a check can be code, it is code.
- An independent assurance pass that fails closed, returning unsupported drafts to the agent rather than to the reviewer's queue.
- Self-consistency on high-stakes discrete outputs; disagreement escalates to a human.
- No artifact passes a gate without a named approver, with edit distance and override rate monitored per reviewer specifically to detect rubber-stamping.

Bias monitoring (R8).

- Defect triage accuracy tracked by module, severity and originating tester; systematic skew is investigated as a defect in the system, not as noise.
- Balanced worked examples across modules and complexity levels, reviewed quarterly.
- AI is explicitly excluded from any use in performance evaluation, staffing or promotion decisions.
- Where estimation is later introduced, outputs are reconciled against actuals quarterly for systematic optimism.

Milestone 12 - AI Governance & Control Framework

3. Monitoring and Audit Processes

What is logged. Every agent invocation records the triggering event, the prompt version, the model and model version, the full retrieval set, the raw output, every validation result, the human decision and identity, and the final artifact version. The log is append-only, immutable and tamper-evident.

The traceability store as the audit backbone. Because lineage is written before human review, the record is complete for rejected drafts as well as approved ones. For any line in any signed FSD the firm can answer: which legacy code it derives from, what the AI originally drafted, who changed it, who approved it, and under which prompt and model version. This is not a reporting feature added afterwards - it is the reason the system is auditable at all.

Continuous monitoring. The metric set in Milestone 8 runs continuously: quality (acceptance rate, edit distance, triage accuracy, override rate), integrity (traceability pass rate, uncited claims, coverage checks), safety (cross-engagement attempts, schema failures) and operational (latency, cost, availability). Safety metrics alert immediately; quality metrics are reviewed on a weekly trend.

Review cadence.

- Weekly - automated quality dashboard to Delivery Leads and the AI Platform Owner.
- Monthly - AI Governance Committee reviews accuracy trends, incidents, override patterns and knowledge base currency.
- Per engagement close - realised hours saved and defect outcomes reconciled against the Milestone 4 model, so the business case is re-tested rather than assumed.
- Quarterly - agent security review covering prompt injection and permission drift; Oracle release impact assessment on the knowledge base.
- Annually - independent assessment against the firm's AI policy and ISO/IEC 42001 control set.

Client-facing transparency. Where a client's content is processed, the engagement provides on request a record of which artifacts were AI-assisted, which human approved each one, and what controls applied. Volunteering this is treated as a commercial advantage rather than a disclosure risk.

Milestone 12 - AI Governance & Control Framework

4. Incident Response Strategy

Phase	Actions
Detect	Automated alerts on safety and integrity metrics; human reports through a single AI incident channel; client-raised concerns treated as detections. Any reviewer may raise an incident without escalating through their Delivery Lead first.
Triage and classify	AI Platform Owner classifies within 1 hour using the S1-S4 severity model in Milestone 8. S1 covers suspected data exposure or cross-engagement contamination; S2 an AI-originated error in an approved client deliverable; S3 systematic quality degradation; S4 availability.
Contain	Default action on S1 or S2 is to suspend the affected agent stage immediately, before root cause is known. Delivery continues on the manual workflow, which remains permanently supported - this is what makes stopping an affordable decision rather than a costly one.
Assess impact	Query the traceability store for every downstream artifact derived from the affected output. Because lineage is complete, blast radius is a query rather than an investigation. For S1, Firm Information Security and Legal determine the notification obligation under the client contract and applicable law.
Notify	S1: AI Platform Owner, Engagement Partner, Firm Information Security, Client DPO immediately; client notified per contractual timeframe. S2: Engagement Manager and Functional Lead within 24 hours, with a client conversation owned by the Engagement Manager. Incidents are disclosed to the client, not managed around them.
Remediate and verify	Correct the affected artifacts; fix the underlying prompt, validator, knowledge base entry or permission; add a regression case covering the failure. The agent stage is not reinstated until the new case passes.
Learn	Root cause review within 10 working days, reported to the AI Governance Committee. Every S1 and S2 produces a permanent regression test. A fix without a test is not treated as a fix.

Milestone 12 - AI Governance & Control Framework

5. Regulatory and Compliance Considerations

Regime	Relevance	Response
India DPDP Act 2023	Applies where client data contains personal data of data principals in India. HCG is typically a data processor acting on client instruction.	Processing limited to contracted purpose; data residency in India by default; breach notification path defined; DPIA completed before the pilot.
EU GDPR	Applies to engagements where client data includes EU personal data, including supplier master data in an AP implementation.	Standard contractual clauses where transfer occurs; regional endpoint pinning; documented lawful basis flowing from the client as controller.
EU AI Act	The system is an internal productivity tool. It is not a high-risk system under Annex III - it makes no decision about a person - but general-purpose AI transparency obligations apply.	Formal risk classification documented and reviewed annually rather than assumed; human oversight and logging already exceed the transparency baseline. Client-facing agents under Initiative 2 would require separate classification.
SOX and financial control frameworks	The most substantive consideration. The outputs configure and document a client's financial reporting systems, so configuration documentation may be used as control evidence.	Segregation of duties preserved - no agent applies configuration; provenance of every AI-assisted control document is available through the traceability store; client's external auditor engaged early where BR100 content supports control testing.
Client contractual terms	Many existing MSAs pre-date generative AI and may not permit processing client data with a third-party model.	Contract review per engagement before onboarding; explicit AI usage clause and client consent obtained; an engagement whose contract does not permit it stays on the manual workflow, with no exceptions.
ISO/IEC 42001 and internal policy	The firm's AI management system and its own responsible AI policy.	Controls mapped to the standard; annual independent assessment; AI Governance Committee is the standing accountable body.

Milestone 13 - 3-Year AI Capability Roadmap

1. Phase 1 - Pilot Implementation (Months 1-6)

Scope. One engagement, chosen for a contract that already permits AI processing and a Delivery Lead who volunteered. Two agents only: the Legacy Analysis Agent and the Design Authoring Agent, limited to logic extracts and FSDs. Test authoring, BR100 generation and defect triage stay fully manual throughout Phase 1.

Why so narrow. Legacy analysis is the biggest bottleneck and the easiest to verify objectively - a consultant can check an extract against source code and be certain whether it is right. Starting where correctness is checkable means the pilot produces evidence rather than opinion.

Activities. Stand up the platform in the firm's tenant; build the prompt library and the deterministic validators; index that engagement's corpus; complete the security review and DPIA; train the pilot team in review as a distinct skill; run in parallel with manual authoring for the first four weeks to establish a like-for-like baseline.

Success criteria - the gate to Phase 2.

- Draft acceptance rate of at least 50% approved with no more than minor edits.
- At least 30% reduction in FSD authoring hours against the measured baseline.
- 100% traceability validation pass rate - no uncited claim reaches a reviewer.
- Zero data incidents of any severity.
- Legacy coverage checks pass on every extract.
- Pilot reviewers report the drafts are worth reviewing - a subjective measure, deliberately included, because a platform consultants resent will not survive scaling.

Decision point. If the criteria are not met, Phase 2 does not proceed on schedule. The remediation path is to extend the pilot, not to scale and hope. This gate is the principal risk control on the whole investment - Rs 45 lakh of build is committed before it, and the remainder after.

Milestone 13 - 3-Year AI Capability Roadmap

2. Phase 2 - Scaled Deployment (Months 7-18)

Scope. Add the Test Authoring Agent and the Delivery Support Agent, bringing the full five-task scope into service. Extend BR100 generation from live instance extracts. Roll out across all ten engagements, onboarding each only after its contract review and Engagement Partner sign-off.

Organisational changes. The responsibility matrix from Milestone 7 comes into force practice-wide. The AI Governance Committee begins its monthly cadence. Review-skills training is added to onboarding for every consultant, not offered as optional enablement. Defect triage moves from L2 assisted to L3 supervised autonomy once accuracy exceeds 85% for four consecutive weeks - a measured promotion, not a scheduled one.

Platform maturation. The knowledge base becomes cross-engagement for practice content - templates, patterns, Oracle documentation - while client content stays hard-partitioned. The regression suite grows with every incident. Prompt performance is tuned against real acceptance and override data rather than against the initial held-out set.

Targets by end of Phase 2. Draft acceptance rate of 70% or better; the full Rs 1.90 crore annual run rate of benefit realised; a 20-25% reduction in defects traced to design ambiguity; BR100 drift reduced from 15% to under 5%; all ten engagements live.

Principal risk in this phase. Scaling adoption across roughly 120 consultants with varying enthusiasm. The mitigation is measurement per reviewer rather than per engagement, so that a team quietly reverting to manual authoring - or quietly rubber-stamping - is visible early and can be addressed as a conversation rather than discovered in a quality incident.

Milestone 13 - 3-Year AI Capability Roadmap

3. Phase 3 - Autonomous Expansion (Months 19-36)

Selective autonomy, earned by evidence. Where an agent has demonstrated sustained accuracy, 100% human review is replaced by risk-based sampling. Test script generation and defect triage move to sampled review at 20%. Daily SIT status reporting is issued automatically with the Test Lead on exception. FSD sign-off, configuration and cutover remain at full human review permanently - these boundaries do not move, in this phase or any later one.

The traceability store becomes a knowledge graph. With three years of engagements captured, the lineage data supports questions no individual could answer: which requirement patterns reliably produce defects, which legacy constructs are consistently mis-specified, which design decisions were later reversed and why. This is the asset the initiative was really building.

Oracle release-impact analysis. A new agent reads each quarterly Oracle Cloud release note against the knowledge graph and predicts which client configurations, integrations and test scripts are affected - converting a recurring manual regression-scoping exercise into a targeted one. This is a capability that only becomes possible once the graph exists.

Initiative 2 brought forward. With a governance model proven over three years of internal operation, the practice can credibly sell Oracle agentic enablement to clients - offering not just configuration of Fusion's bundled agents but a tested framework for guardrails, escalation, audit and incident response. This is the deferred initiative from Milestone 9 re-entering on the strength of evidence rather than ambition.

Explicitly out of scope, and why. Fully autonomous configuration deployment - an agent applying setups to a client environment with automated rollback - is technically foreseeable and is deliberately declined. It requires a trust bar in agent judgement that three years of drafting accuracy does not establish, and it would compromise the segregation of duties that makes the SOX position defensible. It is recorded here as rejected rather than left unmentioned, so that the boundary is a decision the practice has taken and can be held to.

Milestone 13 - 3-Year AI Capability Roadmap

4. Future Capability Enhancements

Multimodal input. Design sessions produce whiteboard photographs, screenshots of legacy forms and annotated process diagrams that are currently transcribed by hand or lost. Multimodal models could read them directly, closing the gap between what is discussed in a design session and what reaches the RTM - the very first handoff in the chain, and the one the current solution does not touch.

Voice-to-artifact from design sessions. Transcribing recorded design sessions into structured requirement candidates with speaker attribution, so decisions are captured as they are made rather than reconstructed afterwards from someone's notes.

Predictive analytics on the knowledge graph. Defect-leakage prediction - flagging requirements whose characteristics historically correlate with UAT defects, at design time. Effort estimation grounded in actuals rather than expert judgement, which is Initiative 3 from Milestone 9 finally becoming feasible because the data now exists.

Evaluator agents. Automated LLM-as-judge scoring of artifact quality against the practice rubric, enabling continuous evaluation at a scale no human review panel could sustain. Introduced as a monitoring aid, never as a replacement for a human gate.

Smaller specialised models. As open-weight models improve, fine-tuned small models for narrow, high-volume tasks such as defect classification could cut inference cost and latency and could run entirely within the firm's tenant - reducing the vendor dependency accepted in Milestone 11 rather than merely surviving it.

The capability that matters most. None of the above is the real enhancement. The compounding asset is the traceability graph itself: a structured, three-year record of how requirements became systems across many engagements. Models will keep improving and will be swapped out. The graph is the thing a competitor cannot buy, and every future capability listed here is a way of querying it.

Milestone 14 - Executive AI Transformation Proposal

Executive Summary

HCG's Oracle Cloud ERP practice delivers about ten Oracle Cloud Financials implementations a year. Five repetitive tasks in the design-to-test chain consume roughly 23,000 consultant hours annually and cause most late-stage quality problems. This proposal recommends four AI agents producing source-cited drafts on a shared traceability store, with a named human approving every artifact at one of six gates.

Dimension	Position
The ask	Rs 1.08 crore in Year 1; Rs 42 lakh a year thereafter.
The return	Rs 1.90 crore a year at steady state. 157% three-year ROI. Payback ~13 months.
Quality case	20-25% fewer defects traced to design ambiguity; BR100 drift from 15% to under 5%.
Governing principle	AI drafts, humans decide. Autonomy covers drafting, never state change in a client system.
Risk control	Phase 1 is a single-engagement pilot with published criteria; only Rs 45 lakh is committed before that gate.
Declined	Autonomous configuration deployment; client-facing agents until governance is proven.

At 30% average time saved rather than 46%, the three-year ROI is still positive at roughly 68%.

Milestone 14 - Executive AI Transformation Proposal

1. Problem Overview and Organisational Context

HCG's Oracle Cloud ERP practice implements Oracle Cloud Financials for enterprises migrating off legacy E-Business Suite R12. A typical engagement runs 12-18 months with 18-25 consultants, producing about 60 RICE objects, 25 BR100 workbooks, 800 test scripts and 600 defects. Work is fixed-fee, so margin depends directly on the hours needed to produce a correct artifact first time.

The problem is manual translation. A requirement travels from design session to RTM row to FSD to BR100 to test script, each handoff a re-typing. Meaning drifts and surfaces late: about 22% of SIT defects trace to design ambiguity rather than coding error. Upstream sits the real bottleneck - reading legacy PL/SQL to establish what the system does, at roughly eight hours per object. This matters on three counts. A defect found in UAT costs many times what clarifying the same ambiguity at design time would have. The people consumed - Functional Leads, Architects, Test Leads - are the binding constraint on concurrent engagements. And Oracle now bundles agentic capability into Fusion Cloud, so differentiation moves to delivery speed and accuracy.

Ten sequential steps run from design session to hypercare; steps 3-8 are in scope and steps 1-2 and 9-10 stay fully human-led. Almost everything in scope is manual - only data loads, integration deployment and regression execution are partly automated, and the specification work around them is still manual.

Named bottleneck	Baseline metric
Legacy comprehension gates everything downstream	8 hrs per object; 4,800 hrs a year
Sequential handoffs, each waiting on the last	FSD elapsed cycle 5-7 days
Review queues at two or three senior reviewers	720 hrs a year of FSD authoring
Test authoring lags design	12% of requirements unscripted at SIT entry
Defect triage congestion at peak SIT	45 min per defect; 18% misrouted
Documentation drifts from the instance	15% of BR100 entries mismatched by UAT
Total, five in-scope tasks	2,300 hrs per engagement; 23,000 a year

Milestone 14 - Executive AI Transformation Proposal

3. Generative AI Opportunity and Redesigned Workflow

Every artifact here is structured text derived from other structured text, against a corpus of approved documents HCG already owns. Retrieval produces house-style drafts without fine-tuning, keeping client content out of model weights.

AI does: exhaustive legacy logic extraction with a source reference and a VERIFIED / INFERRED / UNRESOLVED label on every rule; FSD drafting; BR100 generation from a read-only instance extract; test scripts tagged to requirement IDs; defect classification with a cited root-cause hypothesis; status reports and cutover checklists.

Humans keep: design sessions; the fit-gap decision; verification of flagged legacy items; FSD sign-off; applying configuration; test coverage adequacy; confirming routing; defect closure; UAT sign-off; every cutover decision.

How it changes. The gates are unchanged; what arrives at each one changes. Reviewers start from a cited draft, so work shifts from authoring to judging, and test authoring becomes a consequence of FSD sign-off rather than a following phase. An assurance check sits between agent and reviewer and fails closed.

Prompting. Decomposition forces legacy coverage; few-shot chain-of-thought fixes reasoning and style; tree-of-thoughts evaluates four fit-gap options; ReAct interleaves retrieval with reasoning in the assurance check; self-consistency covers the two outputs where error is most expensive. All outputs are schema-constrained so deterministic code verifies them.

Milestone 14 - Executive AI Transformation Proposal

4. AI Agent Architecture and Automation Design

Agent	Produces	Gate	Approver
Legacy Analysis Agent	Logic extract with references and status labels	1	Functional Consultant
Design Authoring Agent	FSD drafts; BR100 from live instance extracts	2, 3	Functional Lead
Test Authoring Agent	Scripts tagged to requirement IDs	4	Test Lead
Delivery Support Agent	Triage notes, status, cutover checklist	5, 6	Defect / Programme Manager

Four specialised agents rather than one, because a single agent would need the union of all four permission sets - exactly what the governance model forbids.

Planning. Receive event, assemble context, plan, draft with citations, self-check, submit to the assurance pass, escalate after two failed iterations, write lineage, hand to the gate.

Tools. Read-only legacy repository and non-production Oracle instances via REST and BI Publisher; read and gated write to Jira, Confluence and ALM; read/write to the traceability store; partitioned retrieval over the approved corpus. No agent holds write access to an Oracle environment, the document store, or a defect's status.

Memory. Short-term task context discarded on completion; engagement memory (RTM, extracts, decisions, defects) hard-partitioned per client; practice memory of templates and Oracle documentation with no client-identifying content. Binding these is the traceability store - a graph from requirement to extract to FSD section to configuration to test script to defect, each node recording source, prompt and model version, and approver.

Autonomy. Analysis, assurance and test generation run supervised; drafting is assisted; configuration, defect closure, test sign-off and cutover are not permitted at any level. A reviewer may accept unchanged, amend, return for redrafting, or discard and author manually. Execution cost is about Rs 35,000 per engagement, 0.2% of the value released.

Milestone 14 - Executive AI Transformation Proposal

5. Financial Impact Analysis

Repetitive task	Baseline hrs/yr	% saved	Hours saved	Annual saving
Legacy EBS code reverse-engineering	4,800	50%	2,400	Rs 43.20 lakh
FSD authoring	7,200	40%	2,880	Rs 51.84 lakh
BR100 configuration documentation	2,500	60%	1,500	Rs 27.00 lakh
SIT/UAT test script authoring	4,000	55%	2,200	Rs 39.60 lakh
SIT/UAT defect triage and routing	4,500	35%	1,575	Rs 28.35 lakh
TOTAL (blended Rs 1,800 per hour)	23,000	46%	10,555	Rs 189.99 lakh

3-year model	Year 1	Year 2	Year 3	Total
Gross benefit (50% / 100% / 110%)	Rs 95 L	Rs 190 L	Rs 209 L	Rs 494 L
Cost	Rs 108 L	Rs 42 L	Rs 42 L	Rs 192 L
Net	(Rs 13 L)	Rs 148 L	Rs 167 L	Rs 302 L

Three-year ROI 157%; payback about 13 months. Year 1 covers model consumption, platform build, integration, knowledge base, security and change management, of which ~Rs 90 lakh is one-time.

Deliberately excluded. No benefit is claimed for reduced defect leakage, avoided slippage or improved win rates, though all are expected. Including them would double the headline ROI and halve its credibility. A case built on hours saved is one the practice can be held to, and it is reconciled against actuals at every engagement close.

Milestone 14 - Executive AI Transformation Proposal

6. Strategic Alignment

Three initiatives were compared - this platform, a client-facing Oracle agentic enablement offering, and a bid and estimation capability - on feasibility, ROI, implementation risk and strategic impact. This one scored 4.80 against 3.45 and 3.35, and the ranking holds even when weighting is shifted to favour revenue.

Why first. It is the only one the firm can start immediately - no client permission, no new revenue, no data it does not own - and its benefit applies to every engagement. Its failure mode is cheap: revert to manual authoring. It also builds the prerequisites for the other two: the traceability asset that makes estimation credible and the governance competence the client-facing offering would sell.

Practice fit. Margin and senior capacity are the binding constraints and this attacks both. It also strengthens the client's audit position, since lineage and instance-derived documentation are control evidence - a CFO conversation, not only a delivery one.

Benchmarks. Accenture reported generative and agentic AI revenue tripling to about USD 2.7 billion in the twelve months to August 2025 and reorganised five business units into one reinvention services organisation, concluding that capturing AI value meant changing how delivery is organised. Infosys has deployed GitHub Copilot to over 30,000 developers and applies Topaz across the SDLC for automated documentation, testing and legacy modernisation - the same categories targeted here - reporting AI revenues at 8.2% of revenue in Q1 FY27.

What they teach. Both report in incomparable units, none transferable to a per-engagement hours model - which is why this case is built from HCG's own baseline. The lesson is that this layer is already being industrialised by larger competitors: a reason to move now, and not to overstate the advantage. HCG's defensible position is domain depth and governance, not scale.

Milestone 14 - Executive AI Transformation Proposal

7. Governance and Risk Management Framework

Data privacy. All inference is reached over a private network path terminating inside HCG's own cloud tenant, and separately the provider agreement carries zero-retention and no-training clauses. Client content is hard-partitioned by engagement across storage, indexes and logs, with retrieval scoped by engagement identity, so cross-engagement retrieval is structurally unavailable. Retrieval rather than fine-tuning makes deletion on exit complete.

Hallucination. Every asserted rule carries a source reference; uncited claims fail validation before a human sees them. Every extracted item is labelled VERIFIED, INFERRED or UNRESOLVED. Coverage, value comparison and requirement-to-test reconciliation are deterministic code, not model judgement. The assurance pass fails closed and self-consistency covers high-stakes verdicts.

Bias. Triage accuracy is tracked by module, severity and tester so systematic skew is treated as a system defect. Examples are balanced and reviewed quarterly. AI is excluded from performance evaluation and staffing.

Monitoring and audit. Every prompt, retrieval set, model version, output and human decision is recorded append-only. Because lineage is written before review, rejected drafts are as auditable as approved ones. Weekly dashboard, monthly committee, per-engagement reconciliation, quarterly security review, annual independent assessment.

Incident response. Four severity levels. The default on suspected data exposure or an AI-originated error in an approved deliverable is to suspend the agent stage before root cause is known - affordable because manual authoring remains supported. Blast radius is a traceability query, and every serious incident produces a permanent regression test.

The honest risk. Data leakage, cross-engagement contamination, hallucinated content, incomplete coverage, prompt injection, Oracle update drift, skill erosion and vendor dependency all have owners and mitigations. The hardest is none of these: it is rubber-stamp review. A fluent draft invites approval, and if reviewers stop scrutinising, every quality claim here collapses. Edit distance and override rate are monitored per reviewer.

Regulatory. India DPDP Act 2023 and GDPR where applicable; EU AI Act classification reviewed annually; SOX segregation of duties preserved because no agent applies configuration; per-engagement contract review, with any engagement lacking AI permissions staying manual.

Milestone 14 - Executive AI Transformation Proposal

8. Implementation Roadmap

	Phase 1 - Pilot (M1-6)	Phase 2 - Scaled (M7-18)	Phase 3 - Autonomous (M19-36)
Scope	One engagement. Legacy Analysis and Design Authoring agents; extracts and FSDs only.	All four agents, five tasks, ten engagements. BR100 from live instance extracts.	Sampled review for low-risk artifacts. Oracle release-impact analysis. Client-facing enablement.
Why	Legacy analysis is the biggest bottleneck and easiest to verify against source, so the pilot yields evidence rather than opinion.	Benefit compounds only when the whole chain is covered; partial coverage leaves the dependency intact.	Autonomy is earned by measured accuracy, never granted by schedule.
Activities	Platform build; validators; security review and DPIA; four weeks parallel running for a baseline; review-skills training.	Responsibility matrix practice-wide; governance committee monthly; triage promoted to supervised autonomy once accuracy exceeds 85% for four weeks.	Traceability store matures into a knowledge graph; predictive defect-leakage and estimation; Initiative 2 brought forward on proven governance.
Gate	50% acceptance; 30% FSD hours reduction; 100% traceability pass; zero data incidents. Phase 2 does not proceed if unmet.	70% acceptance; full Rs 1.90 crore run rate; 20-25% fewer ambiguity defects; drift under 5%.	Sampled review at 20%; knowledge graph operational across three years.

The gate is the principal risk control. Only about Rs 45 lakh is committed before the Phase 1 decision. If criteria are unmet, the path is to extend the pilot, not to scale and hope.

Permanently out of scope. FSD sign-off, configuration and cutover stay at full human review in every phase. Autonomous configuration deployment is deliberately declined - it would require a trust bar drafting accuracy does not establish and would compromise SOX segregation of duties.



सिद्धिमूलं प्रबन्धनम्
भा. प्र. सं. इन्दौर
IIM INDORE

भारतीय प्रबंध संस्थान इंदौर

INDIAN INSTITUTE OF MANAGEMENT INDORE

