

Video Transcripts

Week 1.2: Impact on Enterprise Value Chains

Video 1: Generative AI as a Value-Chain Transformer

Across enterprises today, AI is no longer confined to isolated tools or experimental use cases. It is increasingly influencing how organisations design, coordinate, and execute work across entire value chains. This session explores how classical AI and generative AI are reshaping enterprise value chains and, in doing so, transforming how organisations create value at scale.

When organisations first encounter generative AI, it is often framed as a productivity enhancer. Employees use it to write faster, summarise documents, generate content, or reduce time spent on routine work. These gains are real and meaningful. However, they represent only a small part of the broader impact. To understand why generative AI matters strategically, we must move beyond individual tasks and examine how value is created end to end.

A value chain describes how an organisation converts inputs into outcomes. It spans internal processes, decision points, systems, and people, from early analysis through execution and customer interaction. When AI is assessed only at the task level, its contribution appears incremental. When assessed at the value-chain level, its impact becomes structural. This is why value chains, rather than tasks, are the appropriate unit of analysis for enterprise AI adoption. Traditional automation focused on optimising discrete activities. Generative AI introduces a different capability. It enables organisations to rethink how information flows between stages, how insights are converted into action, and how decisions are communicated across functions. This marks a shift from task automation to workflow transformation.

Task automation aims to do the same work faster. Workflow transformation asks whether the work itself should be structured differently. Generative AI supports this shift by operating across steps rather than within a single step. It connects analysis, creation, communication, and execution into more integrated workflows. Marketing provides a clear illustration. Traditionally, the marketing value chain moved sequentially through research, strategy, creative development, execution, and performance analysis. Each stage involved different teams and tools, with frequent handoffs and delays. Generative AI enables tighter integration across these stages. Customer and search insights can directly inform messaging. Draft content can be generated, adapted, and tested rapidly. Creative assets can evolve alongside performance feedback. The result is not just faster content creation, but a more responsive and adaptive marketing value chain.

Operations functions experience a similar shift. Operational value chains rely heavily on planning, coordination, and communication. Predictive systems forecast demand or identify bottlenecks, but insights often remain locked in dashboards or reports. Generative AI helps bridge this gap by translating analytical outputs into operational plans, summaries, and clear explanations that teams can act on. When forecasts change, generative AI can help update plans, explain implications, and support scenario discussions, reducing friction between insight and execution.

Finance value chains also benefit. Finance teams rely on forecasting, budgeting, reporting, and decision support. Predictive models estimate outcomes, but generative AI supports interpretation and communication. It can draft management reports, explain variances in plain language, and outline alternative scenarios for leadership review. This allows finance professionals to spend less time preparing materials and more time advising on decisions.

Human resources value chains span recruitment, onboarding, learning, performance, and engagement. Generative AI can assist by drafting job descriptions, personalising onboarding content, generating training materials, and summarising employee feedback. More importantly, it enables personalisation at scale, improving employee experience while freeing HR teams to focus on judgement, culture, and strategic priorities.

Customer support value chains are similarly transformed. Traditionally, support involved ticket intake, classification, response, and resolution across disconnected systems. Generative AI helps unify these steps by summarising issues, drafting responses, and supporting agents during live interactions. Resolution times improve, consistency increases, and human agents remain responsible for final decisions and escalation.

A real-world illustration of value-chain transformation can be seen in Amazon's customer operations. Amazon combines predictive analytics with generative AI to support its order management and customer support workflows. Predictive models anticipate delivery issues before customers raise complaints. Generative AI then supports proactive notifications, suggests compensation options, and assists agents with contextual responses. This integration reshapes the flow of value across the chain rather than simply automating individual tasks. Across these functions, a consistent pattern emerges. Generative AI reduces handoffs, compresses cycles, and connects previously fragmented stages of work. This is what makes it a value-chain transformer. It changes how value is created, not just how quickly tasks are completed.

However, this transformation does not happen automatically. Generative AI requires thoughtful design, clear prompting, and strong human oversight. Outputs must be reviewed, contextualised, and aligned with organisational goals. As generative AI moves closer to core workflows, governance becomes increasingly important.

Leadership plays a critical role in this transition. Rather than asking where generative AI can save time, leaders should ask where it can remove friction across the value chain. Where do handoffs slow progress? Where is insight lost between analysis and action? Where do teams spend effort translating information instead of acting on it? These are the points where generative AI delivers the greatest leverage. Viewing generative AI through a value-chain lens also changes how success is measured. Beyond productivity, organisations should track cycle time, responsiveness, outcome quality, and employee experience. The objective is not simply to work faster, but to coordinate better and execute smarter.

In summary, generative AI represents a fundamental shift in enterprise value creation. When applied at the value-chain level, it enables deeper integration across marketing, operations, finance, HR, and customer support. Used thoughtfully, it becomes not just a tool for efficiency, but a strategic lever for enterprise transformation.

Video 2: Generative AI in Customer-Facing Functions

Customer-facing functions sit at the visible edge of the enterprise value chain. Every advertisement a customer encounters, every message they receive, and every support interaction they experience directly shapes trust in the brand. Marketing and customer support are where organisational intent meets customer perception. As customer expectations rise and digital touchpoints multiply, slow, manual, or fragmented workflows increasingly become a competitive liability.

Generative AI is emerging as a powerful enabler in this space, not by replacing human interaction, but by reshaping how customer-facing value chains operate at scale. Its impact lies not in automating isolated tasks, but in connecting insight, content creation, communication, and response into more integrated customer journeys.

Marketing value chains typically span research, content creation, campaign execution, testing, and performance analysis. In many organisations, these stages operate in silos. Insights generated by analytics teams are passed to content teams, which then hand work over to execution teams. Feedback loops are slow, experimentation is constrained, and learning happens after campaigns are already in market.

Generative AI compresses this value chain by reducing handoffs and cycle times. Marketers can generate content drafts, personalise messaging, and test variations rapidly. Instead of producing a single campaign asset, teams can create multiple versions tailored to different segments, channels, or contexts.

Content generation becomes both faster and more flexible. Blog posts, ad copy, email subject lines, and landing-page text can be drafted in minutes rather than days. More importantly, these drafts can be informed by customer data, search trends, and historical performance, improving relevance from the outset rather than through later optimisation. Personalisation also scales more effectively. Generative AI can adapt tone, format, and emphasis based on audience context without requiring extensive manual rewriting. This allows organisations to move beyond broad segmentation toward more responsive, context-aware communication.

Rapid experimentation becomes the norm rather than the exception. Multiple headlines, creative angles, or calls to action can be tested simultaneously, with performance feedback feeding directly into the next iteration. The marketing value chain shifts from a linear campaign model to a continuous optimisation loop. Customer support value chains experience a similar transformation. Traditionally, support involves issue intake, classification, response formulation, resolution, and follow-up across multiple systems. Agents spend significant time searching knowledge bases, drafting responses, and summarising cases rather than resolving issues.

Generative AI helps unify these stages. At intake, it can summarise customer issues and highlight key details. It can suggest relevant knowledge articles, draft response messages, and provide contextual guidance during live interactions. This reduces friction between understanding a problem and resolving it. Response drafting improves both speed and consistency. Instead of writing replies from scratch, agents receive suggested responses aligned with approved tone, policy, and structure. Agents

retain control, reviewing and adapting responses before sending them, but the cognitive and time burden is significantly reduced.

Knowledge synthesis is another critical benefit. Generative AI can combine information from product documentation, policies, and past cases into clear, actionable guidance. This shortens resolution times and lowers the learning curve for newer agents. Agent-assist capabilities further enhance performance by summarising conversations, highlighting unresolved issues, and suggesting next steps. This allows agents to focus on problem-solving and empathy rather than administrative overhead. Human oversight remains essential, ensuring accuracy, appropriate escalation, and accountability.

A real-world example of this transformation can be seen at Shopify. Shopify uses generative AI across both marketing and customer support to improve experiences for merchants and end customers.

In marketing, AI-powered tools help merchants generate product descriptions, ad copy, and email content tailored to different audiences. This enables small businesses to launch campaigns quickly and experiment with messaging without requiring specialised marketing expertise.

In customer support, generative AI is integrated into helpdesk workflows to summarise issues, suggest relevant documentation, and draft response messages for agents. This improves response times and consistency, particularly during peak demand periods, while keeping humans in control of final decisions.

By embedding generative AI across customer-facing workflows, Shopify improves speed, maintains brand consistency, and scales support without proportional increases in headcount. This illustrates how generative AI reshapes value chains rather than automating individual tasks. Across marketing and customer support, three consistent impacts emerge. First, speed improves significantly as content creation and response drafting accelerate. Second, consistency increases as messaging aligns more closely with brand guidelines and policies. Third, scale becomes achievable without linear growth in effort, enabling personalised engagement and high-quality support at volume.

These outcomes arise not because generative AI replaces human roles, but because it removes friction across the value chain. It connects insight to action and reduces the effort required to move work from one stage to the next.

In summary, generative AI transforms customer-facing value chains by accelerating creation, enabling scalable personalisation, and supporting faster, more consistent customer support. Its impact is visible not just in individual tasks, but in how marketing and support operate end to end.

Video 3: Generative AI in Operational and Financial Workflows

Operations and finance are where plans turn into results. Every forecast, procedure, and report directly shapes how efficiently an organisation runs and how confidently leaders make decisions. Yet these functions are often slowed by heavy documentation, constant coordination, and the effort required to translate analysis into action. Generative AI is emerging as a practical enabler in this space, not by replacing expertise, but by reducing friction in execution-heavy workflows.

Operational value chains typically involve planning, execution, monitoring, and exception handling. These processes depend on standard operating procedures, cross-team coordination, and timely communication. Traditionally, they are slowed by manual documentation, fragmented tools, and delays in responding to change.

Generative AI improves operational execution across three key areas: SOP generation, exception handling, and planning support. SOP generation becomes faster and more adaptable. Instead of manually drafting or updating procedures, operations teams can use generative AI to produce draft SOPs based on existing policies, system data, and operational constraints. These drafts are then reviewed and approved by leaders. SOPs can be tailored to specific locations, roles, or scenarios, reducing ambiguity and improving compliance without increasing administrative burden.

Exception handling is another major area of impact. Operational disruptions—such as supply delays, equipment failures, or sudden demand shifts—require rapid coordination. Predictive systems may flag anomalies, but generative AI helps translate those alerts into actionable guidance. It can summarise the issue, outline response options, and draft clear instructions for affected teams. This shortens response time and ensures more consistent handling under pressure. Planning support also improves significantly. Generative AI assists planners by summarising operational plans, explaining trade-offs, and supporting scenario comparisons. When assumptions change, updated plans and explanations can be generated quickly, keeping teams aligned without repeated manual effort. Overall, this reduces coordination overhead and improves execution speed across operations.

Financial value chains face a similar challenge. They include forecasting, budgeting, reporting, variance analysis, and decision support. While highly analytical, these workflows are also documentation-intensive. Finance teams often spend disproportionate time preparing reports and explaining numbers rather than advising on decisions. Generative AI supports finance functions by enhancing narrative reporting, variance explanation, and decision support.

Narrative reporting becomes more efficient and accessible. Generative AI can convert financial data into structured management commentary, highlighting key trends, risks, and performance drivers. This reduces time spent drafting reports while improving clarity for non-financial stakeholders. Variance explanations are another critical use case. When actual results differ from forecasts or budgets, generative AI can help analyse drivers and draft explanations in plain language. This enables faster review cycles and more focused leadership discussions.

Decision support also improves through better interpretation of financial scenarios. Generative AI can outline implications of different assumptions, summarise trade-offs, and frame options for decision-makers. This shortens the distance between analysis and action, allowing finance teams to focus on judgement and advisory work rather than report preparation. Importantly, generative AI operates under human oversight. Financial accountability, controls, and governance remain firmly with finance professionals.

A real-world illustration of this approach can be seen at Unilever. Unilever applies AI across selected supply chain and finance workflows to improve planning, reporting, and execution. In operations, predictive models identify demand shifts and supply risks. Generative AI then helps produce planning summaries, exception reports, and operational guidance for regional teams. This enables faster responses to disruption while maintaining consistency across markets.

In finance, AI-assisted reporting supports the generation of management narratives and explanations of performance variances across business units. Finance teams spend less time compiling information and more time supporting strategic decisions. The result is reduced coordination overhead and improved enterprise agility.

Across operations and finance, a consistent benefit emerges: reduced documentation and coordination burden. Teams spend less time drafting, summarising, and translating information. Insights move more smoothly from systems to people. Decisions are supported by clearer context and faster communication. This does not eliminate complexity, but it makes complexity more manageable. Generative AI absorbs repetitive cognitive work, allowing professionals to focus on judgement, problem-solving, and strategy.

That said, deployment must be disciplined. Operational and financial workflows are mission-critical. Outputs must be reviewed, validated, and aligned with organisational controls. Clear governance, access management, and accountability are essential. Generative AI should therefore be viewed as an augmentation layer, not an autonomous decision-maker.

In summary, generative AI is reshaping operational and financial workflows by reducing friction in execution-heavy and analytical value chains. In operations, it supports SOP generation, exception handling, and planning coordination. In finance, it enhances narrative reporting, variance explanation, and decision support.

Video 4: Generative AI in Human Resources

Human resources shapes how people experience an organisation, from recruitment through to long-term development. However, HR teams often spend large amounts of time drafting documents, maintaining policies, and producing content at scale. Generative AI is becoming a powerful enabler, helping HR work more efficiently while allowing professionals to focus on people, judgement, and strategy. Rather than replacing human decision-making, generative AI supports the creation, communication, and delivery of HR work across the employee lifecycle, improving speed and consistency.

A core HR activity is job design, which includes job descriptions, role expectations, and skill requirements. These documents are essential but traditionally time-consuming to create and maintain. Generative AI accelerates this process by producing structured role descriptions based on inputs such as responsibilities, required skills, and organisational standards. It can also analyse role patterns, highlight overlaps, and draft progression pathways, which is especially valuable in large organisations. HR professionals still validate and align these outputs with organisational goals, ensuring accuracy and relevance.

Learning and development is another high-impact area. Creating onboarding materials, training content, assessments, and scenarios can be resource-heavy. Generative AI assists by drafting tailored learning materials that reflect role requirements, skill levels, and business needs. It also supports personalised onboarding and enables continuous learning by allowing rapid updates as technology or processes change. While AI speeds up content creation, HR and learning specialists remain responsible

for reviewing and contextualising material to maintain quality and alignment with organisational values.

Generative AI also enhances policy management and internal communication. It can help draft policy documents, simplify complex language, and prepare clear employee-focused summaries. AI-generated announcements, FAQs, and guidance improve communication consistency, particularly during organisational change. Real-world examples, such as IBM, demonstrate how generative AI integrates job design, learning, and communication to improve HR efficiency and employee experience. Across these areas, the benefits are clear: faster content creation, greater consistency, and the ability to scale personalised experiences. With proper oversight, governance, and ethical use, generative AI enables HR teams to shift their focus from administrative tasks to higher-value work such as culture building, workforce planning, and talent strategy.

Video 5: Workflow Compression and Cognitive Leverage

Many organisations do not lack talent; instead, they struggle with slow handoffs, fragmented workflows, and limited access to expert thinking. Traditional work moves sequentially through multiple teams, causing delays and repeated reinterpretation of information. Generative AI transforms this structure by reshaping how work flows and how expertise scales. This shift can be understood through two key concepts: workflow compression and cognitive leverage.

Workflow compression reduces the time, steps, and friction between the start and end of a task. Rather than progressing linearly through analysis, drafting, review, and approval, generative AI enables several of these activities to occur simultaneously or more closely together. It also minimises unnecessary transitions by reducing the need to translate information across formats or teams. Handoffs—often the biggest source of delay—are reduced as AI supports work across functional boundaries. As a result, organisations experience faster cycles, quicker decisions, and work that adapts more readily to changing conditions. Parallelisation further enhances this by allowing multiple drafts, scenarios, or explanations to be generated at once, expanding options without additional effort.

Cognitive leverage refers to scaling expert thinking across the organisation. Expertise is often scarce, and access to specialists can become a bottleneck. Generative AI reflects common expert reasoning patterns, enabling employees to access structured explanations, best-practice frameworks, and decision guidance without waiting for specialist input. This lifts baseline capability, supports stronger decision-making, and allows experts to focus on high-value, complex judgement rather than repeated explanation. It also strengthens learning by exposing employees to higher-quality reasoning patterns more frequently.

A practical example can be seen at McKinsey & Company, where generative AI supports consultants with summarised research, structured presentations, and draft client materials. This compresses workflows by reducing handoffs between research, analysis, and communication teams, while cognitive leverage is achieved by embedding firm-wide frameworks and expertise into AI tools. Together, workflow compression and cognitive leverage form a structural advantage: work moves faster, coordination costs fall, and quality improves. To realise these benefits responsibly, organisations must integrate AI into workflows carefully, apply clear governance, and maintain strong human oversight.

Video 6: Distinguishing Generative AI and Agentic AI

Generative AI and agentic AI are often discussed together, yet they serve distinct roles. Generative AI is reactive: it produces content such as drafts, summaries, explanations, plans, and code when prompted, but it does not hold goals or take action across steps. Agentic AI, by contrast, is designed for action. It sets goals, plans sequences, uses tools and APIs, adapts to new information, and escalates to humans when rules are unclear or risk increases. In essence, generative AI accelerates thinking, while agentic AI accelerates doing.

An airline disruption illustrates this difference clearly. During a storm-related shutdown, a generative model can rapidly create passenger updates, summarise fare rules, suggest alternative routings, or produce call-centre scripts—but it cannot identify affected passengers, check booking eligibility, or change reservations. Its work ends when the prompt ends. An agent, however, interprets the disruption as a trigger. It identifies affected bookings, prioritises vulnerable or high-status passengers, checks fare conditions, searches partner-airline inventory, performs compliant rebooking, issues boarding passes, and sends personalised notifications—escalating ambiguous cases to human agents with a summary and recommended actions.

In practice, generative and agentic AI reinforce one another. The agent executes policy at speed and scale, while generative AI supports communication, drafting messages for different customer segments, preparing empathetic scripts, producing post-incident reports, and highlighting patterns that may prompt policy improvements. This separation of duties matters: relying on generative AI for execution results in good guidance but no action, while giving an agent too much autonomy without clear rules risks over-automation. Effective programmes begin with narrow, well-bounded agents, full observability, and strong governance, expanding autonomy only as confidence grows.

For teams, the nature of work shifts accordingly. With generative AI, humans focus on reviewing drafts, validating logic, and setting tone. With agentic AI, they design policies, guardrails, and exception pathways, stepping in when uncertainty or conflicting rules arise. Together, they narrow the gap between planning and delivery. The simplest way to remember the distinction is this: generative AI explains what should happen; agentic AI makes it happen. Clear boundaries allow both systems to complement each other and deliver reliable outcomes.

Video 7: From Copilots to Goal-Driven Digital Workers

Most people first experience generative AI as a copilot—reactive, prompt-driven, and focused on assisting with individual tasks such as drafting emails, summarising documents, or exploring ideas. This model is powerful for personal productivity, but it relies on humans to drive every step. Copilots do not track long-running objectives, monitor progress, or act when conditions change. As organisations push AI deeper into operational workflows, these limitations become more visible: copilots assist with tasks, but they do not manage or progress work.

Agentic AI represents the next shift. Instead of responding only to prompts, agentic systems operate around goals. A goal sets the outcome to achieve—not simply what content to generate. The system then plans steps, uses tools and APIs, monitors signals, and adapts within boundaries set by humans. This transforms the AI from a passive assistant into a digital worker. For example, in IT operations, a

copilot helps engineers analyse logs or draft incident reports, but an agent monitors alerts continuously, gathers diagnostics, applies approved fixes, updates records, and escalates only when thresholds or confidence levels require human intervention.

This change depends on three capabilities: goal awareness, controlled autonomy, and persistence. The system must understand what success looks like, act independently within defined guardrails, and maintain context across steps rather than resetting after each interaction. When these come together, the AI coordinates work across systems, reacts to changing conditions, and maintains continuity over time. Humans shift from micromanaging steps to supervising outcomes, benefiting from reduced cognitive load, greater consistency, and faster execution.

Agentic AI works especially well in operationally intensive functions—IT service management, finance operations, customer support, and internal workflows—where goals recur and processes are well structured. It does not remove human accountability; humans still define goals, manage exceptions, and set boundaries. The real risk lies not in autonomy but in poorly designed workflows. Strong governance turns autonomy into leverage; unclear processes turn it into confusion. In essence, copilots help with tasks, but goal-driven digital workers take responsibility for progressing work. Used responsibly, they allow organisations to scale execution, reduce friction, and redirect human judgement to the areas where it matters most.

Video 8: Final Thoughts

Generative AI is often first experienced as a tool that makes work quicker or easier, but its impact extends far beyond individual tasks. When viewed through a value-chain lens, AI emerges as a structural force that reshapes how organisations create value, coordinate work, and deliver outcomes. Its role becomes long-term and strategic rather than simply operational.

Across this week, we explored generative AI as a value-chain transformer, shifting focus from task-level productivity to end-to-end impact. We examined how its influence plays out across customer-facing functions, operations, finance, and HR, where a consistent pattern emerged: AI reduces friction, accelerates cycles, and strengthens consistency, while human judgement remains central to oversight, strategy, and decision-making. Concepts such as workflow compression and cognitive leverage highlighted why AI can deliver disproportionate benefits by reducing handoffs, supporting parallel work, and scaling expert thinking.

We then looked at the shift from copilots to goal-driven digital workers, showing the limits of prompt-based assistance and the rise of agentic AI. These systems operate at the workflow level, coordinating actions, monitoring progress, and supporting execution within clear guardrails. This distinction between assistance and partial autonomy helps organisations adopt AI responsibly—allowing AI to progress work while ensuring humans remain accountable and informed.

Finally, we examined AI-native products and services, which place AI at the core of the offering rather than using it as an operational enhancement. This marks a shift in how organisations compete and how products evolve. As you conclude this week, the key message is not to apply AI everywhere, but to apply it purposefully. Real value comes from embedding generative and agentic AI thoughtfully into

value chains, workflows, and offerings. Organisations that succeed will combine technological capability with clear goals, strong governance, and human-centred design.