

Video Transcripts

Week 1.3: Real-World Cases & Strategic Outlook

Video 1: Why Real-World Cases Matter for AI Strategy

Artificial intelligence is often described in terms of limitless potential—promises of higher productivity, competitive advantage, and transformation at scale. However, inside real organisations, the story is far more uneven. Some AI initiatives generate measurable value, while others stall, remain stuck as pilots, or gradually fade away. This gap between AI's theoretical promise and its actual business impact is one of the greatest challenges leaders face today. That is precisely why real-world cases are so important when shaping an effective AI strategy.

In theory, AI models can analyse vast datasets, detect patterns, and automate decisions in milliseconds. In practice, however, the value created by AI depends far less on algorithmic sophistication and far more on how it is embedded within organisational workflows, governed by leadership, and aligned with business priorities. Real-world cases reveal how decisions are made under real constraints—legacy systems, regulatory oversight, limited talent capacity, and constant pressure to deliver results quickly. They show where organisations struggle and why, in many cases, value fails to materialise.

A powerful example can be seen in India's digital payments ecosystem, particularly the Unified Payments Interface (UPI). UPI has grown at extraordinary speed. In FY24, total UPI transactions reached 13,113 crore, representing a 57% year-on-year increase from FY23. With this rapid expansion came a sharp rise in fraud. Fraud cases increased by 85% in FY24, growing from 7.25 lakh cases in FY23 to 13.42 lakh in FY24. The value of fraud also rose significantly, from ₹573 crore to ₹1,087 crore. Even in the early months of FY24–25, more than 6.32 lakh fraud incidents worth ₹485 crore had already been reported.

These figures represent more than statistics—they highlight operational challenges that directly threaten user trust. As transaction volumes surged, early rule-based fraud detection systems struggled to keep pace. Suspicious transactions accumulated in manual review queues, analysts spent hours examining cases individually, and operational throughput slowed considerably. The growing scale of digital payments increased both operational risk and customer friction, creating an urgent need for more sophisticated solutions.

AI-enabled fraud detection systems began to address this challenge. The National Payments Corporation of India, the Reserve Bank of India, and member banks introduced AI and machine learning systems capable of analysing each transaction in real time. These systems evaluate behavioural signals such as transaction history, device identity, usage patterns, and risk scores to detect anomalies instantly. AI-powered fraud monitoring now allows suspicious transactions to be flagged automatically or declined before completion. At the same time, regulators introduced the Central Payment Fraud Information Registry to track fraud cases using AI-driven analytics across the payment ecosystem.

However, the technology itself was not the sole solution. The real impact came from leadership decisions around governance, workflow redesign, and ecosystem coordination. Fraud detection became a shared responsibility involving banks, payment applications, regulators, and cybercrime authorities. Organisations defined clear thresholds for automated blocking and human escalation, strengthened authentication measures, and expanded customer reporting mechanisms such as the national cybercrime portal and the 1930 helpline. As a result, response times were reduced from hours to milliseconds, manual workloads declined, and investigators could focus on complex cases rather than routine alerts. This illustrates the central lesson of case-based learning: successful AI implementation depends less on model accuracy and more on organisational design, workflow integration, and strong leadership alignment.

Video 2: Generative AI in Knowledge-Intensive Work

Knowledge-intensive work inside organisations rarely follows a simple or repeatable formula. Strategy teams analyse markets before making recommendations, legal teams interpret dense regulations before advising clients, and research or analytics teams combine data, experience, and judgement to solve complex problems. Much of this work depends on context and expertise built over time. For many years, technology only supported these activities at the margins. Generative AI is now beginning to change that support role in meaningful ways.

Importantly, generative AI does not replace expert thinking. Instead, it reshapes how experts allocate their time. Across many knowledge-intensive roles, its influence appears most clearly in three areas: faster analysis, improved synthesis, and clearer communication. A large portion of expert effort is typically spent gathering information, reviewing documents, synthesising inputs, and drafting initial outputs. While these tasks are necessary, they are not where the greatest value of expertise lies. Generative AI accelerates these stages, allowing professionals to move more quickly towards interpretation, judgement, and decision-making.

This shift can be seen across consulting, legal services, research, policy, analytics, and corporate strategy. These roles share a common challenge: processing large volumes of information and transforming them into meaningful insight. Generative AI supports this process by reading, summarising, and restructuring content at a scale and speed that was previously difficult to achieve. As a result, the economics of expert work begins to change. Professionals can spend less time assembling raw material and more time refining ideas, evaluating risks, and shaping recommendations.

A clear example of this transformation can be found in the legal sector, including in India. Legal work is highly knowledge-intensive, time-bound, and risk-sensitive. According to a 2025 Thomson Reuters survey, 26 per cent of legal organisations are already using generative AI, up from 14 per cent the previous year, and 95 per cent of legal professionals expect it to become central to legal workflows within five years. Leading Indian law firms such as Shardul Amarchand Mangaldas have introduced generative AI tools to support legal research, contract drafting, and due-diligence reviews. Another major firm, Cyril Amarchand Mangaldas, ran structured pilots involving around 380 lawyers to test multiple AI platforms before selecting a firm-wide solution focused on reducing time spent on first drafts, document review, and repetitive research.

The results demonstrate measurable efficiency gains. Legal process outsourcing firms in India using hybrid human-AI models report turnaround-time reductions of up to 60 per cent for research and document-heavy tasks, while maintaining final accountability with experienced legal professionals. Globally, studies indicate that lawyers using generative AI complete document review and contract analysis between 20 and 50 per cent faster than traditional approaches. In these organisations, the technology is intentionally positioned as an assistant rather than a decision-maker. AI accelerates scale and consistency, while human professionals remain responsible for interpretation, risk assessment, and final advice.

Across knowledge-intensive work, generative AI influences three important dimensions. Throughput increases because professionals can handle more work without proportional increases in effort. Quality can improve as more time is devoted to refining insights instead of gathering material. Cycle times also shorten, allowing organisations to move faster from question to recommendation. However, these benefits only appear when generative AI is integrated thoughtfully into workflows. Without clear guardrails, organisations risk shallow analysis, excessive reliance on AI-generated outputs, and a gradual erosion of expert judgement. Leading organisations address this by defining clear roles: AI supports exploration, drafting, and synthesis, while humans retain responsibility for interpretation, decisions, and accountability.

Video 3: Generative AI at Scale

Most organisations today do not struggle with getting started with generative AI. Pilots, proofs of concept, and small-scale experiments are now common. The real challenge lies in moving beyond curiosity and embedding generative AI into everyday work in a way that delivers sustained business value. This transition is not driven by technology alone; it depends on how adoption unfolds across the organisation. One of the clearest real-world examples of this journey is Microsoft's deployment of Microsoft 365 Copilot, which illustrates how generative AI can move from isolated experimentation to large-scale organisational impact.

Copilot was introduced as part of Microsoft's broader AI strategy in 2023 and, by 2024–25, had been integrated into widely used applications such as Word, Excel, Outlook, and Teams. Crucially, Microsoft did not position Copilot as a separate or standalone tool. Instead, it embedded generative AI directly into the software employees were already using every day. This design decision made adoption far more natural. Early internal experiments saw teams using Copilot to draft emails, summarise meetings, or generate preliminary reports. However, these early pilots produced mixed results—some employees experienced clear benefits, while others rarely used the tools. Curiosity was high, but large-scale impact had yet to emerge.

The shift occurred when organisations began implementing structured deployment. Copilot was integrated into established workflows and supported with governance mechanisms such as enterprise policies on data access, visibility into usage patterns, and clear guidance on when human validation was required. Rather than remaining an open-ended experiment, generative AI became a supported and managed component of everyday work. This structured approach helped organisations move from scattered experimentation to consistent value creation.

The measurable impact of this approach became evident in several deployments. In a major trial across the UK government involving around 20,000 civil servants, participants reported saving an

average of approximately 26 minutes per day using Microsoft 365 Copilot for tasks such as drafting documents, summarising meetings, and preparing reports. Over the course of a year, this equates to roughly 13 working days saved per employee. Microsoft's internal surveys show similar trends: about 70 per cent of early Copilot users reported higher productivity, while 68 per cent said it improved the quality of their work. Users were up to 29 per cent faster at common tasks such as searching, writing, and summarising, and 85 per cent reported reaching a strong first draft more quickly.

Embedding AI in everyday tools fundamentally changes how work flows. Tasks that previously required manual searching, analysis, and drafting can now be assisted or accelerated by AI. As a result, the human role shifts from performing mechanical tasks to validating, interpreting, and improving AI-generated outputs. However, scaling generative AI is not automatic. Adoption often varies across roles and functions. In the UK government study, for instance, about 17 per cent of users reported no measurable time savings, while others experienced significant benefits. This variation highlights the importance of training, role design, and workflow redesign in achieving meaningful impact.

Beyond productivity gains, embedding generative AI into daily work tools has broader strategic implications. AI moves from being an optional add-on to an expected component of how work is performed. Organisations can increase throughput as employees manage more tasks within the same time frame, improve consistency through standardised outputs, and shorten cycle times by generating drafts and summaries rapidly. At the same time, successful organisations invest in governance alongside deployment. Clear policies on data use, transparency, and accountability are essential for building trust—because trust, rather than technology alone, determines whether AI becomes embedded in everyday work or remains an isolated experiment.

Video 4: Early Agentic AI Deployments

Until recently, most enterprise AI systems stopped at providing recommendations. They could analyse data, generate insights, or suggest next steps, but humans still needed to execute the work. Tasks such as creating follow-ups, sending reminders, updating records, and coordinating actions across systems remained manual. This dependence limited how much value AI could truly deliver. Agentic AI begins to change that boundary. These systems are designed to pursue defined goals by coordinating tasks, monitoring progress, and triggering actions across multiple tools. However, in early enterprise deployments, autonomy is carefully constrained, with humans retaining oversight, approval, and accountability.

A practical way to understand this shift can be seen in Indian banking operations, where execution quality directly affects outcomes. Consider loan collections and customer follow-ups in a large bank such as HDFC Bank. Traditionally, once an account entered early delinquency, operations relied heavily on manual coordination. Teams had to monitor account status, draft reminders, send messages through channels like SMS or WhatsApp, update customer relationship management systems, and schedule further follow-ups. Even with modern tools available, people effectively acted as the link connecting multiple systems. As a result, delays, missed follow-ups, and operational inefficiencies were common.

In 2024, HDFC Bank began piloting agentic AI workflows within parts of its retail banking operations. The objective was focused and operational: to ensure consistent and timely follow-up for overdue accounts while escalating complex cases to human officers. The agentic systems monitored account

status in real time, triggered personalised customer outreach using approved templates, automatically updated CRM records, and scheduled additional follow-ups based on predefined rules. When a customer responded or completed a payment, the workflow closed automatically. If a situation involved ambiguity, such as a hardship request or dispute, the system paused and escalated the case to a human officer along with the relevant context.

The impact was measurable. Across hundreds of thousands of overdue accounts each month, routine execution no longer relied on manual coordination. The bank observed a 15–25 per cent reduction in Days Past Due for certain customer segments within six months. More than 60 per cent of reminder messages and CRM updates were handled autonomously, and operations teams reduced the time spent on manual tracking from roughly 70 per cent to under 40 per cent. As a result, staff could focus more on negotiation, resolution, and customer relationships instead of administrative task management.

This example highlights an important pattern in early agentic AI deployments: autonomy remains bounded. Agentic systems operate only within clearly defined scopes and usually under strict confidence thresholds, often above 80 per cent, with explicit escalation rules in place. They do not renegotiate loan terms, alter credit decisions, or override established policies. Human teams retain responsibility and decision-making authority. A similar approach is emerging in telecom service operations, where agentic AI systems monitor ticket queues, detect delays, escalate issues automatically, update stakeholders, and reprioritise tasks according to defined rules. In early deployments, organisations have reported improvements of 20–30 per cent in mean time to resolution, largely because execution became more consistent.

Several lessons emerge from these cases. Agentic AI delivers the greatest value in execution-heavy workflows where the challenge is not deciding what to do, but ensuring that tasks are completed reliably. Trust in these systems depends on transparency, including clear logs, audit trails, and the ability to reverse actions if necessary. Process quality also matters greatly. Agentic AI amplifies existing workflows, so poorly designed processes will become more visible rather than automatically improved. These systems are not yet setting strategy or making complex ethical trade-offs; their autonomy remains conditional and supervised. As organisations move forward, a hybrid model is likely to emerge, where humans define intent, constraints, and judgement, while agentic systems coordinate tasks, monitor progress, and manage follow-through at scale.

Video 5: Strategic Outlook — AI as an Operating Capability

Agentic AI represents a shift from tools that simply help people work faster to systems that begin to shape how work itself moves within organisations. Most technologies automate specific parts of a process while leaving the surrounding organisation largely unchanged. With agentic AI, however, work becomes coordinated, routed, and executed differently. People intervene where judgement and expertise are required, while AI systems manage speed, scale, and orchestration. A clear example of this transition can be seen in the transformation underway at Bharti Airtel, one of India's largest telecommunications providers.

Airtel operates an extensive telecom network where the scale of operations makes purely human-driven monitoring increasingly difficult. As data usage expanded rapidly across 4G and 5G networks,

traditional rule-based operational systems struggled to keep pace. To address this, Airtel launched a multi-year modernisation initiative with Nokia to upgrade its core network infrastructure. By 2025, the company confirmed a strategic shift towards automation, zero-touch workflows, and generative AI-driven service orchestration across many of its service regions in India. The goal was to move gradually towards more autonomous network operations, where service launches and lifecycle management could occur with minimal manual intervention.

At the same time, Airtel expanded its long-term partnership with Google Cloud to integrate advanced AI and machine learning capabilities across its operations. This collaboration focuses on combining telecom data with Google's AI technologies to enable solutions embedded directly into operational workflows. Areas of development include geospatial intelligence based on Airtel's network data, multilingual voice analytics trained across several Indian languages, and generative AI systems designed to improve customer interactions and streamline internal operational processes across mobile, broadband, and digital television services.

Within network operations, these changes are already transforming how work is performed. Instead of engineers being the first to detect problems or open incident tickets, AI-enabled systems continuously monitor network telemetry, identify anomalies, correlate alarms, and trigger operational workflows automatically. Routine issues can often be resolved end-to-end within predefined guardrails. When complex situations arise, the system escalates them to engineers along with detailed contextual information, reducing the time and cognitive effort required to diagnose the problem. Tasks that once required multiple manual steps across dashboards, ticketing systems, and communication channels are increasingly coordinated by AI-driven processes capable of initiating and tracking operational flows.

Importantly, this transition is not happening without oversight. Airtel's approach demonstrates that autonomy grows gradually and is accompanied by strong governance mechanisms. Organisations clearly define which actions AI systems may perform independently, which require human approval, and which must remain fully human-led due to operational or regulatory risk. Automated workflows are supported by logging, observability, and the ability to reverse actions when necessary—critical safeguards in an industry where reliability, security, and compliance are essential. Trust in AI systems is therefore built progressively through transparency and consistent performance rather than through the sudden removal of human control.

The broader strategic implication is that AI is no longer viewed merely as a productivity tool. Instead, it is becoming part of the operational foundation of modern telecom networks. By embedding automation and agentic behaviour into processes such as service provisioning, network assurance, routing, incident response, and customer support, Airtel is redesigning its operational model around AI-enabled execution. Humans continue to define objectives, manage risk, and apply judgement, while AI systems handle coordination, monitoring, and repetitive execution tasks at scale. This shift illustrates how organisations that integrate AI into their operating model can achieve faster cycle times, reduced manual workloads, and more resilient systems—demonstrating that the future of operations lies not in replacing human expertise, but in augmenting it with intelligent systems capable of operating at the scale modern networks require.

Video 6: Implications for Leaders and Organisations

When discussions about artificial intelligence remain focused only on tools, the impact can feel abstract. However, when AI begins to influence how work is assigned, how service issues are resolved, and how accountability flows through an organisation, its implications become very real. Many large organisations are now entering this stage. AI is no longer just improving efficiency; it is reshaping roles, coordination patterns, and leadership responsibilities. A clear example of these organisational implications can be seen at Reliance Jio, one of the world's largest telecom operators, serving more than 450 million subscribers.

At this scale, even small inefficiencies in customer service or network operations can multiply rapidly. Manual coordination alone cannot keep pace with the volume of interactions and operational events. To address this challenge, Jio began embedding AI directly into its operational workflows. Starting around 2018–2019 and expanding further after 2020, the company introduced AI-driven systems across customer service, network monitoring, and service assurance. Today, AI handles a substantial share of first-level customer interactions, including answering queries, creating service tickets, prioritising requests, and routing them to the appropriate teams. At the same time, AI systems continuously monitor network performance, predict potential faults, and trigger corrective actions before customers experience service disruptions.

These changes have significant implications for organisational roles. Tasks that once defined many frontline positions—such as manual ticket handling, status tracking, and routine troubleshooting—are increasingly handled automatically by AI systems. Human employees now focus more on exceptions, complex customer issues, and ensuring successful outcomes. Roles are not eliminated but reshaped, with value shifting from routine execution towards judgement, oversight, and problem-solving. Performance expectations have also evolved. Instead of being measured primarily by the volume of tasks processed, employees are increasingly evaluated based on resolution quality, escalation decisions, and overall customer outcomes.

Managers experience some of the most significant changes. Traditional service management relied heavily on supervision: monitoring queues, chasing unresolved issues, and redistributing work. With AI systems automatically prioritising tasks, triggering follow-ups, and managing escalations, much of this coordination now happens without human prompting. Managers therefore spend less time on operational firefighting and more time defining priorities, setting operational thresholds, and handling unusual or high-risk cases. This shift requires new managerial capabilities, including understanding how AI-driven workflows function, recognising when AI outputs can be trusted, and knowing when human intervention is necessary.

At Jio, coordination between humans and AI systems is deliberately structured. AI operates within clearly defined boundaries. When the system has high confidence in its actions, it proceeds automatically. When uncertainty increases—such as in unusual complaints, regulatory issues, or situations involving reputational risk—the case is escalated to a human decision-maker. Responsibility for outcomes remains firmly with people, even when execution is largely AI-driven. At the organisational level, these workflows also encourage stronger collaboration across departments, since AI-enabled processes often cut across customer service, network operations, and analytics functions.

Jio's experience highlights an important lesson for leaders. The most significant implications of AI are organisational rather than purely technical. Roles shift from execution to oversight, managers move from supervision to judgement, and accountability becomes clearer rather than weaker. These

changes usually occur gradually. Jio introduced AI step by step, beginning with high-volume and well-defined workflows. Trust was built through transparency, logging, and clear escalation procedures before expanding the scope of automation. Ultimately, the challenge for leaders is not simply deploying AI technology but redesigning the organisation around it—ensuring that efficiency gains are balanced with trust, governance, and human judgement at the core of decision-making.

Video 7: A Simple Lens for AI Opportunity Identification

By this stage, AI use cases can begin to feel overwhelming. Examples appear across content generation, analytics, customer service, and operations, making it difficult for leaders to decide where to begin. The real challenge is not recognising that AI can be useful, but identifying opportunities that genuinely matter within the organisation. With limited time, attention, and investment capacity, leaders need a practical way to prioritise where AI can create the most meaningful impact.

A helpful approach is to look less at the technology itself and more closely at how work happens today. High-impact AI opportunities usually emerge from workflows that are under constant pressure—where complexity, repetition, and time constraints intersect. These environments create friction for employees and often limit the speed and consistency of operations. A strong example of this dynamic can be found at DHL, one of the world’s largest logistics companies, which manages millions of shipments across regions, time zones, and regulatory systems every day.

Before introducing AI, DHL’s customer service agents often had to search across multiple internal systems to reconstruct the status of a shipment, identify exceptions, and determine the appropriate response. Each customer inquiry required piecing together information from different sources, which slowed response times and created inconsistency during periods of high demand. DHL recognised this workflow as a critical pressure point—not because it was particularly glamorous, but because it was essential to service quality and operational reliability.

To address this challenge, DHL introduced generative AI into customer service and operational support workflows. AI systems now summarise shipment histories, combine data from multiple systems, and generate response suggestions for agents in real time. Importantly, human employees still retain responsibility for final decisions and communication with customers. The purpose of AI is not to remove human judgement but to reduce the cognitive burden of assembling and interpreting large volumes of information.

This example highlights a simple lens that leaders can apply when identifying AI opportunities. The first element is task complexity. DHL’s customer interactions involve context, interpretation, and the synthesis of fragmented data rather than simple rule-based decisions. These are precisely the conditions where generative AI can provide significant support by organising information quickly and clearly. The second element is volume and frequency. DHL processes customer inquiries continuously at massive scale, meaning that even small time savings per interaction can accumulate into major operational improvements.

The third element is latency sensitivity. In logistics operations, delays can lead to customer dissatisfaction, missed coordination between partners, and cascading operational disruptions. By shortening the time required to move from question to answer, AI improves both efficiency and

service quality. When complexity, high volume, and time sensitivity appear together, AI moves beyond being a productivity tool and becomes a way to transform how workflows operate.

This framework can be applied across many business functions. Financial teams manage forecasting and reporting under strict deadlines, marketing teams combine creative work with fast execution cycles, human resource teams handle large volumes of talent screening and training materials, and operational teams coordinate activities across growing systems and partners. At the same time, the framework also helps leaders rule out weaker opportunities. Tasks that are simple, infrequent, and tolerant of delay rarely produce meaningful strategic returns from AI. The goal is not to apply AI everywhere, but to focus on the areas where it fundamentally changes how work flows through the organisation.

Video 8: Final Thoughts

Across industries, one pattern has become increasingly clear: the value of AI is not created by capability alone, but by how organisations choose to apply, govern, and integrate it into real work. The examples explored throughout this week show that successful AI adoption depends less on technological breakthroughs and more on the operating decisions made by leaders.

A recurring theme is the gap between AI potential and realised business value. Many organisations experiment enthusiastically, yet only a smaller group manage to translate pilots into sustained impact. The difference often lies in clarity of purpose. When AI initiatives are anchored to genuine business problems and embedded into workflows, value compounds over time. When they remain isolated tools, the benefits stay limited. As discussed earlier, generative AI is already reshaping expert and legal work by accelerating research and document-heavy tasks by 20–60 per cent, enabling professionals to work faster without replacing their expertise.

The role of generative AI in knowledge-intensive work highlights this shift clearly. Rather than replacing expertise, generative systems act as force multipliers. They accelerate research, synthesis, and communication, allowing professionals to spend more time on judgement, interpretation, and strategic decision-making. In this context, productivity gains come not simply from automation, but from enabling people to focus their time and attention on higher-value thinking.

Scaling AI reinforces the same lesson. Early success with pilots does not automatically translate into enterprise-wide impact. Organisations that successfully scale AI typically redesign processes, define clear ownership, and establish governance structures early in the journey. Trust, reliability, and integration into everyday tools ultimately determine whether AI becomes part of how work is done or remains an optional add-on used only occasionally.

As AI begins to take on more agentic characteristics, the nature of value evolves further. Early deployments already show AI coordinating tasks, tracking progress, and handling routine follow-ups, although autonomy remains deliberately limited. For example, HDFC Bank implemented agentic AI to automate most routine loan follow-ups, reducing delays by 15–25 per cent and lowering manual workload so that staff could concentrate on more complex cases. In such systems, humans still define the goals, constraints, and escalation rules, ensuring that risk remains controlled while efficiency and consistency improve.

The risks discussed throughout the week reinforce an important insight: AI governance is fundamentally a leadership responsibility. Reliability, accountability, and oversight shape trust in AI systems. Even when AI contributes to outcomes, organisations remain responsible for the decisions made and their consequences. Effective governance therefore enables responsible scaling rather than restricting innovation.

Looking ahead, AI is best understood as an emerging operating capability. Over time it reshapes roles, processes, and decision rights across organisations. Managers shift from overseeing individual tasks to guiding outcomes, while teams increasingly coordinate their work with AI systems as part of everyday operations. This transition is gradual, but its long-term impact on organisational structure and productivity is significant.

To bring these ideas together, leaders can apply a simple lens to identify meaningful AI opportunities. The example of DHL illustrates this clearly. With millions of shipments handled daily, teams constantly interpret data and resolve issues in a high-volume, time-sensitive environment. Leaders can ask: where does AI reduce friction in knowledge work, where does it compress the time between insight and action, and where can carefully constrained autonomy improve consistency without increasing risk?

These questions shift attention away from experimentation alone and towards strategic intent. They help organisations move beyond treating AI as a standalone tool and instead begin building it into how work actually happens.