

Video Transcripts

Week 3.1: Prompting and Applied Techniques

Video 1: Why Prompting Matters in Business Contexts

Welcome to this week on Foundations of Effective and Structured Prompting. Before we dive into techniques and frameworks, let's take a moment to understand why prompting matters, especially in a business context.

When you interact with a Generative AI system, you are not clicking buttons or filling out forms. You are communicating through language. The prompt is the interface. What you write, how you frame it, how specific you are, and what context you provide directly determines the quality of what comes back. In that sense, prompting is not a technical skill. It is a communication skill. And like any communication skill, it can be learned, structured, and improved.

Think about a marketing manager who needs to draft a campaign brief. If they type 'Write me a campaign idea,' the AI will produce something, but it will be generic, untargeted, and disconnected from the brand's tone, audience, or objectives. Now imagine the same manager writes: 'You are a senior brand strategist. Draft a campaign concept for a mid-market B2B SaaS product targeting HR managers in the healthcare sector. The tone should be professional but approachable, and the campaign goal is lead generation for Q3.' The output from that second prompt is in a completely different league, not because the AI changed, but because the instruction did.

This is why organisations are moving towards systematic prompting. Instead of relying on ad hoc instructions, teams are beginning to adopt structured approaches where prompts clearly define context, objective, constraints, and expected format.

For instance, a customer support team may design a prompt that instructs the AI to respond in a professional tone, reference company policy, and provide step-by-step solutions. Similarly, in e-commerce, an agentic AI system may use structured prompts to analyse customer behaviour and recommend personalised products based on browsing history, purchase patterns, and current session signals. These structured prompts ensure consistency, relevance, and alignment with business goals.

As businesses increasingly adopt agentic systems, prompting also becomes dynamic. Agents interpret signals, trigger actions, and generate outputs continuously. In such systems, well-defined prompts enable coordination between multiple agents, ensuring that each action aligns with the customer journey and organisational objectives. Poor prompting in this context can break the flow of interactions, while effective prompting enables seamless, real-time decision-making at scale. This is the shift this module is about: moving from casual, instinctive prompting to systematic, intentional prompting. In a business environment, where outputs carry real consequences, that shift is not optional. It is essential.

In the videos ahead, we will explore exactly how to make that shift, with frameworks, structures, and practices you can apply immediately in your own work.

Video 2: Core Prompting Principles

Now that we understand why prompting matters, it is important to recognise a simple market reality: Most organisations experimenting with Generative AI are not failing because of the technology. They are underperforming because of how they prompt it.

Across functions – consulting, operations, marketing, finance – the difference between average and high-value output often comes down to a few core prompting principles. These are not rigid rules, but disciplined habits that consistently improve outcomes.

The first principle is specificity. In business environments, vague prompts translate directly into rework. A generic instruction like ‘Summarise this report’ often produces an output that still requires editing. In contrast, a precise prompt such as ‘Summarise this report in three sentences, focusing on key financial risks in Q2 and their potential business impact’ produces something closer to decision-ready. The effort shifts from rewriting to refining.

The second principle is context. AI systems do not understand your organisation, your customer, or your constraints unless you tell them. In practice, this is where most enterprise use cases break down. For example, a customer success manager drafting a response without context may receive a generic email. But if the prompt includes the client’s industry, the issue raised, contract value, and relationship sensitivity, the output becomes far more aligned to the situation. Context reduces guesswork – and guesswork is costly in business settings.

The third principle is role and persona. In the market, teams are increasingly using AI not just as a tool, but as a simulated expert. Consulting firms, for instance, often prompt AI to act as a senior strategy adviser when reviewing proposals. The same input, when framed through a role, produces more structured, critical, and commercially relevant output. The role you assign determines the lens through which the problem is analysed.

The fourth principle is output format. In enterprise workflows, outputs rarely stand alone – they flow into slides, reports, dashboards, or client deliverables. If you do not specify the format, you add an extra step of conversion. Asking for a table, a slide-ready bullet list, or a structured framework ensures the output is immediately usable, saving time at scale across teams.

To see how these principles come together, consider a real-world scenario. A procurement head in a mid-sized manufacturing firm is reviewing vendor performance after a year of supply disruptions. A weak prompt might be: ‘Analyse vendor performance.’ This typically results in a generic response that lacks actionable depth.

Now compare that with a structured prompt: ‘You are a procurement analyst for a manufacturing firm with 50+ vendors across logistics, raw materials, and packaging. Analyse vendor performance data from the past 12 months. Identify the top five risk areas, recurring performance issues, and vendors requiring review. Present the output in a table with vendor category, issue, impact level, and recommended action.’ The difference is not just in clarity – it is in decision readiness. The second prompt produces output that can directly inform a review meeting or a management report.

This is the shift organisations are making in practice. Generative AI is not just about generating content – it is about improving the quality and speed of decisions within existing workflows. These four principles – specificity, context, role, and output format – form the foundation for that shift. When applied consistently, they turn AI from a generic assistant into a reliable business tool aligned with organisational needs.

Video 3: Structured Prompt Anatomy

In the previous session, we explored the core principles of effective prompting. In practice, however, organisations face a different challenge: ensuring consistency when multiple teams begin using Generative AI across functions. Marketing, operations, finance, and customer support teams may all use the same tools, yet produce very different outcomes. The variation is rarely due to the technology itself, but to how prompts are structured.

To address this, organisations are increasingly adopting a structured prompt anatomy – a repeatable format that standardises interaction with AI systems. This approach improves clarity, reduces ambiguity, and ensures that outputs are more aligned with business needs. A widely used structure follows five components: role, task, context, constraints, and output.

The first component is role. This defines who the AI should act as and sets the perspective for the response. In business settings, this is critical because different roles approach problems differently. For example, a customer support specialist prioritises clarity and empathy, while a risk analyst focuses on precision and caution. Assigning a role helps shape both tone and depth.

The second component is task. This specifies exactly what needs to be done. In many enterprise use cases, vague tasks lead to generic outputs that require additional editing. A clearly defined task – such as resolving a billing issue or analysing vendor performance – provides direction and ensures the response is relevant. Role and task together establish a strong foundation for the prompt.

The third component is context, which is often where the real business value is created. AI systems do not have access to organisational knowledge unless it is explicitly provided. Context may include customer history, transaction data, operational constraints, or recent interactions. When context is included, the AI can generate responses that are not only accurate but also tailored to the situation.

The fourth component is constraints. In real-world environments, outputs must adhere to specific boundaries such as tone guidelines, word limits, or regulatory requirements. For instance, financial services firms may require outputs to avoid speculative language and comply with strict communication standards. Constraints ensure that responses are both usable and aligned with organisational policies.

The fifth component is output format. In enterprise workflows, AI-generated content is rarely the final step; it feeds into presentations, reports, emails, or dashboards. Specifying the format – such as a table, bullet points, or a structured summary – makes the output immediately usable and reduces the need for reformatting. This is especially valuable when teams are working under time constraints.

Consider a practical scenario. A marketing team in a retail organisation is designing a re-engagement campaign for customers who have not made a purchase in the past 90 days. A simple prompt might ask for a generic message, but a structured prompt defines the role as a digital marketing strategist, specifies the task, includes customer behaviour as context, sets tone and word limits as constraints, and requests a clear output format. The result is not just better content, but content that is directly usable within the campaign.

By applying this structured prompt anatomy, organisations can move from inconsistent experimentation to reliable execution. It enables teams to produce higher-quality outputs, reduces rework, and supports the scaling of AI across business functions.

Video 4: Delimiters and Formats

When one prompts to an LLM, there are many tools that are useful. One of the tools here is called delimiters. Delimiters are like punctuation marks for LLMs.

Several punctuation marks improve the prompt output. In this demo, we will try to understand how it improves the output and what the major punctuation marks are. Some of the punctuation marks or delimiters for LLMs are triple quotes, curly braces, square brackets, hashtags, dashes, labels, colons, and so on.

Let us demonstrate those delimiters on an LLM called Grok. I have a prompt pasted like this: 'Summarise these reviews and tell me if they are positive or negative. Review 1: product arrived late, but it works ok, I am not thrilled but not mad. Review 2: Absolutely amazing, best purchase of the year. Review 3: Total waste of money, it broke after 2 days, do not buy.' This prompt is not free of any delimiter.

The colon is one of the delimiters mentioned here. So let us see the output, how Grok gives an output to this kind of prompt. So, one can see the output here: review 1, sentiment mildly neutral, review 2, review 3, and so on.

An overall summary has been given. Now I have a much more structured prompt with the use of a delimiter. You are a customer success analyst, your task is to analyse customer reviews below within review tags.

So now this is a delimiter, instructions. This colon is a delimiter. 1, 2, 3 give a hint for the delimiter. These quotes are delimiters. So let us see how this prompt is being answered. Now you can see the answer given by the LLM versus the LLM without delimiters has a stark difference. It is much more structured, a little more detailed, and gives key takeaways for the product team as an expert.

One thing to note is that in the new prompt the role is also added. So, in this demo, we have seen that delimiters will enhance the prompt structure, both input and output. Using those in your prompt engineering, in your prompt exercises, will enhance the output and you need not repeat and ask the LLM again and again to give you the output of your choice.

This stops LLM prompt chaining, and you get an output very quickly and to the point. Apart from the usage of delimiters, it is important to know the data formats with which LLMs are most comfortable. The data formats that we are most comfortable with are tabular format data or data in Excel files, data in printed format, and so on.

Usable formats in machine-readable form are xls, csv, tabular format, and so on, that we have already seen. But specific to large language models, there are different types of formats that are used. It is important to know what these kinds of formats are.

Popular formats are text (.txt), csv format. Apart from the delimiters that we have discussed, another important thing to know is the format in which LLMs understand the data. As humans, we understand data in printed format, in audio format, in speech format, and in tabular format. Machines understand data in different formats. For example, audio data is understandable in mp3 format. Text data is understandable in text format, csv format, xls, xlsx format.

Similarly, LLMs understand specific types of formats more. What are those kinds of formats? Popular formats are csv format, json format. Json format is written between curly braces, and double quotes and colons are used to separate the column name and the column values. Similarly, jsonl format is a variant of the json format.

Another format is YAML format. YAML format does not have curly braces, but it has colons and dashes to separate between lists of values. Some efficient storage formats are also nowadays very popular across LLMs.

It is called the parquet format. Parquet formats are not human-readable, but they are column binary format and their schema contains string integer. Apart from the parquet format, there is a pickle format, which looks similar to json, but has subtle differences. Then we have Apache Arrow format, and we have specific formats related to NumPy and PyTorch, tensors also.

Video 5: Zero-Shot Prompting

Imagine a customer operations manager in an e-commerce company during a peak sales period. Orders are high, delivery timelines are tight, and customer queries are increasing by the hour. In the middle of this, a customer raises a complaint about a delayed shipment and requests an update. The manager turns to a Generative AI tool and types a simple instruction: 'Write a professional email to a customer apologising for a delayed delivery and providing reassurance.'

Within seconds, the AI generates a clear, polite, and usable response. The tone is appropriate, the structure is correct, and the message can be sent with minimal editing. No examples were provided. No additional guidance was given. The AI relied entirely on its existing knowledge of how such emails are typically written.

This is zero-shot prompting in practice. A direct instruction is given, and the AI produces a response based on what it has already learned. For routine, well-understood tasks like drafting emails, summarising content, or generating basic marketing copy, this approach is both fast and efficient.

Now consider a second situation involving the same manager. At the end of the week, she is asked to review customer complaints and identify patterns that may indicate operational issues. She enters a prompt: ‘Analyse customer complaints from the past week.’

This time, the output is far less useful. The response is broad, somewhat generic, and lacks clear direction. It does not prioritise issues, quantify impact, or highlight specific problem areas. While the AI attempts to respond, it does not fully align with what the manager needs for decision-making.

The difference between the two outcomes is not the tool – it is the nature of the task. In the first case, the task was standard and clearly defined. In the second, it was ambiguous and required deeper interpretation, structure, and context. Without additional guidance, the AI defaults to general patterns, which may not meet business requirements.

This highlights the core limitation of zero-shot prompting. It works well when expectations are clear and the task is common. However, as soon as the task becomes more complex – requiring prioritisation, structured analysis, or business-specific context – the lack of guidance becomes a constraint.

In real organisational settings, zero-shot prompting is often used as a starting point. Teams rely on it for quick outputs, early drafts, or simple interactions. But as workflows become more critical and decisions more consequential, prompts need to evolve. Additional context, constraints, or examples are introduced to improve accuracy and relevance.

The key, therefore, is not just knowing how to use zero-shot prompting, but knowing when it is sufficient. It enables speed and efficiency for routine tasks, but must be complemented with more structured approaches when precision and reliability are required.

Video 6: Few-Shot Prompting – Concept

Consider a content lead at a fast-growing retail brand managing hundreds of product listings across an e-commerce platform. The brand has a very specific voice—premium, minimal, and benefit-driven. Every product description must follow this tone because it directly impacts conversion.

Initially, the team tries using Generative AI with a simple prompt: ‘Write a product description for a leather backpack.’ The output is usable, but inconsistent. Some descriptions are too generic, others too promotional, and many do not match the brand’s tone. Even when the prompt is structured with role, task, and constraints, the results still vary. The team spends significant time editing to maintain consistency.

To solve this, the content lead changes the approach. Instead of only giving instructions, she provides two existing product descriptions that the brand has already approved. The prompt now includes: ‘Here are two examples of our product descriptions,’ followed by the samples, and then, ‘Write a product description for a leather backpack in the same style.’

The difference is immediate. The AI begins to mirror the sentence structure, tone, vocabulary, and level of detail seen in the examples. The output is not just correct – it is on-brand and consistent.

Editing effort drops significantly, and the team can now scale content production across hundreds of products.

This is few-shot prompting in practice. Instead of relying only on general instructions, the AI is guided using a small number of examples that demonstrate the expected pattern. These examples act as a reference point, helping the model understand not just what to do, but how to do it.

Now consider a second scenario in customer support. A company wants all responses to follow a specific structure: acknowledge the issue, provide resolution steps, and close with reassurance. Simply instructing the AI to 'respond professionally' leads to variation. However, when the prompt includes two or three sample customer queries and ideal responses, the AI begins to replicate that structure reliably across new cases.

This highlights where few-shot prompting becomes most valuable. It is particularly effective when the task requires a consistent format, tone, or reasoning style – something that is difficult to define through instructions alone. The examples remove ambiguity and reduce interpretation gaps.

In real business environments, this approach is widely used in marketing, customer service, and content operations – especially where brand voice and consistency are critical at scale. It allows teams to maintain quality while significantly increasing output volume.

Compared to zero-shot prompting, few-shot prompting offers greater control and reliability. It reduces variability, improves alignment with expectations, and ensures that outputs are not just usable, but consistent across use cases.

Video 7: Few Shot – Sentiment Analysis

In this demo, we will see how few-shot prompting dramatically improves the sentiment classification accuracy. Till now, we have understood or gone through the examples of delimiters. Also, we have gone through the output formats, how LLM understands different formats. The prompting that we have done till now was zero-shot prompting.

And it is interesting to note that zero-shot prompting will have a follow-up in one-shot prompting and many-shot prompting. The zero here implies that you are plainly giving instructions to the LLM to answer the query. While in one-shot prompting, you give one example and one output.

That means you give a label to data. In few-shot prompting, you give multiple examples and outputs and ask the LLM to answer your instruction based on these examples or, in other words, learn from the examples that you have given, and then you give new data. In this way, we will improve the output of the LLM and also attune it to our specific standards.

Let us now look at a demo of few-shot prompting. It will dramatically improve sentiment classification. So, we are taking one business use case in customer feedback.

I have opened ChatGPT and inputted my few-shot prompt. As you can see, the instruction is clear and is delimited by hashtags. Classify the sentiment of the following customer feedback as positive,

negative, or mixed. Use the exact same classification style and criteria shown in the example below. This is the example of few-shot. This is your first-shot example.

Example feedback: The product works exactly as described, and delivery was superfast. The sentiment is positive. The feedback: The app keeps crashing every time I try to log in. Terrible experience. This is negative. And the product quality is good, but the customer support response is very slow. This is mixed. Now, the new data is added.

The delivery was on time, but the item has a small scratch, and the colour looks different from the website images. And now, we are asking the LLM to classify this output as positive, negative, or mixed based on these three different examples and their outputs given to it. So, let us see.

It gives a simple output that the sentiment is mixed. This kind of prompting becomes extremely important when you try to use GPT for different machine learning tasks. In this case, classification. It will not give you long sentences, verbose text, and so on. Your point-to-point reply will be provided. We have seen from this example that few-shot prompting exactly understands the tone, and it is very precise with the answer, and your answer also matches because you provide several examples to learn from.

This may not be very appealing on a web-based interface, but it becomes extremely crucial when you try to download an open model or when you try to use an API for a closed model to develop specific business use cases for your internal as well as external clients to develop a sentiment classification algorithm.

Unlike before the LLM advent, these kinds of tasks would be trained on data. You needed to do several iterations of this classification algorithm, but LLMs can shorten that life cycle and give you precise and reliable sentiment classification.

Video 8: Few-Shot vs Zero-Shot Prompting

The prompting that usually people try is a very basic prompt, and due to some reasons we call that a zero-shot prompt. It is curious to understand why it is zero shot because there must be something called one shot, two shots, three shots, many shots, and so on. Yes, we have a terminology called few-shot prompting.

So, in this demo, we will try to explore the difference between few-shot prompting and zero-shot prompting. In zero-shot prompting, you give a prompt, you structure it using delimiters, you structure it using role and other things. But in few-shot, one-shot, two-shot prompting, we try to give them candidate examples and their answers so that the LLM output becomes more consistent with whatever our specification is.

So, let us go to an LLM model web interface and try to see the difference. I have opened ChatGPT, and I will put a prompt. The prompt is having delimiters. The instruction is 'Extract key information from the following text in clean JSON format.' The text is given to be this, and ChatGPT will extract the information in the JSON format. So, this is the zero-shot prompt because an example is not given.

Now I have pasted a few-shot prompt. Let us see the difference between the prompt, the zero-shot versus the few-shot. You can see the instructions are almost the same, and the instruction follows the exact structure in the example below.

Now this is one shot. The text example is given: 'Amit Patel is 28 years old,' and so on, and the output, the sample output, how the LLM should respond, as an example, it is given. Another shot. So, if the prompt is only stopped here, then it will be called one shot, but we have two shots. We have a Sneha Reddy marketing team lead, and its output should be given like this. Now we have here our new text.

So, based on this one positive example, negative example, missing data example, we want the output in a JSON format, and now we have seen the difference between one-shot and few-shot prompt. Now we will see the difference between their output. So, I will hit the enter button on ChatGPT.

It is generating the output in this manner. So, this is the output of the few-shot prompt, and this is the output of the zero-shot prompt. You can clearly see the difference that this few-shot prompt is much more structured and it is aligning with the structure that we are given. Here there are six entries in the JSON format, and here there are different entries. Some of the entries, some of the column names, are not desirable or we do not want. We want it in this format.

Therefore, this shows that few-shot prompting is a much more improved version of zero-shot prompting. So, in this demo, we have seen the difference between the few-shot prompt as an input, how few-shot prompt is structured vis-a-vis how zero-shot prompt is structured. Also, the difference between the output of the zero-shot prompt and the few-shot prompt. We get a much more structured, aligned output in this few-shot prompt as compared to the zero-shot prompt, even in the frontier models like ChatGPT

Video 9: Recap and Prompting Cheat Sheet

You have now covered the core building blocks of effective prompting. Before you move on, let me bring everything together into a quick recap you can refer back to whenever you need it. Prompting is a foundational skill that determines how effectively organisations can use generative and agentic AI systems. The key principles remain consistent across use cases. Prompts must be specific, include relevant context, define a clear role or persona, and specify the expected output format. These elements reduce ambiguity and ensure that the AI generates responses aligned with business objectives. For example, instead of asking an AI tool to 'create a report,' a structured prompt that specifies the role of a business analyst, defines the task as analysing quarterly sales, includes regional data context, and requests a concise summary with insights will produce a far more useful output.

A structured template helps standardise this approach. The sequence of role, task, context, constraints, and output ensures that no critical element is missed. In a customer support use case, a prompt may define the role as a support specialist, assign the task of resolving a delivery issue, include customer history as context, set constraints such as a polite tone and concise explanation, and specify the output as a short paragraph followed by next steps. This structure improves consistency across interactions and allows teams to scale AI usage effectively.

Few-shot prompting builds on this by adding examples. When tasks require a specific tone, classification logic, or format, providing two or three examples can significantly improve output quality. For instance, a marketing team can include sample product descriptions to guide the AI in maintaining brand voice. Similarly, a data analysis task can include labelled examples to ensure consistent categorisation. These examples act as a reference, helping the model replicate the expected pattern.

A practical approach is to begin with zero-shot prompting for simple and well-defined tasks. If the output lacks precision or consistency, additional structure and few-shot examples can be introduced. This step-by-step progression balances efficiency with control and prevents unnecessary complexity in prompt design.

The next step is application. Choose one real business task such as drafting a customer email, generating a product recommendation, or summarising a report. Apply the structured prompt template and refine it by adding examples if required. This hands-on practice will help translate prompting concepts into practical capability, enabling more reliable and impactful use of AI systems in everyday workflows.

Video 10: Final Thoughts

Prompting has emerged as a core capability in working with generative and agentic AI systems. Throughout this week, the focus has been on building a structured and reliable approach to prompting. The journey began with understanding why prompting matters in business contexts, where clarity directly impacts output quality, consistency, and decision-making. The discussion then moved to core principles such as specificity, context, role definition, and output structure, which form the foundation of effective prompts.

Structured prompt anatomy introduced a repeatable framework that ensures completeness and reduces ambiguity. Concepts such as delimiters, instruction design, and output control further strengthened the ability to guide AI systems with precision. Zero-shot prompting provided a starting point for simple tasks, while the concept of few-shot prompting demonstrated how examples can improve accuracy and consistency in more complex scenarios. The comparison between these approaches highlighted when to prioritise speed and when to prioritise control.

The key takeaway is that prompting is not a one-time activity but an iterative process, as we learned in all our demonstrations. It requires refinement, testing, and alignment with business objectives. As organisations adopt agentic systems, structured prompting enables better coordination, scalability, and real-time responsiveness.

The next step is application. Start using structured prompts in everyday tasks, observe the outputs, and refine continuously. This practice will help build confidence and ensure that AI systems deliver meaningful and reliable results across business functions.