

Video Transcripts

Week 5.2: Disruptive Innovations of Gen AI in the Realm of Finance, HR, Strategy and Leadership

Video 1: Subscription to Intelligence Models

Hello learners, and welcome back. In the previous week, we examined how Generative AI acts as a disruptive force – reshaping industries, redefining value creation, and blurring traditional boundaries across functions.

In this week, we extend that discussion to the organisational and strategic level. We will explore how these disruptions impact finance models, talent structures, and long-term competitive advantage, and what it means to operate as an AI-native enterprise.

What if instead of buying software tools, organisations could simply subscribe to better decisions? Traditionally, organisations purchased software licences. They installed tools, collected data, and internally generated insights. The software provided the capability, but the analytical work remained with the organisation.

Generative AI is changing this model. Instead of selling the tool, many companies now sell the insights produced by the AI system itself. Organisations increasingly subscribe to continuously generated intelligence rather than owning the underlying technology.

A strong example comes from AlphaSense. Investment analysts traditionally spent hours reviewing earnings transcripts, analyst reports, and regulatory filings. AlphaSense applies AI to analyse millions of documents and automatically identify themes, risks, and sentiment changes across companies and industries.

Clients subscribe not simply to access documents, but to receive interpreted signals and insights that accelerate investment research.

Another example appears in industrial operations. C3 AI provides predictive maintenance intelligence for sectors such as energy, manufacturing, and aviation. Its AI models analyse sensor data from equipment and identify early indicators of mechanical failure.

Instead of purchasing monitoring software, companies subscribe to an intelligence service that continuously predicts operational risks and recommends actions.

The legal sector offers a third example. Harvey AI, built on large language models, supports law firms by analysing contracts, summarising regulations, and assisting with legal research. Firms subscribe to ongoing AI-generated insights that augment the work of legal professionals.

Across these examples, the economic model changes. Revenue shifts from one-time software licences to recurring subscriptions, and competitive advantage increasingly depends on data quality and model performance.

In essence, Generative AI enables organisations to sell continuous intelligence rather than static tools. For business leaders, this represents a significant shift in digital business models. The companies capturing the most value from AI are often those that transform expertise, data, and algorithms into subscription-based intelligence services that support decision-making over time.

Video 2: Outcome-Based Pricing

Let's look at another disruptive business model emerging in the age of Generative AI – outcome-based pricing.

Traditionally, organisations charged customers based on inputs. Software companies sold licences, consultants billed by the hour, and service providers charged for usage. In these models, customers paid for access to tools or expertise regardless of whether the final business outcome improved.

Generative AI allows organisations to price services differently – based on the measurable results delivered by AI systems.

A clear example comes from healthcare technology company Viz.ai. Stroke treatment is extremely time-sensitive; every minute of delay can significantly increase the risk of brain damage. Viz.ai uses AI to analyse CT scans and detect signs of stroke within seconds.

When a potential case is identified, the system immediately alerts specialist doctors through a mobile platform, accelerating treatment decisions. Hospitals adopting the platform often measure value through reduced treatment times and improved patient outcomes, making pricing models increasingly tied to the clinical impact created by the AI system.

A similar approach appears in marketing technology. Persado uses AI to generate and optimise marketing language across email campaigns, advertisements, and digital channels. By analysing consumer behaviour and linguistic patterns, its AI produces messaging designed to maximise engagement.

Companies using Persado often evaluate success through improvements in conversion rates and campaign performance, allowing pricing structures that link directly to measurable marketing outcomes rather than simply paying for access to the tool.

Industrial analytics provides another example. Companies such as Uptake analyse sensor data from industrial equipment to predict maintenance failures. Instead of selling monitoring software alone, deployments may tie pricing to cost savings achieved through reduced downtime or improved operational efficiency.

Across these examples, the economic logic changes. Organisations are no longer purchasing technology; they are paying for measurable business results enabled by AI.

For business leaders, outcome-based pricing represents a powerful shift. Generative AI enables providers to align revenue directly with the value they create – transforming AI from a technology purchase into a performance-driven business partnership.

Video 3: Cost Structure Inversion

For decades, many global organisations reduced costs through labour arbitrage. Work was moved to lower-cost regions while the underlying process remained the same. Industries such as business process outsourcing, customer support, and back-office operations scaled by employing large teams in lower-wage markets.

Generative AI is now changing this economic logic through what can be described as AI arbitrage. Instead of relocating work geographically, organisations reduce costs by automating cognitive tasks with AI systems. The cost advantage comes not from cheaper labour, but from requiring less labour altogether.

A clear example can be seen in customer support. The customer messaging platform Intercom introduced an AI support agent capable of answering common customer queries, resolving product questions, and guiding users through troubleshooting steps.

Previously, companies had to expand support teams or outsource call centres as customer volumes grew. With AI handling routine enquiries, organisations can manage thousands of conversations simultaneously without increasing support headcount. The cost structure shifts from labour-heavy service operations to AI-enabled automation.

Document processing provides another example. Companies such as Rossum apply AI to automatically read and interpret invoices, purchase orders, and financial documents.

In many finance departments, staff previously spent significant time manually entering invoice data into accounting systems. AI-powered document understanding can now extract and validate this information automatically, allowing organisations to process large volumes of invoices with far fewer administrative resources.

A similar productivity shift is occurring in software development. Platforms like Replit integrate AI coding assistants that help developers generate code, detect errors, and suggest improvements. Rather than scaling development teams linearly, companies can increase output by augmenting developers with AI capabilities.

Across these examples, the economic structure changes. Labour arbitrage reduced costs by relocating work. AI arbitrage reduces costs by reducing the amount of human labour required.

For business leaders, this represents a major shift. Competitive advantage will increasingly depend not on where work is performed, but on how effectively organisations deploy AI to automate and augment knowledge work.

Video 4: New Competitive Moats

In traditional industries, competitive advantage often came from scale, brand reputation, distribution networks, or patents. While these still matter, Generative AI is creating new types of competitive moats.

Today, some of the most powerful advantages come from proprietary data, specialised AI workflows, and fine-tuned models.

First, consider proprietary data. AI systems become far more valuable when trained on unique datasets that competitors cannot easily access. A strong example is LexisNexis, a global provider of legal intelligence. For decades, the company has built a vast database of legal cases, statutes, and regulatory documents. By applying AI models to this proprietary corpus, LexisNexis can help lawyers quickly identify precedents, analyse legal risks, and summarise complex case law. Because this dataset has been curated over many years, competitors without access to similar legal archives cannot easily replicate the same analytical capability. In this case, data itself becomes the competitive moat.

Second, competitive advantage can emerge from specialised AI workflows and prompts. Marketing platform Jasper provides AI tools for generating marketing copy, advertising campaigns, and brand messaging. However, the real value lies not only in the language model, but in the structured prompt frameworks and marketing workflows embedded within the platform. These frameworks encode best practices used by professional marketers, allowing users to generate consistent and brand-aligned content. Over time, these curated prompt structures become a form of intellectual property that improves the quality of AI outputs.

Third, organisations build advantage through fine-tuned models. Biotechnology company Insitro uses machine learning to accelerate drug discovery. The company trains AI models on specialised biological datasets and experimental research, enabling the models to identify complex patterns in genetic and pharmaceutical data. These continuously refined models become increasingly specialised and difficult for competitors to replicate without similar research data and expertise.

Across these examples, the source of competitive advantage shifts. In the Generative AI era, sustainable moats increasingly come from unique data assets, domain-specific AI workflows, and continuously improving models. For business leaders, the lesson is clear: adopting AI tools alone is not enough. Long-term advantage comes from building AI capabilities and intelligence assets that competitors cannot easily reproduce.

Video 5: Capital Allocation in the AI Era

As Generative AI becomes a strategic capability, organisations face a critical question: how should they invest in AI? Leaders must decide whether to build AI internally, buy external solutions, or partner with specialised providers. Each option involves different trade-offs in cost, speed, and long-term control.

Let us begin with the build strategy. Some organisations choose to develop AI models internally when the technology directly supports their core business advantage. A notable example is Bloomberg, which developed its own large language model called BloombergGPT.

The model was trained on decades of proprietary financial data, including market reports, filings, and financial news. Because the training data is highly specialised, the model can perform tasks such as financial document analysis, sentiment detection, and information extraction with greater domain accuracy.

By building the model internally, Bloomberg retains control over its AI capabilities and strengthens its competitive position in financial intelligence.

However, building AI systems is expensive and requires specialised talent. As a result, many organisations adopt a buy strategy. For example, companies often use platforms such as DataRobot to deploy machine learning models for forecasting, risk analysis, and operational optimisation. Instead of building AI infrastructure from scratch, firms can quickly implement predictive analytics tools and integrate them into existing business systems. This approach significantly reduces the time and resources required to adopt AI capabilities.

A third approach involves strategic partnerships. Pharmaceutical company Pfizer has collaborated with technology partners such as IBM to apply AI in biomedical research and drug discovery. These partnerships allow pharmaceutical scientists to combine their domain expertise with advanced AI tools, accelerating research without building the entire technology stack internally.

Across industries, organisations increasingly use a portfolio approach: building AI where it creates strategic differentiation, buying solutions for common capabilities, and partnering for specialised innovation.

For finance and strategy leaders, capital allocation in the AI era requires balancing experimentation, speed, and long-term capability development to ensure that AI investments generate sustainable competitive advantage.

Video 6: Redefining Roles and Job Architectures

What if one of the biggest impacts of Generative AI is not the elimination of jobs – but the redesign of how work itself is structured?

As organisations adopt Generative AI, they are increasingly shifting from thinking about jobs as fixed roles to analysing work at the task level. This approach is known as task decomposition.

Instead of asking whether AI can replace an entire role, companies examine the individual tasks within that role and determine which can be automated, which can be augmented by AI, and which still require human expertise.

A useful example comes from Deloitte, which has been integrating Generative AI tools into consulting workflows.

Traditionally, consultants spent a significant portion of their time gathering information, reviewing lengthy reports, synthesising market research, and preparing presentation decks for clients. While

these activities were essential, they were also highly time-consuming, particularly in the early stages of consulting engagements.

With Generative AI tools, many of these tasks can now be accelerated. AI systems can summarise hundreds of pages of industry reports in minutes, generate structured outlines for client presentations, and draft early versions of analytical insights.

Instead of starting from a blank page, consultants now begin with AI-generated drafts that help them move faster through the research and documentation phase.

Importantly, this does not eliminate the role of consultants. Instead, it reshapes how their time is allocated.

Consultants spend less time compiling information and formatting slides, and more time focusing on interpreting insights, designing strategic recommendations, and engaging directly with clients.

Junior consultants who once spent large portions of their time on manual research increasingly act as AI-assisted analysts, validating and refining outputs produced by AI systems.

This illustrates a broader shift in job design. Rather than replacing entire roles, Generative AI reallocates tasks between humans and intelligent systems.

For organisations, the implication is clear. Workforce planning must increasingly focus on skills, tasks, and human – AI collaboration, rather than relying only on traditional job titles.

In the Generative AI era, competitive organisations are not simply automating work – they are re-architecting how work itself is performed.

Video 7: From Headcount to Capability Density

For decades, organisational growth followed a simple formula: more work required more people. As companies expanded operations, they hired additional employees to increase productivity, manage processes, and deliver services. In this model, organisational capability was often measured by headcount.

But Generative AI is beginning to challenge that logic.

Instead of focusing primarily on workforce size, organisations are now shifting toward capability density – the amount of expertise, analytical power, and productivity available within a team.

When employees are supported by advanced AI tools, smaller teams can deliver the level of output that previously required much larger workforces.

A powerful example comes from Moderna, one of the pharmaceutical companies that rapidly scaled its research capabilities during the development of mRNA-based vaccines.

Drug discovery and development traditionally require large research teams. Scientists must analyse enormous volumes of biological data, review scientific literature, and run complex simulations to identify viable drug candidates – activities that typically take months or even years.

Moderna has invested heavily in AI and machine learning systems that help scientists analyse genomic data, identify biological patterns, and model potential drug targets.

These tools allow researchers to evaluate massive datasets and generate hypotheses far more quickly than with traditional research workflows.

The result is not simply faster analysis – it fundamentally changes how research teams operate. Scientists can test more ideas, explore more biological pathways, and accelerate experimentation cycles without proportionally increasing team size.

In effect, the capability of each researcher increases dramatically.

This exemplifies the concept of capability density. When AI systems amplify human expertise, every employee gains access to far greater analytical power, enabling smaller teams to solve problems that once required significantly larger organisations.

For leaders, the strategic implications are significant. Organisational strength will increasingly depend not on how many employees a company has, but on how effectively human expertise is augmented by intelligent systems.

In the Generative AI era, competitive advantage will belong to organisations that maximise capability per employee – not simply grow headcount.

Video 8: Talent Strategy in an AI-Native Firm

Talent strategy in an AI-native firm begins with a simple shift in emphasis from routine execution to judgement, systems thinking, and responsible use of machine intelligence. The market is already moving in this direction, and the strongest signals are coming from firms that have paired adoption with large-scale capability building and clear operating guardrails.

Accenture's multiyear investment in advanced AI has been matched by workforce action and measurable revenue impact, reporting tripled advanced AI revenues in fiscal year 2025 and training in AI fundamentals at massive scale.

The firm grew its AI and data talent base from forty thousand to seventy-seven thousand while training more than five hundred and fifty thousand employees, a signal that breadth of literacy and depth of expertise now move together in hiring and development.

PwC has embedded generative AI into day-to-day consulting and audit work while rolling out enterprise-grade tools to its people. It is OpenAI's first reseller for ChatGPT Enterprise and has been equipping over one hundred thousand employees in the US and UK, underpinned by a firmwide upskilling programme that has already engaged the vast majority of its workforce.

The message to candidates is clear that AI literacy, responsible use, and domain judgement are baseline expectations.

In financial services, JPMorgan is treating AI as a core skill and not a niche speciality. New hires receive prompt engineering training, and the bank estimates more than a billion dollars of value across hundreds of use cases, while leaders describe a roadmap toward agentic systems embedded in everyday work.

This is accompanied by a structured, bank-wide curriculum so different roles learn how to interrogate and validate models in context.

Indian IT leaders are also recasting talent at scale. TCS has trained more than one hundred and fifty thousand employees in generative AI and created an internal experience zone to accelerate hands-on proficiency across tools and models, while Wipro committed a one-billion-dollar programme to AI and has already trained well over two hundred thousand employees on fundamentals and responsible use.

The toolchain is evolving as well, which changes job design. Amazon Q Business is now generally available and is being integrated into everyday applications so non-technical employees can query enterprise knowledge and automate multi-step workflows, a shift that elevates skills in orchestration, verification, and data stewardship across roles.

Across these cases, the hiring bar is moving toward people who can frame problems, test model outputs, manage risk, and design workflows that combine humans and AI.

Upskilling, governance, and measurable impact are no longer differentiators but the price of entry for an AI-native organisation.

Video 9: Competitive Advantage in an AI-Saturated World

When most firms can license similar models and cloud services, technology parity arrives quickly. Advantage comes from how well organisations integrate AI into the operating system of the business – decision rhythms, data flows, governance, and the speed with which insights become action. Zara illustrates this better than almost anyone in retail.

Zara's historic strength was speed. AI has turned that speed into a system. The company runs a real-time 'AI inventory and trend engine' that combines natural-language trend scraping, computer-vision pattern recognition, RFID telemetry, and hyper-local demand forecasting, enabling design-to-store cycles in roughly two weeks at global scale. Zara augments social and runway signals with daily qualitative inputs from store managers via handheld devices. Designers and planners in Arteixo act on a live signal picture rather than seasonal assumptions, which shortens feedback loops and improves hit rates.

Proprietary context is Zara's durable moat. Every garment carries RFID, creating a single, real-time stock view across more than 5,600 stores. AI-driven allocation now governs the majority of initial production decisions at SKU-colour-size granularity, while robotics accelerate restocking and reduce

error. This stack is tied to logistics algorithms that optimise routes and warehouse assignment, reducing last-mile emissions and cycle time – evidence that the system is designed for both speed and cost.

The signal side is equally technical. NLP models scan millions of posts to identify micro-trends before they peak. Computer vision detects recurring silhouettes and palettes from street and runway imagery. Behaviour mapping from Jetlore converts preferences in fit, colour, and style into structured attributes that feed design, buying, and merchandising decisions. Crucially, these AI outputs are reviewed and acted upon by cross-functional teams, preserving human judgement where context matters and preventing blind automation.

The lesson here is clear. In an AI-saturated world, organisations differentiate through execution architecture. Firms build proprietary data flows that competitors cannot easily access, compress sensing-to-production cycles through human-AI collaboration, and align governance so models can influence decisions rather than simply generate presentations. Information, rather than inventory, becomes the competitive edge – and the more a system learns from its own operations, the harder it becomes to replicate.

Video 10: Regulatory and Ethical Disruption

As generative AI moves into mainstream operations, two areas now dominate the risk agenda: intellectual property and synthetic media. Both are evolving rapidly and require practical governance rather than broad principles alone.

On intellectual property, a pivotal reference point is the 2025 UK High Court ruling in Getty Images v Stability AI. Getty alleged that millions of its copyrighted images were scraped to train Stable Diffusion. The court largely dismissed copyright claims – partly because model training occurred outside the UK and because learned weights were not treated as copies of images – while finding limited trade mark infringement where older model versions produced outputs containing Getty watermarks.

The judgment clarified some issues but left core questions unresolved, including whether training on protected data is itself infringing and how creators should be compensated when their work contributes to model performance. Organisations deploying generative AI should therefore assume increased scrutiny around dataset provenance, licensing, and output filtering, even when training occurs outside a particular jurisdiction.

Synthetic media presents a different category of risk. AI-generated images, video, and voice can enhance marketing, education, and customer engagement. However, the same tools can also enable impersonation and deception. Voice cloning has already been used to imitate executives and trigger fraudulent payments, while deepfakes can spread rapidly enough to create reputational and market damage before corrections are issued.

Controls gaining traction include provenance and watermarking systems, layered content authentication, model-side safety filters designed to suppress impersonation, and enterprise playbooks for rapid verification and takedown procedures. Risk-based AI regulation and platform governance policies are also increasing pressure for transparency and auditable controls.

What should leaders do now? First, organisations should inventory where models interact with copyrighted or rights-managed material and implement contractual and technical controls for ingestion, fine-tuning, and output moderation. Second, synthetic-media threats should be incorporated into fraud, security, and communications runbooks, including mandatory call-back verification for high-value requests and media provenance checks.

Third, organisations should create an internal ‘model governance pack’ that documents dataset origins, licences, filtering approaches, and output testing procedures. This allows firms to demonstrate due diligence as regulatory expectations continue to strengthen.

Trust and accountability are increasingly becoming operational capabilities. Organisations should treat them as integral components of product and process design rather than as after-the-fact compliance requirements.

Video 11: Commoditisation of Intelligence

What happens when ‘intelligence’ becomes something any organisation can buy off the shelf? That question defines one of the biggest strategic risks in the age of Generative AI.

Hello learners, and welcome to this session on the Commoditisation of Intelligence.

As generative AI becomes widely available through cloud platforms, open-source models, and ready-made APIs, intelligence is no longer limited to organisations with specialist teams or expensive infrastructure. Today, almost any business – from a start-up to a multinational – can access the same class of language models, reasoning tools, and automation systems.

This ease of access creates a quiet but significant shift. When many organisations use similar models trained on similar public data, their outputs can begin to resemble one another. Marketing content may follow similar patterns. Product descriptions may sound alike. Analytical summaries may rely on the same style of reasoning embedded within the underlying model.

A clear example is customer service automation. Tools such as Zendesk AI use pre-trained language models to classify tickets, suggest replies, and generate automated responses. Because these systems rely on shared components – including general-purpose embeddings, intent recognisers, and structured reply templates – customer interactions often follow similar phrasing and problem-solving patterns across different organisations.

The technical reason behind this standardisation is straightforward. Most plug-and-play AI tools rely on pre-trained parameters shaped by public datasets. Without additional fine-tuning or proprietary data, the system generates outputs that reflect the statistical patterns within those datasets. This means organisations using out-of-the-box models risk producing content or decisions that feel generic.

The strategic risk is that when many organisations rely on the same AI capabilities, differentiation becomes more difficult. Competition can shift away from originality and value creation towards cost

or speed. Organisations that depend only on general AI tools may find their advantages narrowing over time.

However, this outcome is not inevitable. The key to maintaining differentiation is combining AI with what is unique to the organisation itself. This includes proprietary data, internal knowledge, domain expertise, and distinctive workflows. When organisations integrate their own operational data into models – such as customer history, product performance, or specialist terminology – they shape outputs in ways that competitors cannot easily replicate.

In an AI-enabled economy, the challenge is not simply using intelligence but ensuring that the intelligence reflects the organisation's unique strengths. The goal is to move from generic intelligence towards context-rich, organisation-specific intelligence – the kind that remains defensible even when powerful AI tools are available to everyone.

Video 12: The AI-Native Enterprise of 2035

If every decision in your company could be sensed, modelled, and acted on in seconds, what would you redesign first? That is the question the AI-native enterprise of 2035 will answer – by building firms where human judgement and intelligent systems co-own the work.

AI-native enterprises begin with a fundamentally different operating architecture. Instead of workflows driven by managerial escalation, decisions are increasingly powered by real-time optimisation engines embedded directly into the flow of work. One of the clearest early examples comes from UPS, whose ORION route optimisation system uses telematics, traffic data, and behavioural patterns to eliminate up to 100 million unnecessary miles per year, saving 10 million gallons of fuel and generating an estimated USD 300–400 million in value annually.

These systems do not sit passively on dashboards. They operate continuously – an early signal of how enterprises will shift from periodic reviews towards algorithmic control loops.

The second defining feature is large-scale human – AI collaboration. AI-native organisations do not replace expertise; they amplify it. DeepMind's AlphaFold, for example, unlocked access to more than 200 million protein structure predictions and became essential scientific infrastructure within only a few years, earning a Nobel Prize in 2024 for its impact on biology and drug discovery.

This pattern – AI handling complexity at machine speed while humans interpret and apply insight – will become standard across fields ranging from research and development to diagnostics and engineering.

Leadership itself also transforms in AI-native firms. Instead of directing every decision manually, leaders increasingly design the rules, safeguards, and data foundations that govern autonomous systems. Siemens illustrates this shift through Senseye Predictive Maintenance, which combines machine-learning-driven fault detection with generative AI capabilities that allow maintenance teams to interrogate machine states conversationally, accelerating diagnosis and reducing downtime for industrial operators worldwide.

Leaders in such environments focus on model oversight, data quality, and ensuring that algorithmic decisions remain aligned with operational intent.

By 2035, the AI-native enterprise will be defined by intelligent workflows, expert – AI co-creation, and leadership models designed for algorithmic environments. The organisations that succeed will be those that redesign their structures early, build proprietary data advantages, and combine automation with strong human judgement.

Ultimately, AI-native does not mean less human – it means systems that allow people to operate at their highest level, with intelligence embedded wherever work takes place.

Video 13: Final Thoughts

This week, we explored how Generative AI is reshaping organisations through business model innovation, organisational redesign, and strategic shifts.

We saw how firms are moving towards intelligence-led models, with subscription and outcome-based pricing replacing traditional approaches. Value creation is increasingly driven by AI arbitrage, where automation, data, and decision systems scale impact. At the organisational level, roles are being redefined through AI augmentation, enabling greater capability with leaner teams and shifting the focus from headcount to capability density.

We also examined how long-term advantage will depend on proprietary context, speed, culture, and responsible AI use.

Stepping back, across this module, we explored Generative AI as a disruptive force – one that goes beyond improving existing processes to fundamentally redefining how organisations create and capture value.

From collapsing industry boundaries and intelligence-centric firms to new customer interfaces, operational models, and organisational structures, the shift is systemic. The key takeaway is that in an AI-saturated world, advantage will not come from access alone, but from how effectively organisations integrate AI into strategy, decision-making, and execution.

Generative AI, therefore, is not just enabling change – it is driving enterprise-wide reinvention.