

Video Transcripts

Module 7: Agentic AI in Business Functions I – Sustaining innovations

Video 1: What is Agentic AI in Business Terms

Welcome to this module. Sustaining innovations refer to improvements that enhance existing processes without fundamentally altering them. In the context of Generative AI, this includes refining capabilities such as content creation or data analysis within tools that organisations already use. Agentic AI builds upon this idea by introducing autonomous execution of multi-step workflows. Rather than simply supporting tasks, it can plan, decide, and act across business functions such as procurement or finance, delivering operational efficiency while maintaining established organisational structures.

Building on the fundamentals of Agentic AI explored earlier, this module shifts focus from theory to practical application. Previously, attention was given to how agentic systems reason and utilise tools to complete tasks. Now, the emphasis is on how these capabilities are embedded within corporate environments. The transition is from understanding how an AI system “thinks” to examining how it actively enhances workflows, optimising processes such as procurement and finance to achieve immediate operational improvements.

To illustrate this, consider the role of a procurement manager in a UK-based manufacturing firm facing a sudden logistics disruption in Europe. In a traditional digital setup, AI might assist by highlighting the issue or helping draft communications, but it remains a passive tool dependent on human direction. The responsibility for resolving the issue, such as identifying alternative suppliers or adjusting production schedules, still rests entirely with the manager.

Agentic AI transforms this dynamic by shifting from assistance to autonomy. In the same scenario, an agentic system would independently identify the disruption, evaluate alternative suppliers within defined parameters, arrange expedited shipments, and update production timelines—all without waiting for human input. This demonstrates its capacity to plan, decide, and execute across multiple steps, effectively taking proactive action rather than merely supporting decision-making.

This capability is particularly important for maintaining process continuity, a common challenge in modern operations where workflows often stall due to exceptions or disruptions. Agentic systems are designed to navigate uncertainty, ensuring progress even when conditions are imperfect. A practical example can be seen in Ocado’s automated fulfilment centres, where agentic systems coordinate thousands of robots in real time, dynamically adjusting routes and inventory handling based on live demand. Such systems do not replace human expertise but enhance it by reducing executional friction, allowing professionals to focus on higher-value judgement and strategic decision-making.

Video 2: Task → Workflow → Agent Transformation

Automation in modern business is progressing through a clear evolution, moving from isolated task execution to connected workflows and ultimately to autonomous, agent-driven systems. In the earliest stage, organisations relied on task automation—tools designed to complete a single, repetitive action, such as extracting contact details. While efficient, these systems operate without awareness of the wider business context, requiring human involvement to carry outputs forward and limiting their overall value.

Workflow automation marked the next step in this evolution by connecting individual tasks into structured sequences. For instance, capturing a lead’s email could automatically trigger a follow-up message. Although this improves efficiency, such workflows depend on rigid, rule-based logic. When unexpected or complex inputs arise, these systems struggle to adapt, often breaking down and requiring manual intervention. This creates “glue work”, where professionals must step in to bridge gaps, reducing productivity and slowing processes.

Agentic AI represents a significant advancement by replacing rigid workflows with goal-oriented execution. Rather than following predefined steps, these systems are guided by high-level objectives and can adapt their actions accordingly. For example, instead of sending generic outreach, an agentic system can analyse a prospective client’s context, identify their needs, and craft a tailored engagement strategy. It determines who to contact, how to communicate, and when to follow up, demonstrating a more intelligent and flexible approach.

This shift from task support to workflow ownership fundamentally changes how organisations allocate human effort. By delegating coordination and exception handling to autonomous systems, managers are freed from supervising processes and can focus on strategic priorities. As a result, adopting agentic AI is not simply an incremental improvement but a foundational step towards scalable innovation, enabling organisations to enhance efficiency while allowing human expertise to be applied where it delivers the greatest value.

Video 3: Human-in-the-Loop Design

One of the most important design decisions when deploying Agentic AI is determining where human involvement should remain. Full autonomy is not always appropriate, particularly in business contexts where accountability and oversight are essential. The objective is to carefully define which decisions can be delegated to the agent and which require human approval, ensuring that the system operates efficiently without introducing unnecessary risk or delays.

Human oversight is especially important in situations where errors could have serious, irreversible, or regulated consequences. For example, in financial services, an agent can autonomously collect and analyse transaction data, detect anomalies, and prepare reports. However, the decision to submit a report to a regulatory authority typically requires a compliance officer’s approval. In this model, the agent improves efficiency by handling preparation, while the human retains responsibility for the final decision.

A central consideration in designing such systems is the balance between speed and risk. Requiring human approval at every stage offers strong control but reduces the speed and efficiency gains that Agentic AI can deliver. Conversely, allowing full autonomy increases speed but may expose the organisation to greater risk in high-stakes scenarios. To manage this, organisations must assess workflows based on the potential impact of errors and how easily those errors can be reversed.

Practical applications of this approach can be seen in large-scale operations such as Amazon's fulfilment network. Agentic systems are trusted to manage routine decisions like inventory routing and shipment scheduling because these actions are frequent, low-risk, and easily corrected. However, more significant decisions—such as major contractual commitments or responses to large disruptions—still require human review before execution.

To support this balance, organisations rely on structured escalation frameworks. These define when an agent should pause and involve a human, the urgency of that escalation, and the information needed for effective decision-making. By formalising these processes, businesses ensure consistency and efficiency while maintaining governance. Ultimately, human-in-the-loop design is about aligning autonomy with risk, enabling organisations to benefit from Agentic AI while preserving control and accountability.

Video 4: Autonomous Campaign Management

Marketing campaigns have traditionally required constant human oversight, with teams responsible for launching, monitoring, and adjusting performance. This process is time-consuming and limited by how quickly individuals can interpret data and respond. Agentic AI is changing this model by shifting from supportive tools to systems that can manage campaigns autonomously.

In this approach, an AI agent takes end-to-end ownership of campaign execution. It launches campaigns, tracks results in real time, and makes immediate adjustments to budgets, targeting, and messaging. This creates a continuous feedback loop where insights are instantly acted upon, removing delays between analysis and implementation.

A major advantage of this model is autonomous A/B testing. Instead of waiting days or weeks for results, agentic systems rapidly generate and test multiple variations, reallocating resources to the most effective options while discarding weaker ones. This significantly accelerates optimisation and improves overall campaign performance.

Examples such as Persado and Albert AI highlight this shift in practice. These systems continuously refine messaging and manage media spend across platforms without manual input. As a result, marketing teams can move away from routine adjustments and focus on strategic priorities such as brand positioning and creative direction.

Video 5: Customer Support Resolution Agents

Customer support is undergoing a major shift, moving from systems that simply respond to queries to those that manage the entire resolution process. Traditionally, support has been fragmented, involving multiple steps such as responding with templates, escalating issues, and manually closing cases. This approach is slow, dependent on human coordination, and often leads to a poor customer experience where individuals must repeat information. Agentic AI replaces this reactive model with end-to-end resolution agents.

These agents act as proactive problem-solvers rather than basic chatbots. When a customer raises an issue, such as a delayed order, the agent goes beyond providing standard information. It investigates the root cause, checks internal systems, and takes appropriate actions, such as initiating a refund or arranging a replacement if certain conditions are met. This allows the system to achieve outcomes independently without requiring step-by-step human intervention.

An important capability of these systems is intelligent escalation. While many issues can be handled autonomously, more complex or sensitive cases are escalated to human specialists. The key difference is that the agent provides a complete context during the handover, allowing the human to focus on solving the problem rather than gathering information. After the intervention, the agent can resume control to complete follow-ups and closure.

In practice, this model can be applied across industries. In banking, agents can monitor suspicious transactions and take preventative action, while in retail, they can manage returns from initiation to completion. By transforming support into an autonomous resolution function, organisations can eliminate delays, improve consistency, and free human teams to focus on complex decisions and empathetic interactions.

Video 6: Sales Pipeline Management

In many organisations, sales professionals spend a significant portion of their time on administrative tasks rather than engaging with customers. Activities such as updating CRM systems, sorting inbound leads, and drafting follow-up emails often take priority before meaningful sales conversations can even begin. Agentic AI changes this dynamic by transforming the CRM from a passive data repository into an active system that manages and progresses the sales pipeline.

One of the most immediate benefits is seen in autonomous lead qualification. Instead of simply recording a prospect's interaction, an agentic system evaluates it in real time by analysing external data sources and identifying intent signals. It scores and prioritises leads, routes high-quality opportunities to the appropriate salesperson, and filters out low-value prospects. This ensures that the pipeline remains focused and efficient without requiring manual triage.

Agentic systems also improve follow-up consistency, which is a common challenge in sales. Rather than relying on individuals to remember and execute multiple touchpoints, the system tracks each prospect's stage and automatically initiates timely, relevant communication. Messages are tailored based on prior interactions and sent at optimal times, with responses seamlessly logged and used to adjust future engagement steps.

Overall, agentic AI removes much of the administrative burden that limits sales productivity. By handling lead qualification, maintaining follow-up discipline, and updating records automatically, it enables sales professionals to concentrate on high-value activities such as building relationships and closing deals, where human judgement and interpersonal skills remain essential.

Video 7: Workflow Automation Across Functions

Workflow automation across functions is one of the most impactful applications of Agentic AI in business operations. Traditionally, processes are split across departments such as sales, finance, operations, and logistics, each operating within separate systems. Coordination between these functions often relies on manual handoffs, which can introduce delays, inefficiencies, and reduced visibility across the overall workflow.

Agentic AI addresses this challenge by enabling intelligent, end-to-end coordination through autonomous agents. Rather than functioning in isolation, these agents operate across the entire process, ensuring that information flows smoothly and actions are triggered in the correct sequence. By continuously observing signals and interacting across systems, they create a unified and uninterrupted workflow.

A clear example is the order-to-dispatch process. When a customer places an order, the system records it and an agent verifies the details instantly. Billing is then triggered to generate an invoice and confirm payment, after which logistics processes begin to coordinate dispatch. Each stage is automatically connected, removing the need for manual intervention and eliminating the delays commonly seen in traditional workflows.

A similar approach can be applied to customer onboarding. When a new customer signs up, agents coordinate tasks across marketing, sales, and support functions. They create profiles, initiate communication, assign ownership, and trigger onboarding steps automatically. This ensures consistency and speed, even though multiple teams remain involved in the process.

In summary, Agentic AI transforms fragmented, cross-functional workflows into coordinated and automated systems. By ensuring seamless task progression, improving transparency, and reducing manual effort, it enables organisations to operate more efficiently and effectively at scale.

Video 8: Continuous Compliance Monitoring

Compliance has traditionally been a reactive function, relying on periodic audits and manual reviews that often identify issues only after they occur. In highly regulated sectors such as global banking, this approach is no longer sufficient, as it cannot keep pace with the scale and complexity of modern operations. Agentic AI is transforming compliance by shifting it from retrospective checks to continuous, real-time monitoring.

In this new model, compliance agents act as always-on observers, continuously aligned with evolving regulatory requirements. Rather than waiting for scheduled reviews, they monitor operations around the clock and adapt instantly to regulatory updates. For example, when a regulator such as

the Financial Conduct Authority introduces a policy change, the system automatically adjusts its monitoring criteria, ensuring consistent enforcement across all relevant processes.

A strong example of this transformation can be seen in large global institutions such as HSBC. Operating across numerous jurisdictions, the organisation previously faced the challenge of manually screening massive volumes of transactions. With agentic systems, this process has become intelligent and proactive, where transactions are analysed in context by cross-referencing customer data, behavioural patterns, and risk indicators. When potential financial crime is detected, the system can take immediate action, such as holding transactions and preparing detailed reports for human review.

A defining feature of agentic compliance is its ability to take autonomous corrective action in lower-risk situations. For instance, if a required disclosure is missing, the system can identify and correct it without human input. Similar applications are seen in industries like pharmaceuticals, where agents monitor production conditions and trigger corrective protocols when deviations occur. This ensures that compliance becomes an integrated, continuous function rather than a delayed, manual process.

Video 9: Financial Close Acceleration

The monthly financial close is one of the most time-sensitive processes in any organisation, traditionally requiring finance teams to manually reconcile accounts, consolidate data, and resolve discrepancies across multiple systems. This process is labour-intensive, prone to human error, and heavily dependent on coordination between departments. Agentic AI is transforming this function by shifting from manual effort to autonomous orchestration of the financial close.

At the core of this transformation are intelligent agents that actively manage reconciliation rather than simply storing data. Instead of finance professionals logging into various banking and enterprise systems to match transactions, these agents connect to data sources independently. They extract relevant information, identify mismatches, and highlight specific exceptions for human review in real time, significantly reducing manual effort.

This approach is particularly valuable in large multinational organisations such as Unilever or Vodafone, where multiple subsidiaries complete financial closing simultaneously. Using platforms like BlackLine, agentic systems monitor the progress of reconciliation tasks across regions, identify outstanding items, and automatically escalate them to the relevant teams. What once required extensive manual follow-up across different time zones is now handled continuously by the system.

As a result, organisations experience a substantial reduction in the time required to complete the financial close. Processes that previously took several working days can be completed much faster, allowing finance teams to move beyond routine data handling. With the operational workload reduced, professionals can focus on validating results, analysing financial performance, and delivering meaningful insights to support strategic decision-making.

Video 10: Employee Onboarding Agents

First impressions play a crucial role in shaping an employee's experience, sense of belonging, and speed to productivity. Traditionally, onboarding has been a fragmented process involving multiple departments, disconnected systems, and extensive manual coordination. Employee onboarding agents address this complexity by integrating these activities into a seamless, intelligent workflow that ensures consistency and efficiency from the outset.

One of the earliest challenges in onboarding is document collection. New employees are required to submit various forms and identification documents, and this is often managed manually through emails with limited tracking. An onboarding agent streamlines this process by sending personalised requests, monitoring submission status in real time, and issuing targeted reminders. It can also validate submitted documents against predefined criteria and flag incomplete or incorrect entries for human review, reducing administrative effort.

System access provisioning is another critical aspect that typically requires coordination across IT, HR, and management. In traditional setups, delays and errors are common due to manual requests and approvals. An onboarding agent automates this process by triggering access rights based on the employee's role, department, and start date. Platforms such as ServiceNow demonstrate how these workflows can be coordinated efficiently, ensuring that new hires have the necessary tools and access from their first day.

Training workflow orchestration further enhances the onboarding experience. Instead of providing generic training lists, onboarding agents create personalised learning paths tailored to the employee's role and seniority. They enrol employees in relevant modules, track completion, and send reminders for pending tasks. If critical training is not completed on time, the system can escalate the issue to the appropriate manager, ensuring compliance and progress.

To summarise, onboarding agents transform a traditionally manual and disjointed process into a cohesive and automated experience. By improving coordination, reducing administrative workload, and ensuring consistency, they accelerate time to productivity while maintaining a high-quality onboarding journey, even as organisations scale.

Video 11: Internal Helpdesk Automation

Most organisations rely on a constant flow of internal requests, ranging from IT support to HR queries. Traditionally, these are managed through ticketing systems where human agents manually categorise, prioritise, and assign each issue. This triage process often leads to delays, inconsistencies, and inefficiencies in service delivery. Agentic AI is transforming this model by shifting from basic ticket handling to autonomous internal helpdesk execution.

In this new approach, systems act as active problem-solvers rather than passive record-keepers. For example, agentic solutions such as Atlassian's Rovo agents go beyond chat-based assistance to function as "AI teammates". When an employee submits a request, the agent analyses the intent and takes action. If access to a software tool is requested, the system can verify the employee's role, check compliance with security policies, and provision access with minimal human intervention.

The impact of this shift is substantial. Agentic systems can manage the full lifecycle of a support ticket, from refining its description to adjusting priority based on urgency and routing it to the correct specialist with complete context. They can also monitor service levels and automatically escalate unresolved issues when necessary. This ensures faster response times, greater consistency, and improved user experience, even across non-technical functions such as finance and HR.

Overall, this represents a move from simply recording problems to actively orchestrating solutions. By automating routine workflows and administrative steps, agentic systems reduce the operational burden on support teams. This allows human experts to focus on more complex, high-value issues that require critical thinking and judgement, while employees benefit from faster and more reliable support.

Video 12: Control, Auditability & Trust in Agentic Systems

As agentic AI becomes more deeply embedded across functions such as procurement, finance, HR and compliance, organisations must address how to maintain oversight of systems that act autonomously. This introduces a critical need to balance autonomy with control. Achieving this depends on three core capabilities: comprehensive decision logging, clear explainability of actions, and well-defined accountability frameworks.

Decision logging forms the foundation of any reliable agentic system. Every action taken by an agent must be recorded in a structured and retrievable format, capturing not only what was done, but when it occurred, what data informed it, and what outcome followed. This creates a transparent audit trail that allows organisations to review performance, identify errors, and demonstrate compliance when needed.

This level of traceability is essential in high-stakes environments. For example, if a financial services agent flags and holds a suspicious transaction, the organisation must be able to clearly justify that decision. Detailed logs enable teams to reconstruct the exact reasoning process, including the signals and rules that triggered the action, ensuring accountability when reviewed by regulators or challenged by customers.

Explainability adds another layer of trust by addressing why an agent acted in a particular way. Systems must present their reasoning in clear, human-readable terms so that non-technical stakeholders can understand and validate decisions. This is particularly important in regulated industries, where transparency is required. Tools and frameworks designed for AI explainability support this by making decision logic more accessible and interpretable.

To summarise, effective oversight of agentic AI depends on combining robust logging, clear explainability, and structured accountability. Together, these capabilities ensure that organisations can benefit from autonomous systems while maintaining transparency, governance, and trust in their operations.

Video 13: Final Thoughts

Welcome to the final session of this module on Agentic AI in Business Functions, focused on sustaining innovations. Throughout the module, we have explored how organisations are evolving from isolated automation to connected workflows and increasingly autonomous agents, while maintaining the importance of human oversight to ensure control and accountability.

We examined how agentic AI is being embedded across core business functions. In marketing and sales, it enables personalisation and stronger pipeline management, while in operations it reduces delays and improves execution by ensuring processes continue seamlessly without manual intervention.

The module also highlighted practical applications across areas such as procurement, finance, compliance, and HR. In each case, agentic systems take over repetitive and time-sensitive tasks, improving speed, accuracy, and consistency, while enhancing both operational performance and employee experience.

The key takeaway is that agentic AI is already reshaping how organisations operate today. As we move forward, the focus shifts beyond optimising existing processes to exploring how these systems can drive deeper transformation, creating new business models and redefining competitive advantage.