

Video Transcripts

Week 12.1: AI Governance

Video 1: Why AI Governance Matters in 2026

AI governance has become a strategic necessity as organisations increasingly adopt generative and agentic AI systems. These technologies can create content, support decision-making, and perform tasks with limited human intervention. As AI capabilities expand, risks now extend beyond model accuracy and bias to include system behaviour, unintended actions, and wider organisational impacts. Governance provides the structure needed to manage these evolving risks responsibly.

Real-world incidents highlight the importance of governance. Air Canada was held accountable when its chatbot provided incorrect refund information, while Samsung faced concerns after sensitive information was shared through generative AI tools. These cases demonstrate that AI-related errors can lead to legal, financial, security, and reputational consequences. As AI systems become more autonomous, the potential impact of such mistakes grows significantly.

The need for governance is especially critical in sectors such as healthcare and financial services. Generative AI can assist with clinical documentation and decision support, but inaccuracies and hallucinations can still occur. Similarly, autonomous financial agents have demonstrated unintended behaviours when pursuing specific objectives. Human oversight, validation processes, and clear accountability remain essential to ensure AI systems operate safely and reliably.

Strong governance enables organisations to innovate with confidence while maintaining trust and compliance. Companies such as Coca-Cola have successfully scaled AI adoption through approval workflows, content controls, and audit trails. At the same time, increasing regulatory expectations and lessons from past failures, such as Amazon's biased hiring tool, reinforce the importance of transparency and accountability. As AI becomes more influential, governance provides the foundation for responsible and sustainable AI adoption.

Video 2: Ethical Principles of Responsible AI

Responsible AI is built on key ethical principles that guide how AI systems are designed, deployed, and monitored. Fairness ensures that systems do not produce biased or discriminatory outcomes, while transparency helps users and stakeholders understand how AI-driven decisions are made. These principles are essential for building trust and ensuring that AI systems operate in a responsible and reliable manner.

Accountability ensures that organisations clearly define ownership for AI outcomes across development, deployment, and operational teams. When responsibilities are clearly assigned, issues can be identified and addressed more effectively. Transparency and accountability work together to improve compliance, reduce disputes, and strengthen confidence in AI-powered processes.

Privacy and safety are equally important components of responsible AI. Organisations must protect sensitive data through strong governance practices, including data protection, anonymisation, and access controls. Safety focuses on ensuring that AI systems perform reliably and minimise the risk of harm, particularly in environments where errors can have significant operational or customer impacts.

Ethical failures can quickly become business risks, affecting compliance, reputation, and customer trust. Bias, lack of transparency, weak accountability, or poor data protection can lead to serious consequences. By embedding ethical principles throughout the AI lifecycle—from design and development to deployment and monitoring—organisations can manage risks effectively, maintain user confidence, and support the responsible adoption of AI at scale.

Video 3: AI Governance Evolution Across AI Types

AI governance has evolved significantly as AI systems have progressed from traditional models to generative and agentic systems. Traditional AI was primarily focused on analysing structured data to produce forecasts, classifications, or recommendations. Governance during this stage centred on model accuracy, data quality, and fairness. Risks were generally limited to incorrect outputs, such as inaccurate demand forecasts leading to overstocking or stock shortages. These risks were typically contained within defined operational boundaries and could be addressed through validation and retraining.

The rise of generative AI introduced a new layer of complexity. Unlike traditional systems, generative AI creates content such as text, images, and code. As a result, governance expanded beyond output accuracy to include content quality, reliability, and alignment with organisational standards. For example, AI-generated learning materials have occasionally contained incorrect information or misleading explanations. Such issues can affect user trust and outcomes, making content validation, human review, and controlled prompt design essential governance measures.

Generative AI has also created risks in specialised domains such as legal documentation. While these systems can assist with drafting contracts and policies, they may generate content that does not comply with legal requirements or organisational standards. This can expose organisations to compliance and operational risks. To address these challenges, businesses must implement domain-specific guardrails and ensure that subject matter experts review critical outputs before they are used in practice.

The emergence of agentic AI represents a further shift in governance requirements. Agentic systems can plan tasks, make decisions, and take actions with limited human intervention. In areas such as logistics, these systems can optimise schedules and routes using real-time data. However, incorrect decisions can have far-reaching consequences across operations and supply chains. The risk is no longer limited to inaccurate outputs but extends to actions performed directly by the system.

Agentic AI also raises important questions about accountability and control. When autonomous systems execute tasks independently, organisations must clearly define responsibility for outcomes and establish appropriate oversight mechanisms. Traditional governance controls alone are no longer sufficient. Effective governance now requires real-time monitoring, human intervention points,

escalation processes, and clearly defined decision boundaries. As AI systems become increasingly autonomous, adaptive governance frameworks are essential to maintain control, accountability, and trust while enabling innovation.

Video 4: Core Generative AI Risks

Generative AI has introduced powerful capabilities across enterprise functions, but it also brings important risks that organisations must manage carefully. Unlike traditional systems, generative AI can create content that appears accurate and convincing even when it is incorrect. This phenomenon, known as hallucination, occurs because the model generates responses based on patterns rather than verified knowledge. In enterprise environments, inaccurate outputs can affect decision-making, create compliance issues, and reduce trust in AI systems.

Another significant risk is bias amplification. Because generative AI models learn from large datasets that may contain historical biases, they can reproduce or even strengthen those biases in their outputs. For example, AI-assisted recruitment tools may favour certain candidate profiles based on patterns present in training data. Such outcomes create ethical concerns and may expose organisations to legal and regulatory risks. Regular bias testing, careful dataset selection, and ongoing monitoring are essential to minimise these issues.

Privacy leakage and intellectual property exposure present additional challenges. Generative AI systems may inadvertently reveal sensitive or confidential information contained in training data, creating security and privacy concerns. Similarly, AI-generated content can closely resemble existing copyrighted material, potentially leading to infringement claims and reputational damage. These risks are particularly relevant in industries that rely heavily on proprietary information and creative content.

These risks are often interconnected and can occur simultaneously, making effective governance critical. Organisations address them through measures such as prompt design standards, human review of critical outputs, continuous monitoring, and clear usage policies. Ultimately, accountability for AI-generated content rests with the organisation deploying the system, not the technology itself. By proactively managing hallucinations, bias, privacy concerns, and intellectual property risks, organisations can deploy generative AI in a way that is reliable, secure, and aligned with enterprise standards.

Video 5: Security & Misinformation Risks

As generative and agentic AI systems become more deeply integrated into enterprise workflows, they also increase exposure to security and misinformation risks. These threats often originate from external actors who exploit AI behaviour, user inputs, or connected systems. Unlike traditional cyber risks, AI-related threats can influence not only data security but also decision-making, operational processes, and organisational trust. Managing these risks is therefore a critical component of AI governance.

One of the most significant security threats is prompt injection, where attackers manipulate inputs to override a system's intended instructions. In AI-powered assistants, malicious prompts can cause systems to disclose sensitive information or perform unauthorised actions. Another major concern is data exfiltration through connected tools and applications. When AI systems are linked to databases, document repositories, or external platforms, weak access controls can allow users to retrieve information beyond their authorised permissions. These risks highlight the importance of input validation, access management, and strong system safeguards.

Misinformation at scale presents an equally serious challenge. Generative AI can rapidly create large volumes of realistic but misleading content, including fabricated articles, synthetic reviews, and manipulated narratives. Because such content often appears credible, it can be difficult to detect and may influence both public opinion and organisational decision-making. For enterprises, misinformation risks extend to internal knowledge systems, customer communications, and public-facing content, making ongoing monitoring and validation essential.

The impact of these risks is closely tied to brand reputation and customer trust. Incorrect, biased, or harmful AI-generated outputs can lead to public complaints, reputational damage, and loss of confidence. The risks become even greater in agentic AI systems, where manipulated inputs may trigger automated actions across connected environments. To address these challenges, organisations must implement layered controls such as prompt filtering, input sanitisation, access restrictions, audit logging, and content validation. Effective governance helps ensure that AI systems remain secure, reliable, and trustworthy while supporting safe enterprise adoption.

Video 6: Agentic AI Risks (Autonomy Layer Risks)

Agentic AI introduces a new level of capability by enabling systems to plan, make decisions, and execute actions with limited human intervention. Unlike traditional AI systems that primarily provide recommendations or generate outputs, agentic AI can take actions independently. While this improves efficiency and scalability, it also introduces autonomy-related risks that require careful governance and oversight.

One of the most significant risks is autonomous action failure. Agentic systems rely on goals, rules, and contextual information to make decisions. If the underlying logic or data is flawed, the system may take incorrect actions repeatedly and at scale. Because these systems can operate continuously without waiting for human approval, mistakes can quickly escalate and affect critical business operations.

Another key challenge is the breakdown of coordination between multiple AI agents. In complex enterprise environments, different agents may be responsible for planning, decision-making, and execution. If they operate using inconsistent data or conflicting assumptions, the overall workflow can become misaligned. Such coordination failures can lead to operational inefficiencies, financial losses, and unpredictable system behaviour.

Goal misalignment presents an additional risk. Agentic systems optimise for the objectives they are given, but poorly defined goals can result in outcomes that conflict with broader organisational priorities. A system may successfully achieve a target while creating unintended negative

consequences, such as compromising quality, customer trust, or brand reputation. This highlights the importance of defining objectives alongside appropriate constraints and business context.

The absence of effective human oversight can further amplify these risks. To manage them, organisations must establish clear decision boundaries, embed human review at critical points, and implement continuous monitoring and escalation mechanisms. Strong governance ensures that autonomous systems remain aligned with organisational goals, operate responsibly, and deliver value without compromising accountability, trust, or operational control.

Video 7: EU AI Act & Risk Classification

The EU AI Act establishes a structured, risk-based approach to AI governance, helping organisations manage AI systems according to their potential impact. Rather than applying the same requirements to every application, the Act classifies AI systems into different risk categories. This framework enables enterprises to address compliance, accountability, and governance requirements in a proportionate manner while promoting responsible AI adoption.

At the lower end of the framework are minimal-risk and limited-risk systems. Applications such as spam filters and recommendation engines face relatively few obligations, while user-facing systems like chatbots must clearly disclose that users are interacting with AI. These transparency requirements help build trust and ensure that AI-enabled interactions are clear and understandable.

The most significant obligations apply to high-risk AI systems used in areas such as recruitment, credit assessment, healthcare, and critical infrastructure. Because these systems can influence employment opportunities, financial services, and access to essential resources, organisations must ensure data quality, conduct bias testing, maintain detailed documentation, provide human oversight, and continuously monitor performance. These measures help reduce risks and improve accountability.

The EU AI Act also affects how AI systems are designed and deployed. Governance requirements such as transparency, auditability, data governance, monitoring, and human intervention must be considered from the earliest stages of development. Before implementation, organisations need to assess the risk level of each application and apply the appropriate controls. By aligning governance requirements with risk levels, the Act supports responsible AI deployment while maintaining compliance, trust, and regulatory alignment.

Video 8: NIST AI RMF & OECD Principles

The NIST AI Risk Management Framework and the OECD AI Principles provide organisations with structured guidance for managing AI risks and ensuring responsible deployment. As AI systems become more complex and widely used, these frameworks help organisations establish consistent governance practices. Together, they support the development of AI systems that are trustworthy, reliable, and aligned with both organisational objectives and societal expectations.

The NIST AI Risk Management Framework follows a lifecycle-based approach built around four core functions: govern, map, measure, and manage. These functions help organisations establish accountability, understand AI use cases and risks, assess system performance, and implement appropriate controls. A key feature of the framework is its emphasis on continuous monitoring, recognising that AI systems evolve over time and require ongoing oversight rather than one-time validation.

This approach is particularly valuable in applications such as fraud detection and predictive maintenance. In digital payments, continuous monitoring helps identify emerging fraud patterns and detect model drift. In manufacturing, regular assessment of AI models ensures that systems remain effective as operating conditions change. By applying structured risk management throughout the AI lifecycle, organisations can improve reliability and reduce operational disruptions.

Complementing the NIST framework, the OECD AI Principles provide an ethical foundation centred on fairness, transparency, accountability, robustness, safety, and human-centred design. While NIST focuses on operational risk management, the OECD principles define the values that responsible AI systems should uphold. Together, these frameworks enable organisations to build effective governance programmes, manage AI risks proactively, and deploy generative and agentic AI systems with greater confidence, trust, and control.

Video 9: Final Thoughts

AI governance is no longer optional. As AI systems evolve from providing assistance to taking autonomous actions, the associated risks and responsibilities increase significantly. Organisations deploying generative and agentic AI at scale must establish governance frameworks that promote responsible, secure, and trustworthy AI adoption. Ethical principles such as fairness, transparency, accountability, and safety should guide every stage of the AI lifecycle to build trust and reduce risk.

We also explored how governance requirements change across different types of AI. Generative AI introduces risks including hallucinations, bias, misinformation, data leakage, and intellectual property concerns. Agentic AI adds further complexity because autonomous systems can make decisions and take actions independently, increasing the potential impact of errors. Without appropriate oversight and controls, these risks can affect operations, compliance, customer trust, and organisational reputation.

Finally, we examined key frameworks such as the EU AI Act, the NIST AI Risk Management Framework, and the OECD AI Principles, which provide structured approaches to managing AI-related risks. These frameworks help organisations assess risks, implement controls, ensure accountability, and maintain compliance. The key takeaway is clear: responsible AI requires proactive governance rather than reactive fixes. Organisations that embed governance into the design, deployment, and monitoring of AI systems will be better positioned to build trust, maintain control, and create sustainable long-term value.